

A Web-Based Weather Service for Wind Sports

Jakob Aarøe Dam

Supervised by Olivier Danvy

Master's Thesis



3 July 2009
Department of Computer Science
Aarhus University
Denmark

Abstract

We describe the foundations, design, and implementation of a mashup that assists practitioners of wind sports. The problem consists of three parts: (1) obtaining weather forecasts and weather observations from weather resources; (2) creating Web Services that serve representations of relevant data; and (3) creating a Web Service client that presents the data in a comprehensible matter for users.

We present the theoretical foundations for creating scalable Web Services. The theoretical foundations consist of Roy Fielding's architectural style, Representational State Transfer (REST), and of Leonard Richardson and Sam Ruby's concrete Web Service architecture, the Resource Oriented Architecture that is subject to the architectural constraints of REST.

We use the Google App Engine (GAE) as the platform for the Web Service. Using GAE for geographical spatial data demands that space-filling curves are applied to map from multi-dimensional data, e.g., latitude and longitude, to a one-dimensional value that can be indexed by GAE. Geohash is presented and used in a scalable solution that serves location-based information on a proximity basis. In addition, we present a solution that circumvents the geohashing limitations regarding proximity searches across boundaries.

We illustrate how the Web Service is used by describing the design and implementation of a location-based Ajax front-end that presents weather information relevant in the context of wind sports.

Acknowledgements

This last half year creating my ‘Speciale’ has been a fantastic ride uniting two of my passions: computer science and kite surfing. I am grateful to Olivier Danvy for making it possible by seeing potential and believing in my thesis, even before I did myself, and for giving keen advice and comments on the dissertation.

I owe a great deal of thanks to Brian Jensen and Jens Rimestad for proof reading this dissertation. I also thank them and Jacob Sloth Mahler-Andersen for fruitful visions, comments, and feedback on the design and functionality of this weather service for wind sports and, for many great hours kite surfing on the water, however rarer these hours have become lately.

I owe great thanks to Michael Glad, for setting up the infrastructure to run a Web Service here at DAIMI and for bringing this infrastructure back online outside working hours and with an outstanding celerity when a software update caused it to fail; to John Kruuse for conjuring up the otherwise hidden article “Virtues of Haversine” from the archives of Statsbiblioteket; to Gerth Stølting Brodal and Lasse Deleuran for listening to me talking about geohash and for pointing me towards space-filling curves as the natural generalization of geohash; and to Mogens Dam for always being helpful and for expertly introducing me to spherical geometry.

I am also thankful to the providers of weather data: the Danish Coastal Authority, the Danish Meteorological Institute, the National Oceanic and Atmospheric Administration, and WeatherBug; without them my thesis would have been unrealistic and unprovable.

Finally, thanks to my family for their never-ending love and support, and to my girlfriend for her support, for bearing with my long hours of work, and for bringing into my life happiness, love, and perspective.

*Jakob Aarøe Dam,
Århus, July 3, 2009.*

Typographical conventions

The following conventions are used in this dissertation:

- In each chapter, *italic* fonts emphasize words or mark terms that are described in the glossary. The first occurrence of each such term is in *italics*. The subsequent occurrences are not emphasized.
- **teletype** fonts indicate computer code either inlined in the text or in listings; code is everything from concrete program code to urls and file names. In listings (...) indicates that code is omitted for brevity. In addition, most Python imports are left out for brevity.

Sometimes in listings we reproduce a transcript of a session with the Python interpreter, as in Listing 1.

```
>>> import geohashneighbors as ghn
>>> import geohash
>>>
>>> london = geohash.Geohash((0, 51))
>>> str(london)
'u1040h2081040'
>>> ghn.neighbors('u10')
['u12', 'u13', 'u11', 'u0c', 'u0b', 'gbz', 'gcp', 'gcr']
```

Listing 1: *Transcript of a Python session*

Contents

Abstract	iii
Acknowledgements	v
1 Introduction	1
1.1 The Problem	2
1.2 Related Work	4
1.2.1 Geographic Information Systems	4
1.2.2 Web Technologies	6
1.2.3 A Similar Web Application	6
I The Foundations	7
2 Web Services	9
2.1 Web Services	9
2.2 Software architecture and style	10
2.3 REST	10
2.4 ROA: A RESTful architecture	11
2.4.1 Addressability	11
2.4.2 Statelessness	12
2.4.3 Uniform interface	13
2.4.4 Connectedness	15
2.5 Caching	16
2.6 Summary	17
3 Web Service clients	21
3.1 Ajax: the style for modern Web Service clients	21
3.2 JSON	22
3.3 Summary	23
4 Google App Engine	25
4.1 What is the Google App Engine?	25
4.2 An example for the impatient	26
4.3 Architecture	27
4.3.1 Model	28
4.3.2 View	30
4.3.3 Controller	30
4.4 Putting it together	31
4.5 Further GAE prerequisites	31
4.5.1 Unique Properties	32
4.5.2 Pagination	33
4.5.3 Timezones	34
4.6 Summary	36

II	The Web Services	37
5	Resources and their Representations	39
5.1	Spots	39
5.2	Weather Stations	41
5.3	Forecast Points	43
5.4	Sorting out the resource requirements	45
5.4.1	Calculating proximity points with haversine	45
5.5	Summary	55
6	The We Love Wind Web Service	57
6.1	Implementation of Models	57
6.1.1	AbstractGeoHash model	57
6.1.2	Spot model	59
6.1.3	AbstractWeatherData Model	61
6.1.4	WeatherStation Model	62
6.1.5	ForecastPoint model	64
6.2	Implementation of Controllers	66
6.2.1	REST controller	67
6.2.2	Spots controller	70
6.2.3	Weather stations controller	72
6.2.4	Observations controller	72
6.2.5	Forecast points controller	73
6.3	Implementation of Views	74
6.4	Enhancing Performance	74
6.5	Summary	76
7	The DAIMI Forecast Web Service	77
7.1	Model	78
7.2	Controller	78
7.3	Summary	79
III	The Clients	81
8	Weather data	83
8.1	Getting weather data	83
8.2	Weather Observations	84
8.2.1	Scraping DMI and DCA	84
8.2.2	Scraping WB	85
8.2.3	Integrating with the Web Service	88
8.3	Forecasts	89
8.3.1	Retrieving forecasts	89
8.3.2	Extracting data from GRIB2 files	90
8.3.3	Integrating with the Web Service	91
8.4	Summary	92
9	Main user interface	93
9.1	Design	93
9.2	Finding the location of the client	94
9.3	Reducing the data set	96
9.4	Seamless points loading	97
9.5	Coordination between Ajax calls	98
9.6	Drawing a circle with Google Maps	100

9.7 Summary	102
10 Interface for mobile devices	103
10.1 Design	103
10.2 Sorting out the resource requirements	104
10.2.1 Location-based mobile services	105
10.2.2 Content Negotiation	105
10.2.3 Local time	107
10.2.4 Geocoding	108
10.2.5 Displaying static Google maps	109
10.3 Implementation	110
10.3.1 Controllers	110
10.3.2 Map Controller	111
10.3.3 Info controller	112
10.3.4 Views	113
10.4 Summary	113
 IV Conclusion & Perspectives	 115
 V Appendix & Glossary	 119
 Glossary	 122
 A Appendix A	 123
A.1 Rejseplanen.dk screenshots	123
A.2 Cross-site scripting with eval	124
 B Appendix B	 127
B.1 Implementation of Views for Mobile Devices	127
B.1.1 Base mobile template	127
B.1.2 Search template	128
B.1.3 Info template	128
B.1.4 Map template	130
 C Appendix C	 131
C.1 Comparing Distances Calculated with Haversine and Cosines	131

"Everybody talks about the weather, but nobody does anything about it."

Mark Twain (1835 – 1910)

1

Introduction

It is our thesis that online weather data can be harvested and combined with new data in a mashup to assist the practitioners of wind sports. Such a mashup is a global weather service that provides users with relevant wind information tailored for the individual users. In particular, this means that 1) wind data is synthesized and transformed to weather information in a way that is relevant for wind sports; 2) only wind information relevant for the user's location is served; and 3) wind information is accessible through the means of its user such as a standard computer or a cell phone.

A mashup is commonly defined as follows:

Definition 1.1 *"A web mashup is a web page or application that combines data from two or more external online sources." [?]*

This dissertation describes the foundations, design, and implementation of such a mashup, We Love Wind, for wind sports. We split the problem of constructing the We Love Wind mashup up into three parts:

1. We give the foundations for developing mashups. This consist of foundations for the server-side of mashups manifested as RESTful Web Services (Chapter 2), the client-side of mashups manifested as RESTful Web Service clients (Chapter 3), and an infrastructure for running the Web Service, the Google App Engine (Chapter 4).
2. We design the resources of the Web Service (Chapter 5) and build two RESTful Web Services. Because of restrictions in the Google App Engine (GAE) the Web Service consists of two Web Services: one located at the GAE (Chapter 6), and one located at the Department of Computer Science, Aarhus University (Chapter 7).
3. Finally, we create clients of the Web Services that maintains weather data in the Web Services (Chapter 8), presents data from the Web Service apt for standard computers (Chapter 9) and mobile devices (Chapter 10).

When all the three parts are done we have integrated all the Web Services shown in Figure 1.1 and arrive at the goal: a mashup that assists practitioners of wind sports.

A prototype of the mashup is running at:

- <http://welovewind.com/> for standard computers; and
- <http://m.welovewind.com/> for mobile devices.

Prerequisites: The reader is expected to be somewhat familiar with Python, JavaScript, HTTP and UML.

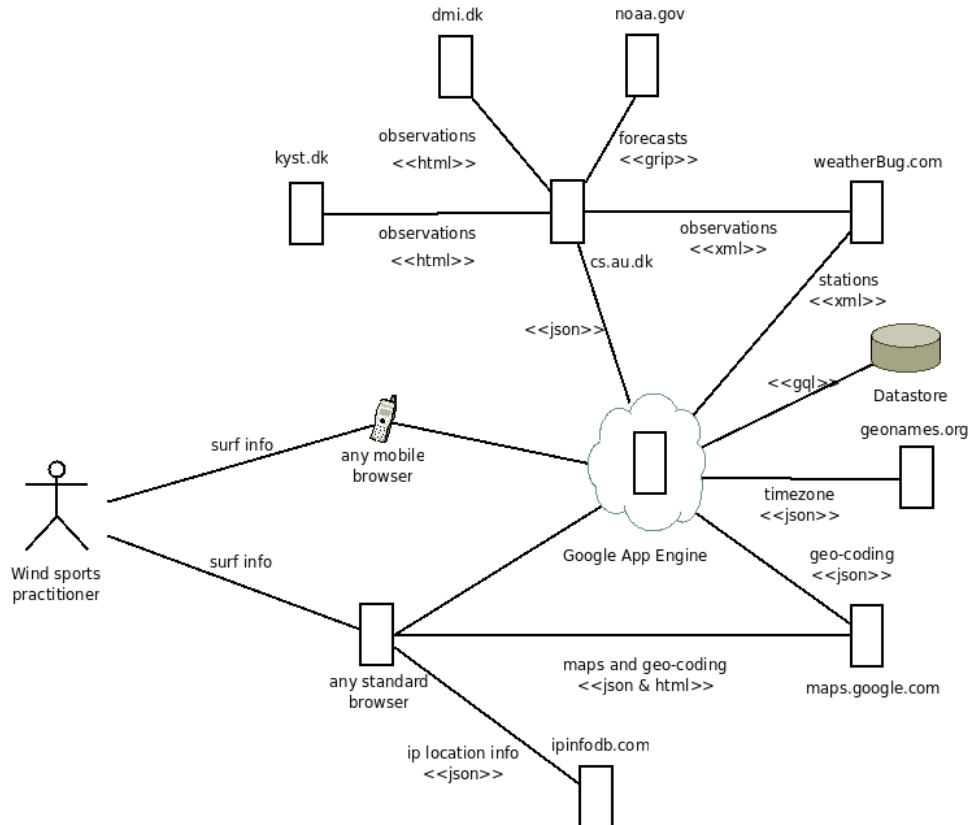


Figure 1.1: Architectural Overview

1.1 The Problem

This section gives an overview of the problem and the problem domain by presenting the potential users of the system and the task they currently perform. Finally, we present three scenarios which are visions of how the mashup will assist practitioners of wind sports

Users

Wind- and kitesurfers are users of the application. Little statistic data exist about this segment of the population. However, a 1994 survey [?, p.27] shows that there are in the proximity of one million windsurfers in the US, where 59.6 per cent are male, and 56.6 per cent of the windsurfers are in the age interval 25 - 44. An article from 2004 speculates that windsurfing is a growing sport with over 20 million windsurfers worldwide mostly between 15 and 40 years of age.¹

¹<http://www.westcountrynow.com/main/articles/display.cfm?r=0.15720977&ref=417>

Tasks

When wind- and kitesurfers are on call for surfing they continuously check available weather data for favorable weather conditions. Favorable conditions mean suitable wind speed in a suitable direction. The suitable direction is defined by the location of the surf spot (a suitable place for wind- and kitesurfing). It is dangerous to surf in offshore wind conditions, unless the surfer is accompanied by a boat to fetch the surfer drifting to the sea. Suitable wind speed is about 6m/s and above. The users rely on two types of weather data: forecasts and observations. The two types of data are located at different sources.

Forecasts: The currently most detailed source for wind forecasts in Denmark is the wind service of the Danish Maritime Safety Administration² (DMSA). The page, shown in Figure 1.2, contains a map with a wind forecast for the current hour in the form of arrows and colors that indicate wind velocity. The page provides options to access forecasts for the following hours up to 48 hours in the future from the time of forecast calculation. The calculations are started every six hours.

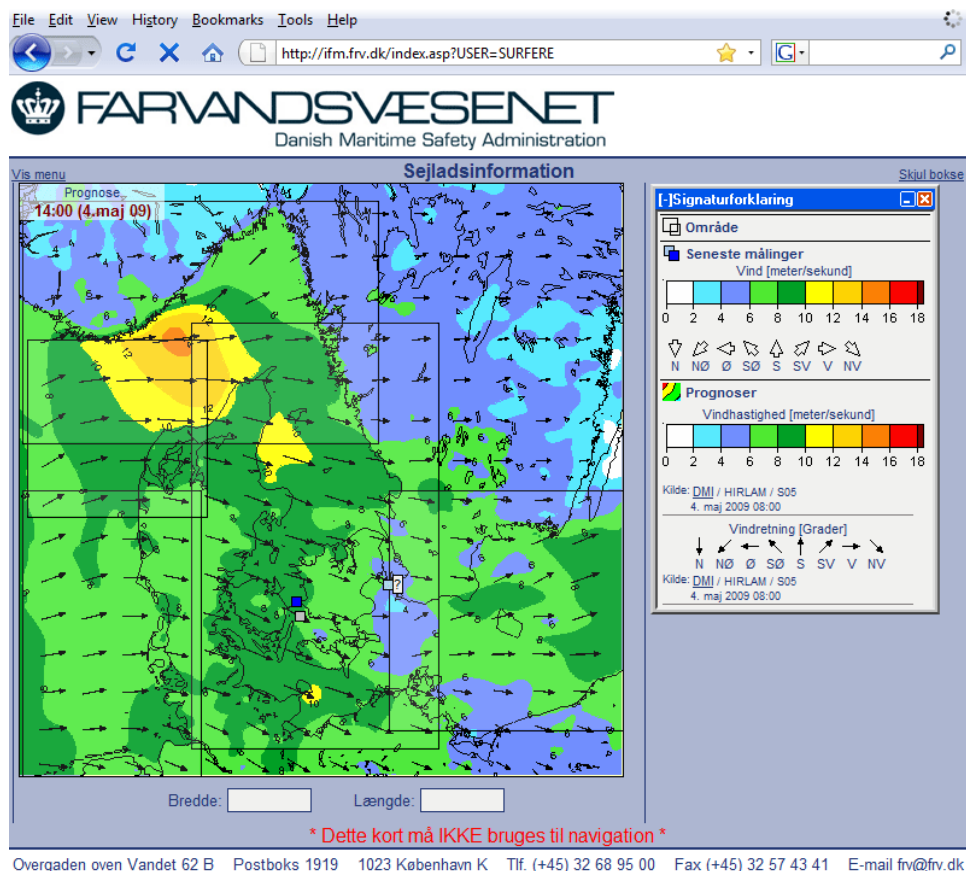


Figure 1.2: DMSA wind forecast page

The task users execute with DMSA's service is mapping between the colors and arrows on the page, and relevant surf spots known to the surfer. Surfers execute that task at least one time every day when they want to surf.

²<http://ifm.frv.dk/index.asp?USER=SURFERE>

Observations: Sources for current wind conditions are the Danish Coastal Authority³ (DCA), the Danish Meteorological Institute⁴ (DMI), and WeatherBug⁵ (WB).

The service at DMI⁶ shows a map with some of the weather stations and their current measurements. Detailed information about a weather station is accessed from a standard browser by clicking on it. The detailed information includes a graph of the wind measurements for the last 24 hours. Users map between the conditions at weather stations and spots to check if the current conditions allow for surfing. This can include an investigation of the graph to see whether the wind is picking up or slacking down.

Scenarios

We now present real world scenarios illustrating the benefit of using the We Love Wind mashup (scenarios are small stories that tell what people do with the Web site and not how they do it).

Scenario 1 *Jacob has just got off work. He wants to know if he can surf in the proximity.*

Jacob accesses the We Love Wind site from his mobile. The site shows that currently and the next 4 hours it is possible to surf at a surf spot near Aarhus, he decides to go.

Before rigging his kite Jacob accesses the page again and checks the current observations to get a feeling of what size of kite he should use.

Scenario 2 *One morning, Jens wants to know whether the day is going to be a surf day. He accesses the We Love Wind site from any device (mobile device or standard computer). The page tracks his location and the page shows tailored spot forecasts for the area. He looks at the forecasts and decides to plan for surfing in the afternoon.*

Scenario 3 *Brian is on holiday in Skagen. He accesses the We Love Wind site to check out the forecasts and the current observations for surf spots near Skagen. The site does not have a surf spot for his location.*

In order to get observations and forecasts for the surf spot he decides to input a new spot into the site. Afterwards, the site shows observations around the spot and wind forecasts for the surf spot.

1.2 Related Work

Working with mashups is a cross-disciplinary area. Mashup construction includes concepts and technologies from a wide and diverse area. This includes Web technologies and geographic information systems. The following sections present the important related work in those categories. In addition, we present another mashup similar in concept to the We Love Wind mashup.

1.2.1 Geographic Information Systems

Geographical information systems are systems that associate geographical points with information.

An important aspect of these systems is serving geographic data effectively. Geographic data are multi-dimensional – it typically has a latitude and a longitude

³<http://www.kyst.dk>

⁴<http://www.dmi.dk>

⁵<http://www.weatherbug.com/>

⁶<http://www.dmi.dk/dmi/index/danmark/borgervejrr.htm?param=wind&map=undefined>

and maybe an altitude – and therefore usual one-dimensional database indices are inappropriate [?, p.173].

A space-filling curve is a curve that parses through each point in a square (it can be generalized to any dimension). Space-filling curves map values from a multi-dimensional space to a one-dimensional value. Space-filling curves therefore support one-dimensional indexing methods. Examples of space-filling curves for multi-dimensional indexing are the Hilbert space-filling curve, and the z-order [?, p.199].

Figure 1.3 shows a z-order space-filling curve. Cells in the grid are recursively assigned a one-dimensional value, called the z-value, that indicate its location in space. The z-value is calculated by cyclically taking one bit from each coordinate (in the figure, x and y respectively) of the cell and appending that bit to the z-value [?, sec.9]. Thus, the coordinates are interleaved in the z-value. Decoding a z-value to a point in two dimensions goes as follows: given a z-value such as 11111₂ (in the figure the cell most South-East), the space is recursively divided in two starting from the left in the bit string; even bits pertain to the y-axis, odd bits pertain to the x-axis.

The values of a space-filling curve have the property that the proximity of points in space is preserved to some extent [?, p.199], and that usual database indices, such as B-trees [?], can index the values of the space-filling curve since they are total ordered.

	x:	0	1	2	3	4	5	6	7
		000	001	010	011	100	101	110	111
y: 0	000	000000	000001	000100	000101	010000	010001	010100	010101
1	001	000010	000011	000110	000111	010010	010011	010110	010111
2	010	001000	001001	001100	001101	011000	011001	011100	011101
3	011	001010	001011	001110	001111	011010	011011	011110	011111
4	100	100000	100001	100100	100101	110000	110001	110100	110101
5	101	100010	100011	100110	100111	110010	110011	110110	110111
6	110	101000	101001	101100	101101	111000	111001	111100	111101
7	111	101010	101011	101110	101111	111010	111011	111110	111111

Figure 1.3: Z-values for two dimensional z-order curve [?]

Geohash is an algorithm to convert between latitude and longitude coordinates and a hash value. The term geohash is un-described in the literature the only source being a Wikipedia article [?]. Geohash values have the essential property that points in the proximity of each other often have similar geohash prefixes; therefore, a database can benefit from a string index on the geohash value of points when performing searches for points in the proximity of each other.

Geohash is reminiscent to z-order space-filling curves. Geohash divides space in 32 squares recursively; z-order divides space in 4 squares recursively. The interleaving of two-dimensional points in z-values is the same scheme applied in geohash.

Geohash is used in [?]; this implementation includes an algorithm that returns adjacent locations in the grid. The approach is un-documented, the implementation being the only reference.⁷

[?] presents a technique to find adjacent location in grids that takes time for the worst case proportional to the length of the encoded string. [?] have improved the algorithm to find adjacent neighbors in $O(1)$, constant time, in the worst case. The techniques were immediately applicable for geohash if the number of squares – i.e., chars in the base string – in geohash were a power of four.

1.2.2 Web Technologies

Mashups rely on Web technologies which is an ever evolving area. In this dissertation we use a subset of Web technologies: Hypertext Transfer Protocol (HTTP), Hypertext Mark-up Language (HTML), ECMAScript (JavaScript), JavaScript Object Notation (JSON), and Google App Engine (GAE) to name a few. We introduce the technologies on the go in the dissertation.

1.2.3 A Similar Web Application

windguru.com is a weather service specialized for wind sports. The site uses global weather forecast data, and finer detailed data for North America, and Europe. windguru.com is interesting since it is very detailed and enable users to setup customized pages with forecasts of their own choice in terms of location and forecast model.

⁷Unfortunately, the available example uses JavaScript functions that are not widely supported. The code uses `Array.prototype.indexOf` which is not defined in the ECMAScript Language Specification. As a result many browsers are unable to execute the page. The author has, however, successfully opened the page in Google Chrome and Opera 9.

Part I

The Foundations

2

Web Services

In this chapter we lay out foundations for the architecture of Web Services. We introduce the Resource Oriented Architecture (ROA): a concrete architecture for Web Service that is constrained by REST. REST builds on the concepts of software architecture and architectural style; in order to define and understand ROA and REST it is necessary to define these two concepts first. We introduce the necessary terminology for Web Services, Software Architecture, REST, and ROA. Finally, we give a short intro to HTTP caching.

2.1 Web Services

Definition 2.1 *A Web Service is a software system that support machine-to-machine interaction over a network via HTTP.*

The definition is a relaxed version, removing the requirements to use *WSDL* and *SOAP*, of the W3C definition [?]. In addition, *support* is interchanged with *designed for* making any Web site accessible by computers also fall into the category of Web Services.

The Web Services specifically designed for machine-to-machine interaction typically fall into one of the two categories: RESTful Web Services and Web Services based on WSDL, SOAP, and the *WS-* stack* – known as *Big Web Services*. Creating just the smallest Web Service with Big Web Services demands an excessive amount of plumbing. In a Big Web Service implementation several things must be defined: the message formats (in SOAP), data type definition (in XML Schema), and an interface definition (in WSDL). The plumbing is due to creating a tight contract forcing everything into a standard.

There, however, already exist a standard that has the plumbing in place, Web Services based on REST with HTTP and Uniform Resource Identifiers (URI). Web Services based on REST have the great advantage that the means to work with them are simple. On the client side, the Web Services are accessed with existing browsers; on the server side, the Web Services are implemented with usual technologies to create dynamic Web sites. There are scenarios where a RESTful approach probably won't do the job, such as a distributed transaction protocol [?]. A Web Service under the constraints of REST, however, is the approach with the lowest entry barrier,

and the approach that work with and not against the nature of the Web, therefore, we decide to travel light, with REST.

2.2 Software architecture and style

Definition 2.2 *A software architecture is the structure(s) “of a computing system, which comprise software elements, the externally visible properties of those elements, and the relationships among them” [?, p.21].*

The *de facto* definition of software architecture, Definition 2.2, tells that software architecture is first and foremost an abstraction of the system. The focus is not on internal properties, but on visible aspects of the system and how the components in the system cooperate.

An important thing about the software architecture is that it is a basis for archiving quality in the system it pertains to. A software architecture always facilitates and prohibits different *quality attributes*. Ensuring certain quality attributes in the system is done by constraining the architecture in ways known to affect these; usually facilitated by the use of architectural styles. We define an architectural style, based on the definition in [?, p.24], as follows:

Definition 2.3 *An architectural style is a general solution to a problem that describes element and relation types and constraints on their use in the software architecture.*

The most important aspect of architectural styles (styles from now) are that an architecture that satisfy a style exhibit well known quality attributes imposed by the style [?, p. 25].

2.3 REST

REST is a style designed for systems that handle *distributed hypermedia* [?]. REST defines a set of architectural elements and constraints on the use of them that facilitate quality attributes relevant in a context of distributed hypermedia.

REST is the sum of applying several styles for network-based architectures; Table 2.2 on page 19 summarizes the styles; the summary includes a short description of the styles and their effect on a system in terms of quality attributes.

The most characteristic style for REST is uniform interface. It defines several architectural elements and constraints on the use of them. A pivotal architectural element of the uniform interface is the resource.

Definition 2.4 *A resource is any information that can be named [?, p.125].*

Resources are uniquely identified by a resource identifier. In terms of the Web, a resource is any information identified by a *URI*. Clients that request a resource get a representation of the resource – not the resource itself – in a format depending on the capabilities of the client. Interaction with resources are performed through their representations by passing the representations around in messages and applying the uniform interface on them. Messages are self-descriptive, which decouples them from the underlying transport layer; enabling intermediate servers to cache messages, forward messages, etc.

REST stands for Representational State Transfer. The following quote explains the meaning of it:

The name “Representational State Transfer” is intended to evoke an image of how a well-designed Web application behaves: a network of Web pages forms a virtual state machine, allowing a user to progress through the application by selecting a link or submitting a short data-entry form, with each action resulting in a transition to the next state of the application by transferring a representation of that state to the user. [?, p.116]

An architecture that follows the architectural style of REST is called RESTful. A RESTful architecture exhibits well known quality attributes chosen by the designers of REST. Since REST is a model of how the Web *should* work [?, p.116], a RESTful architecture will facilitate quality attributes deemed relevant for the Web.

2.4 ROA: A RESTful architecture

A concrete RESTful architecture is the Resource-Oriented Architecture [?, ch.4]. ROA is designed specifically for Web applications; ROA therefore shows how REST is applied using the intrinsic technologies of the Web, such as HTTP and URI. ROA hereby turns the constraints of REST concrete. By using HTTP many of the styles from REST are already satisfied:

- Client-server: in HTTP only the client is able to make requests, while the server is ready to receive them.
- Cache: HTTP provides several caching mechanisms, presented in Section 2.5.
- Code on demand: via HTTP it is possible to transfer programs, e.g., JavaScripts or Java applets, to the client and execute them locally.
- Layered system: HTTP is the medium for client-server interaction; what technology, platform, etc., the server is based on is hidden for the client.

That many of the styles of REST are already satisfied comes as no surprise: REST is the style for HTTP and URI. ROA ignores the styles a Web application cannot affect and instead focus on only four styles; each is described in the following sections.

Throughout this section we examine an existing Web application to give examples of violations of the styles and their consequences. The examples illustrate some of the advantages of a RESTful Web Service by examining the inverses. The inverses are so common that we denote them as *anti-patterns*. The Web service examined is `rejseplanen.dk`; a journey planning Web Application for bus and train trips in Denmark.¹

2.4.1 Addressability

Since ROA is based on REST it is centered around the resource architectural element. In ROA all interesting information provided by the server is a resource. Resources are identified by a URI with a reasonable name reflecting the concept of the resource [?, Ch.4]. Data on the server not exposed by a URI, are not resources. This constraint, known in ROA as addressability, gives several advantages:

- The provider of data is separated from the consumer of data which decouples the application. The decoupling introduces support for arbitrary composition of the resources at the client side – known as mashups.
- Caching mechanisms in HTTP have a larger basis to work on, because they have a larger contact area with the Web application.
- Clients can link to and bookmark specific pages in the Web application.

¹`rejseplanen.dk` uses HAFAS planning software which is in use in 16 countries <http://www.hacon.de/hafas/>

Addressability anti-pattern

Some Web applications do not expose their interesting information through URIs. An example of a web application that violates addressability is `rejseplanen.dk`.

Journey plans from `rejseplanen.dk` is shown in Figure 2.1. We arrived at the page after filling out a search form on the front page, entering the information that we want to go from Aarhus to Copenhagen, and submitting the form.

In the figure we notice that the URI of the journey plan page is unchanged when compared to the main page. This makes it impossible to link to the journey plan, since the URI in the address bar does not reflect the page currently visited; therefore it is not addressable. The only address ever exposed in the address bar is the URI of the main page. The reason is that the page uses frames. From a usability perspective, the use of frames was discouraged already in 1996, because it destroyed addressability [?]. The addressability violation has a severe impact on the application: bookmarking the journey plans page is disabled, and the provider and consumer of data are coupled.

What if the frames are removed? The site expose a search resource, namely the search engine interface, and not the interesting data behind it. This is contrary to similar search services, e.g., Google that exposes their data. A Google URI like `http://www.google.com/search?q=addressability` returns a representation of the directory of the 10 most relevant pages containing the term addressability.

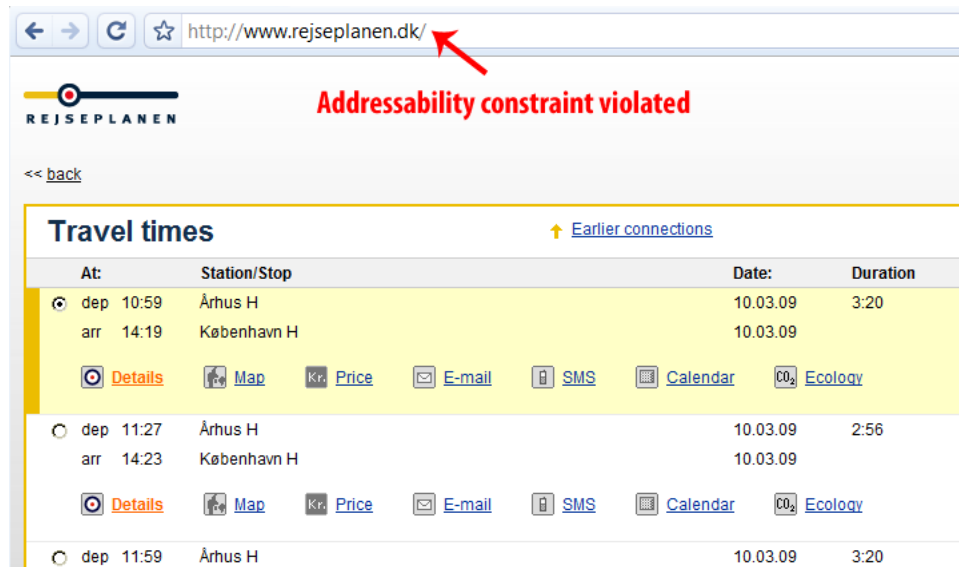


Figure 2.1: Addressability violation; [accessed 11-March-2009]

2.4.2 Statelessness

The style of statelessness is enforced in ROA by using HTTP. By design HTTP is stateless; a workaround is needed to introduce state. An example of this is cookies [?, p.145]. Cookies usually contain a hashed value referring to a data structure on the server containing *application state*. Since application state is saved on the server the stateless style is violated. Cookies used in this way circumvent the quality advantages introduced by statelessness since servers must use resources to store client state and replicate state across servers in a clustered environment.

Statelessness anti-pattern

Many Web applications use cookies to support session state, e.g., [rejseplanen.dk](http://www.rejseplanen.dk). Storing application state for each user demands an excessive amount of resources on the server. Resources are regained by garbage collecting session data structures after a short amount of idle time – typically less than 20 minutes. Re-accessing a site after the 20 minutes idle time causes an error because the data structure is removed; Figure 2.2 shows what happens when we re-access the journey plans: it causes a failure degrading usability.

In the previous section we criticized the page for violating addressability, but now it makes sense that the page is not addressable. The statelessness violation totally removes the option to make the application addressable since all pages are invalidated after a short amount of time.

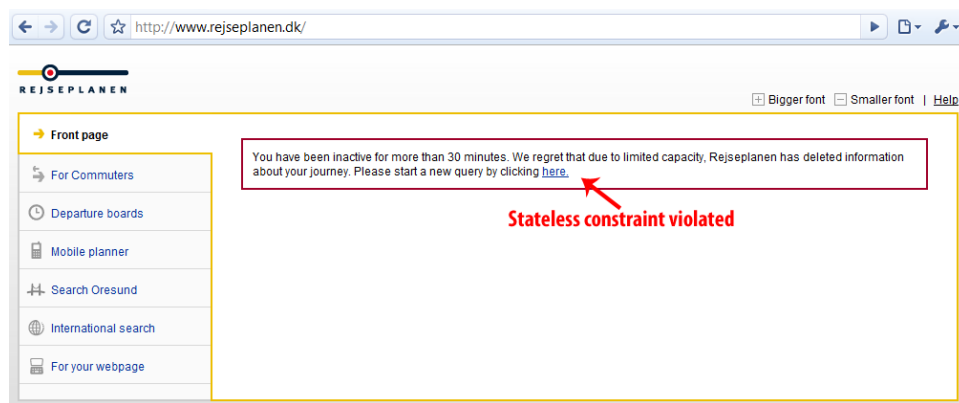


Figure 2.2: *Stateless violation; [re-accessed 11-March-2009]*

2.4.3 Uniform interface

Clients of RESTful Web Services interact with a resource, through HTTP and the URI of the resource. ROA demands the proper application of the HTTP methods [?, ch.4]. HTTP defines eight methods that can be performed on a resource each with well defined semantics [?]. The semantics of the important methods are given in Table 2.1; it is indicated whether the methods are idempotent and whether they are safe. Safe and idempotence are terms defined in [?]. Safe denotes the property that a method is free of side effects that the user can “be held accountable for”. Idempotence denotes the property that “the side-effects of $N > 0$ identical requests is the same as for a single request.” The definition of idempotence is unclear unless side effects are interpreted as the side effects the user can be held responsible for, e.g., GET is an idempotent operation but in many cases a GET request changes the server state by writing to the server log, causing some caching, incrementing a page counter etc. Interpreting side effects as user responsible side effects allow side effects, such as incrementing a counter and writing to the log, and does not violate the property of idempotence; in addition the safe methods become a proper subset of the idempotent methods.

There exist a problem in supporting the uniform interface of HTTP on the Web. (X)HTML forms only support two HTTP methods: GET and POST [?, sec.17]. A workaround is to overload the POST method, by adding, e.g., a `_method` parameter to the query parameters, and route to the corresponding implementation on the server. To PUT or DELETE a resource from a Web browser a form then needs to POST to the following URIs respectively:

```
{resource_uri}/?_method=PUT  
{resource_uri}/?_method=DELETE
```

HTTP Method	Semantics
GET	Retrieve the representation of a resource pointed to by the URI (safe).
HEAD	Same as GET except the response only consist of header fields (safe).
PUT	Create or update the resource identified by the URI in the HTTP message, with the representation of the resource located in the body of the HTTP message (idempotent).
DELETE	Delete the resource identified by the URI in the HTTP message (idempotent).
POST	Post the data located in the body of the HTTP message as a new subordinate of the resource identified by the URI in the HTTP message (un-idempotent).

Table 2.1: *HTTP Methods and their semantics* [?, sec.9]

Uniform interface anti-pattern

The query interface of `rejseplanen.dk`, which consists in an HTML form, is an example of an illegal use of the HTTP methods in terms of their semantics. The form uses `POST` to send the query, however, it makes no sense for the server to create a subordinate resource, in reaction to the request.

Browsers notice the use of `POST` and some, e.g., Google Chrome and Safari, ask for a confirmation when using back / forward buttons, ruining the user experience, as shown in Figure 2.3. The warning is perfectly valid if `POST` is used correctly, namely for posting some data at the server. In this case, however, it is a misuse of the method; the method is safe doing nothing besides retrieving data and therefore `GET` should be used.

Another unfortunate side effect of the misused `POST` method is that, by definition, clients must invalidate their cached representation when `POST` is used [?, sec. 13.10], taking away the advantage of caching.

There are two probable explanations for the illegal use:

1. Web crawlers refrain from taking the search link because it uses `POST`; reducing use of resources on the server. `rejseplanen.dk`, however, has disabled indexing of all their site, see Appendix A.1; therefore, resource consumption by friendly bots are not a problem.
2. Internet Explorer imposes a maximum length on URIs to 2,083 characters² although no limit exists in the HTTP specification [?]. The limitation does not exist for `POST`, since values in a `POST` message does not go into the URI.

The length of the encoded form is 1562 characters, This is without the from-city, the to-city, and a optional via-city. Thus there are over 500 characters left to encode only 3 places, surely enough, also rendering this explanation invalid.

²<http://support.microsoft.com/default.aspx?scid=KB;en-us;q208427>

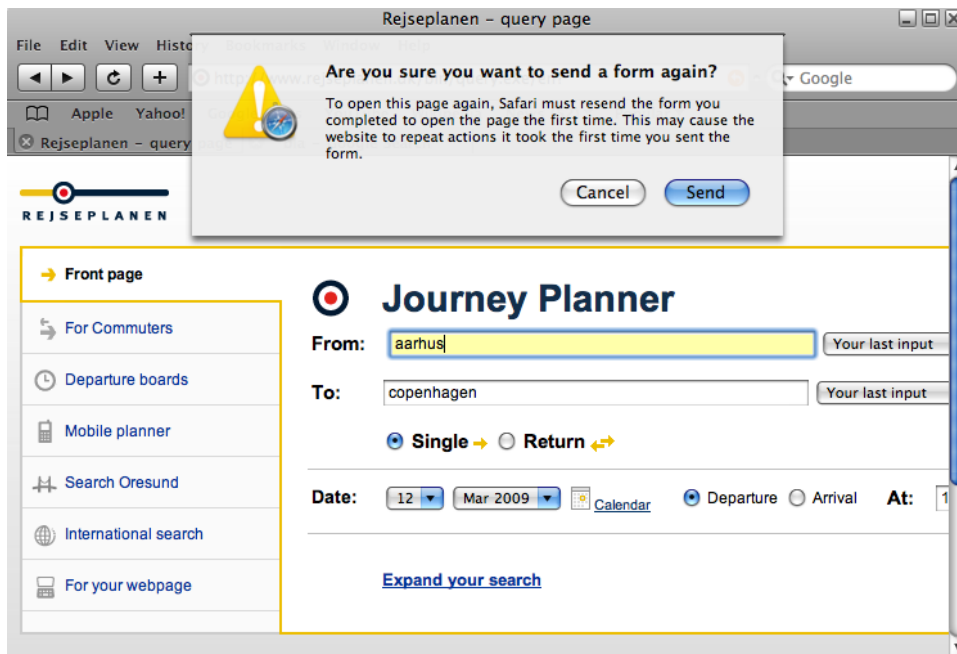


Figure 2.3: Warning message in Safari because of uniform interface violation; [accessed 12-march-2009]

Both reasons are invalid reasons for using POST. The explanation probably lies in `rejseplanen.dk` ignoring the HTTP specification, causing unfortunate side effects for the Web application.

2.4.4 Connectedness

The style of connectedness states that resources in a Web Service should be linked together.

Connectedness anti-pattern

The human Web is a myriad of connected pages in contrast to many Web Services. In Web Services, connectedness is often forgotten [?, p. 96]; disallowing users “to progress through the application by selecting a link or submitting a short data-entry form” breaking REST. The problem is easily illustrated by an example: contrast the JSON representation of weather stations in Listing 2.1 with the representation in Listing 2.2.

```
{
  "next": "/api/weather_stations/?geohash_offset=swgv3sxf003p"}
  "items": [
    {
      "name": "Skagen",
      "lat": 57.723952722299998,
      "lon": 10.5997467041016,
      "timezone": "Europe/Copenhagen",
      "type": "DCAWeatherStation",
      "uri": "/api/weather_stations/dk/skagen/",
      "uri_extern": "http://www.kyst.dk/sw3029.asp",
    }
  ]
}
```

```

    "uri_observations": "/api/weather_stations/dk/skagen/
      observations/",
    "uri_scrape": "http://www.kyst.dk/custom_asp/defaultjs.
      asp?id=3003&targetStation=1100"
  }, (...)
]
}

```

Listing 2.1: *Connected weather stations representation*

The former is a connected resource containing links to other resources. The latter is disconnected. A client of the disconnected Web Service must know all the relevant URIs. Clients are bound to break if the Web Service changes its URIs. A Web Service violating connectedness results in a close coupling between the client and the server.

In a RESTful, i.e., connected Web Service links between resources are available in the representations, such as `uri_scrape` and `observations` in the connected example. The added level of indirection decouples clients from the server since they are now able to use references, instead of direct links, to resources.

```

{
  "page": 1,
  "has_next": true,
  "items": [
    {
      "name": "Skagen",
      "lat": 57.7239527223,
      "lng": 10.5997467041,
      "type": "DCAWeatherStation",
      "timezone": "Europe/Copenhagen"
    }
  ]
}
(...)]]

```

Listing 2.2: *Unconnected weather stations representation*

2.5 Caching

Caching is about one overall thing: reducing the amount of traffic to transfer over the network. Caching is not a requirement in ROA; however, caching is pivotal in reducing the latency of Web Services. This section presents the foundations of HTTP caching.

HTTP caching is a set of mechanisms defined in [?, Sec. 13] that indicate to clients and servers which resource representations are cacheable and when they are stale.

HTTP caching is applied by setting appropriate HTTP headers in the responses from the server. Advanced clients, such as browsers, perform caching based on these headers. The content in the caching headers can be seen as validators; additional client requests trigger the validators. Only if the validators are true the client should update its resource representation.

The caching style of REST (i.e., responses are “implicitly or explicitly labeled as cacheable or noncacheable” [?, p.121]) is satisfied in any HTTP application: if no caching directives is set responses may be cached, however, it is not expected [?, sec.13.4].

There exist three headers to cease explicit control over HTTP caching:³ **Cache-Control**, **ETag**, and **Vary**; each are described in the following.

Cache-Control

Cache-Control is a header field that controls the valid duration of cached responses. The field can specify that a cached resource is valid for a relative time duration. In that time duration all contact with the origin server on additional requests for the same resource is bypassed.

A header response that defines that the *time to live* of a cached representation is 3600 seconds and applies for all clients, looks like the following:

```
Cache-Control: public, max-age=3600
```

A response can also be marked as non-cacheable. The **Cache-Control** field is the concrete means to do the marking by setting value to **no-cache**. There is a subtlety in **no-cache**: clients are still allowed to cache responses, however, before applying the cache they must contact the server to validate that the cached version corresponds to the server's version. Validation puts forth a requirement of identifying different versions of a resource, manifested in the **ETag** field.

ETag

When time to live of a representation is overdue the client might still use the cached version.

ETag is a unique identifier attached to resources in requests and responses. With **ETag** servers validate whether clients' local cached representation is stale by comparing the **ETag** of the local representation against the **ETag** of the resource at the server. The server only sends back a representation upon mismatch of the **ETags**.

Note: Uniqueness of **ETags** are only in terms of identifiers under the same URI; therefore, e.g., the last time of modification is a valid **ETag** value, assuming that no two updates of the same resource can be processed in the same time frame.

The gain of using **ETags** is a reduction of the amount of data which are sent over the network; the total number of requests and responses remain unchanged.

Vary

Since **ETags** are used to compare a client's representation of a resource against the **ETag** of the resource at the server, the cache might serve the wrong representation, but with the correct **ETag**. A common example: given a Web Service that does *content negotiation* on the type of user agent, e.g, a mobile user agent or a desktop user agent, a cached response must only be served to clients of the same type. The **Vary** header field comes to the rescue and indicates to the client the header fields that can vary the representations of the resource. The following line indicates how this is indicated in the HTTP header:

```
Vary: User-Agent
```

2.6 Summary

In this chapter, after having introduced the terms of REST and ROA, we showed what advantages a RESTful Web Service provides. The advantages were illustrated by giving examples of a Web Service that violates the styles of REST, and by

³We ignore the older caching headers from HTTP 1.0: Expires and Last-Modified.

discussing the disadvantages the REST violations cause. Finally, we looked at HTTP caching which is important for increasing the performance of Web Services.

Architectural style	Description	Effect
Client-server	The system is split up into the provider of a service and the consumer of it.	<i>Performance</i> and <i>modifiability</i> benefit since separating the concerns makes the server and client simpler and localize changes.
Stateless	Client state is disallowed on the server therefore all requests must be standalone.	Server <i>performance</i> increases, since the server can free resources immediately after requests are processed. Network <i>performance</i> decreases, since messages must be self contained and therefore demands more space. <i>Availability</i> and <i>modifiability</i> benefit, since requests can be routed to any server containing the service, removing issues of coordination in a clustered environment and losing state on server crashes.
Cache	Results of a request are marked as cacheable or non-cacheable.	When caching is applied, <i>performance</i> increases due to less network communication and less load on the server. However, the complexity of the application increases, and data staleness also becomes an issue.
Uniform interface	Components are restricted to a general interface.	<i>Modifiability</i> benefit since the overall system architecture is simplified and decoupled.
Layered system	The system is split up into hierarchic layers, where layers only use functionality from the layer directly beneath it.	<i>Modifiability</i> increases, as underlying information is hidden and therefore reduces coupling. <i>Performance</i> decreases since communication must pass through several layers.
Code on demand	Clients are enabled to download code from the server and execute it locally.	<i>Performance</i> increases, since code is executed locally avoiding a round-trip to the server. Load is taken off the server and put on the client. <i>Usability</i> increases since the application reduces latency resulting in a more responsive application. <i>Modifiability</i> benefit since features are easily modified and added to the client

Table 2.2: *Architectural Styles of REST (quality attributes are in italics)*

3

Web Service clients

A Web Service client is the consumer of a Web Service. According to Definition 2.1 the consumer can be any HTTP client at the other end of the line, e.g., an Ajax client. In this chapter, we introduce what is probably *the* most popular style for Web Service clients: Ajax. Ajax is presented in the light of REST, leaving important things behind, such as the `XMLHttpRequest` object. In addition, we introduce JSON; a simple data format for interchanging data between servers and clients.

3.1 Ajax: the style for modern Web Service clients

Ajax is a style for a set of Web technologies that narrow the gap between the “richness and responsiveness” of desktop applications and Web applications [?]. Initially, the technologies used in Ajax were incorporated in the name itself. Ajax is short for Asynchronous JavaScript + XML; however, in many newer Web applications the XML has been replaced by JSON.

In Ajax applications, an Ajax engine on the client side extends the HTTP request-response cycle by communicating asynchronously with the server. Traditionally, in server-side Web applications, a client’s request result in the server crafting and returning a unique HTML page to the user. The Ajax style, however, allows for putting the logic of crafting the page on the client¹, and lets the server handle pure *business logic*. The approach has the following advantages:

- the server is relieved of compositioning Web pages;
- the presentation of data at the client is separated from the producer of it; and
- the Web application can be made more rich and responsive.

Most server-side applications that customize pages violate the REST style of statelessness: *application state* is put on the server in order to create customized pages effectively. Maybe the biggest advantage of the Ajax style is that it can circumvent the problem of customization by placing application state where it belongs in terms of REST, in the client.

¹The Ajax style also allows for small page updates of a site crafted at the server; this is, however, not the focus of this section.

The benefits of the Ajax approach over a usual server-side approach are numerous:

- it is possible to create dynamic customized sites with a stateless server;
- static files, such as the Ajax client, are cached at the client after the user's first access to the page, therefore, on additional requests only new business data is loaded from the server. Compare this to a server-side approach where a unique page is created for every user.
- business data is isolated boosting cache hits on intermediate servers for all clients. In addition, an optimal caching strategy for each type of component can be applied. Compare this to a server-side strategy where any change in the site must invalidate the cache.

Thus, an Ajax approach supports REST while at the same making customized Web sites possible. Hereby the Ajax approach inherits one of the main quality attributes of REST: scalability [?, p.116]

The drawbacks of adopting the approach are that 1) the application needs to be implemented using several languages: one for the server-side Web Service, and JavaScript for the client-side; 2) clients need to have JavaScript enabled, and 3) JavaScript is the language of implementation making well-known Web application frameworks – and all their tools such as debugging tools – useless for anything else than creating Web Services.

3.2 JSON

JavaScript Object Notation (JSON) is a data interchange format. JSON is a subset of JavaScript and described as “The Fat-Free Alternative to XML” [?] and indeed since its introduction Web developers and Web Services have adopted JSON as the simpler and cleaner data interchange format with companies such as Yahoo and Google embracing JSON [??]. One of the reasons JSON is widely adopted is due to its simplicity: the syntax can be described with a context-free grammar in just a few lines:

```
<Literal> ::= <Object> | <Array> | string_literal | numeral |  
            boolean_literal | null  
<Object>  ::= {} | { string_literal: <Literal>} |  
            { <Object>, <Object>}  
<Array>   ::= [] | [<Literal>] | [<Literal>, <Literal>]
```

In addition, JSON is efficiently used within JavaScript, and JSON can circumvent the same origin policy. A JSON response is de-serialized into JavaScript objects by using the JavaScript `eval()` function:

```
var data = eval('(' + json_response + ')');
```

The `eval()` function takes a string as parameter and evaluates it as if it were JavaScript code directly written in the document. Since JSON is a subset of JavaScript [?] JSON can be de-serialized with `eval`. However, that approach is subject to *cross-site scripting* if the response contains user generated data. See Appendix A.2 for an example of using a raw `eval` to inject and execute JavaScript code. Instead of using `eval` to de-serialize data we use a JSON parser².

The same origin policy is a set of security restrictions imposed by browsers. They prohibit access to methods and properties of documents served from other

²<http://www.json.org/json2.js>

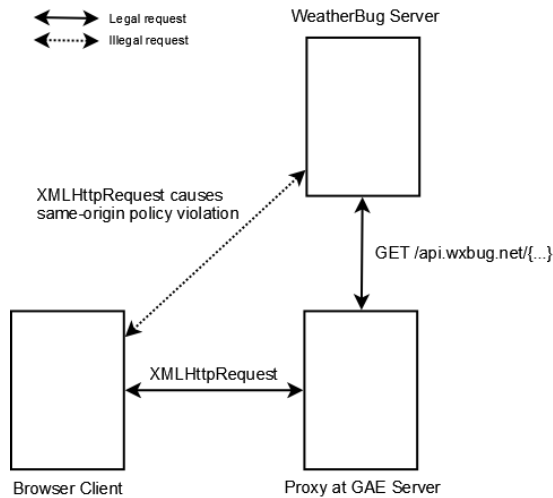


Figure 3.1: *Circumventing the same origin policy with a proxy*

origins compared to the document where it is embedded. Origin is defined by most browsers as the hostname and port number [?]. Therefore, e.g., external Ajax requests and Document Object Model (*DOM*) access to external content in frames are impossible: the content must stem from the same domain to do this.

A workaround to enable access to cross-domain data is setting up a proxy that tunnels all request to a server at another domain through the server hosting the current site. An example where we could have set up a proxy to circumvent the problem is for XML weather data resources from WeatherBug. The resources are not in JavaScript format which makes it impossible to load data directly from WeatherBug into the client. Therefore we could setup a proxy that tunnels all request to the WeatherBug API through the GAE server, as shown in Figure 3.1.

Another reason for the success of JSON is due to an exception in the same origin policy: it does not apply for externally valid JavaScript loaded via the HTML `<script>` tag. Script tags are untouched by the same origin policy, and therefore many Web Services provide data in JSON format. Client access JSON Web Service by dynamically creating `<script>` tags and inserting them into the document; upon insertion the browser fetches the external content into the document. Now how do we know when that `<script>` tag has finished loading? If the Web Service is able to wrap the JSON object in a user-chosen function we have JSONP [?]. With JSONP that user-defined function is called automatically when the `<script>` tag is loaded.

3.3 Summary

In this brief chapter we introduced Ajax from a RESTful perspective, which included an overview of the advantages of the Ajax style where the server is a pure holder of business data.

"The issue is no longer where the information lives - what server, what application, what database, what data center. It's actually now all about putting information to work."

Carly Fiorina (1954-)

4

Google App Engine

When creating Web applications there is an overhead going from development into production with the application. The overhead includes setting up several elements: a Web server, a database, scripts, monitoring, etc. In this chapter, we introduce and describe the Google App Engine (GAE) which removes the overhead. We describe the GAE architecture, and based on the description we present snippets of Web application code and assemble them into a working GAE application.

4.1 What is the Google App Engine?

GAE is a server and development platform, created by Google. GAE lets programmers focus on crafting the application and eliminates many server issues. GAE reduces the time to deploy and automatically scales as the application grows by providing access to the Google infrastructure. The infrastructure consists of a serving infrastructure and several Google services, such as *Google accounts* and the datastore.

The serving infrastructure is encapsulated away from the application programmer. Behind the scenes, the infrastructure adjusts the resources allocated to serve the application, depending on the load on it. Outsourcing the serving infrastructure to Google has the advantage that issues about serving the application are moved to experts at Google. Server issues include things such as setting up the server, securing it, making it scale, etc.

Before the programmer can focus entirely on creating the application, there are still issues regarding both bringing the application online and debugging the application. Google has put an effort into solving these issues by creating a local development environment and scripts to bring the application online instantly.

The advantages come at a cost; a GAE application is limited in numerous ways [??]. In February 2009, the limitations were quite severe:

- only one language was supported, namely Python;
- scheduling tasks were impossible, since HTTP requests were the only means to start a process; and
- long-lasting processing was impossible, since requests could not last longer than 10 seconds.

Google App Engine is an emerging technology; since the first writing, five months ago, the GAE is changed in numerous ways: Google has since added Java support¹, scheduling of cron jobs², however, limited to 20, and increased the request duration to 30 seconds³, and etc..

The limitations are now mitigated, however, essential tasks still *cannot* be done. The limitations mean that many GAE applications demand workarounds; thus, when to use GAE is a subjective matter that depends on the type of the application.

On the other hand, GAE is free and makes it easy to get started with one's application. Therefore, for a prototype Web application that fit almost within the limits of GAE, the advantages outweigh the disadvantages, and make GAE suitable.

4.2 An example for the impatient

Before digging into the GAE architecture we show a 'Hello, world' example for GAE. The interesting point about the example is how little code is needed to run the example on the server.

The 'Hello, world' example consists of only two files, given in Listing 4.1 and Listing 4.2.

- `app.yaml`, a configuration file containing handlers. Handlers are a list of *URI* patterns with descriptions of how requests to URIs are handled. In this case the server will execute the script `main.py` for all URIs.
- `main.py`, a Python *CGI* script that output a string containing a content header and the content, which the server sends back to the client.

```
application: helloworld
version: 1
runtime: python
api_version: 1

handlers:
- url: .*
  script: main.py
```

Listing 4.1: *app.yaml*

```
print 'Content-Type: text/html'
print                                     # end of headers
print '''
<html>
  <head>
    <title>Hello, world</title>
  </head>
  <body>
    <h1>Hello, world</h1>
  </body>
</html>
'''
```

Listing 4.2: *main.py*

¹<http://code.google.com/p/googleappengine/issues/detail?id=1>

²http://groups.google.com/group/google-appengine/browse_thread/thread/552e9ab4a97abc48?hl=en

³<http://code.google.com/p/googleappengine/issues/detail?id=6>

The example is run by starting the development server, with the directory of the two files, as argument. Now accessing `http://localhost:8080` with a browser outputs `Hello, world`.

4.3 Architecture

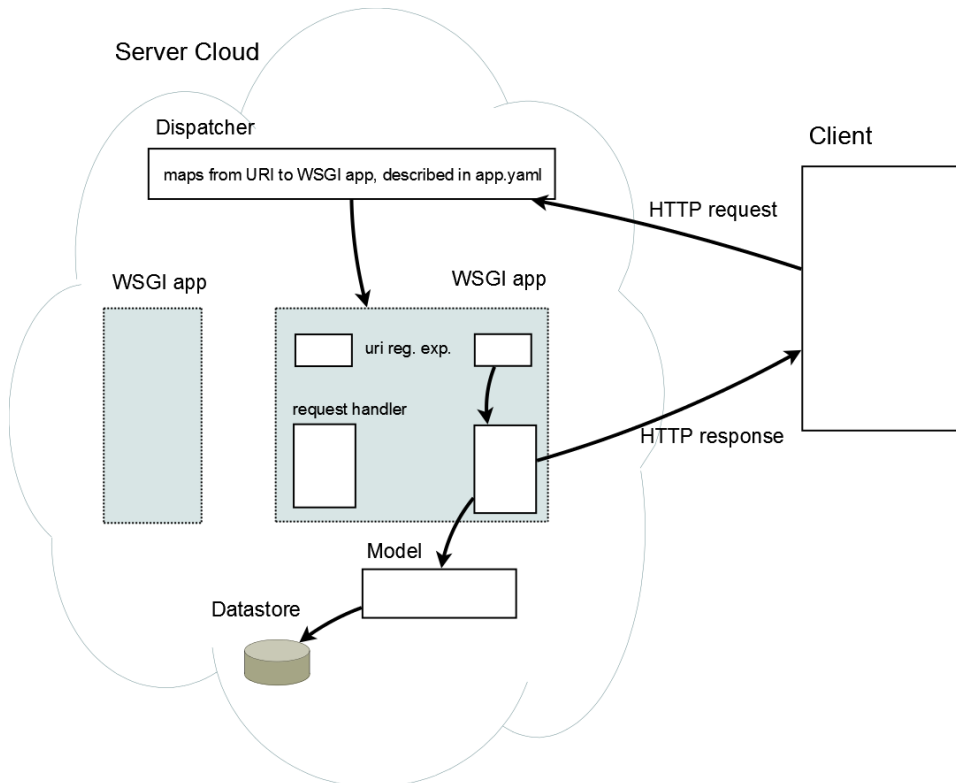


Figure 4.1: *GAE architecture*

The basic architecture of GAE is illustrated in Figure 4.1. Starting from the client, on the right, incoming URI requests are routed to a corresponding *WSGI* application – the blue squares in the figure – somewhere in the server *cloud*. Inside the *WSGI* applications the requests are routed again to a request handler, which potentially interacts with the data models, to get relevant data. At last the server returns a string to the client.

The GAE comes with its own Web application framework called *webapp*. The Web server supports any *CGI*-compliant Python application, like *webapp* or *Django*. However, since *Django* is tied closely together with relational databases, the *Django* framework cannot directly be used in GAE. There are different helpers to work around the database impedance mismatch [??]. For the application in this thesis we are using the *app-engine-patch* that adds *Django* support to GAE.

Crafting Web pages in *Django* follows the *request-response pattern* from HTTP. A *Django* application contains request handlers which are Python functions. When a request arrives in the *Django* application, it is routed to a request handler and wrapped in a request object which is given as argument to the handler. Routing is setup explicitly by regular expression.

Each request handler must return a response object. The handler creates the response string either explicitly, within the handler, or implicitly with templates.

Templates are used to separate the presentation logic from the request handlers. Using templates, the architecture can be seen as consisting of three layers, fitting into the Model-View-Controller style (MVC), shown in Figure 4.2.

The style consists of three architectural components:

- the Model that contains the data and the persistence logic;
- the View that takes the data from the controller and produces a textual response (HTML, JavaScript, etc.); and
- the Controller that receives the requests and connects views and models.




Figure 4.2: *Model-View-Controller architectural pattern*

The depicted pattern is a variant of the original Model-View-Controller pattern [?]. Instead of the model pushing updates to the view, the controller pulls the data from the model, and passes model data to the view. Therefore there are no connection between the model and the view, like in the original MVC pattern.

In the following sections, the GAE and Django is described in terms of the MVC pattern.

4.3.1 Model

The models of the application declare the format of, and manipulate the data of the application. The models are usually stored in a relational database. The “database” in GAE is special; it is not a normal relational database, but a database based on Google’s BigTable technology [?], called datastore. The best way to think of the datastore is as a database of objects.

GAE comes with an API for datastore interaction, called Datastore API. The API provides functionality to declare datastore entities, and declare their properties, directly in Python, based on the *active record* pattern. Creating a data model with the Datastore API is easy. The datastore data models are created reflectively, on their first import, by the GAE from Python classes that inherits from the `db.Model` class. This means a data model is a Python class. The content of the data model is declared by defining class attributes of the type `Property` in the model class. An example model that represents a surf spot is given in Listing 4.3.

```
from google.appengine.ext import db
```



```

class Spot(db.Model):
    name = db.StringProperty(required=True)
    point = db.GeoPtProperty(required=True)

```

Listing 4.3: *Spot model*

The spot model uses the built-in property-classes to declare the properties of the spot data model. The point property uses a richer semantic type `GeoPtProperty` used to store geographical points consisting of a latitude and longitude. The API includes many such richer types, such as `UserProperty` and `EmailProperty`.

References between tables in the datastore is defined by a reference property variable in the referring class. The referenced entity is accessed as if it were just an object reference. The GAE automatically creates indexes for the references in order to increase the performance of reference look-ups.

Queries

Queries can be defined in two ways: 1) it is possible to define queries in a SQL-like language called GQL, and 2) with a query object where the query is built with methods. An example of the latter is the `all` method in Listing 4.6. GQL and query object are equally expressive, however, comparing GQL with SQL is easier with GQL. The GQL syntax is given in Listing 4.4.

```

<gql> ::= SELECT [* | __key__] FROM <kind>
        [WHERE <condition> [AND <condition> ...]]
        [ORDER BY <property> [ASC | DESC]
        [, <property> [ASC | DESC] ...]]
        [LIMIT [<offset>,<count>]]
        [OFFSET <offset>]

<condition> ::= <property> {< | <= | > | >= | = | != } <value>
<condition> ::= <property> IN <list>
<condition> ::= ANCESTOR IS <entity or key>

```

Listing 4.4: *GQL grammar*

GQL is only concerned with fetching data. A GQL query always restricts by kind and zero or more property filters are applied. GQL is limited in comparison to SQL; notice the lack of JOINS, ORs, projections, and aggregation functions in the grammar. An important restriction is that when inequality filters are applied, they can only be on the same single property, e.g., a query with inequality filters on both a latitude and longitude is invalid.

GQL is subject to the design choice of scalability; to achieve scalability Google removed the aforementioned functionality. The result is queries which

1. avoid sorting records in memory [?];
2. avoid expensive JOINS; and
3. benefits from distributed servers;

This lack of expressive power is not a significant problem, however, it demands a paradigm shift, because a relational SQL approach cannot be used; e.g., the aggregate count function, unavailable in GQL, can be implemented by maintaining an integer in the datastore of the count. Of course this could also be implemented by fetching all objects into a list and doing a count of the objects in the list, but

this is exactly what is not meant to be done, because it demands a lot of processing; in addition, the data set of queries is restricted to 1000 entities, making the result invalid when that upper bound is hit.

4.3.2 View

The view layer in Django consist of Django templates. A Django template consists of

- regular text (HTML, JavaScript, etc.);
- programming constructs for sticking in values in the template, always resident inside `{{ }}`; and
- control flow constructs called block tags `[?]`, always resident inside `{% %}`.

When a template is rendered a dictionary is passed to it, the values are stuck into the template, properties are resolved, and a text document is returned.

A template that will render all spots (given to it in a dictionary) in JSON format is given in Listing 4.5. The template takes advantage of the `for` block tag to iterate through all the spots in a given dictionary. For every spot the template outputs attribute name value pairs. The `if` block tag uses `forloop.last` to check whether the current iteration element is not the last, and outputs a comma if so; without the check it would return invalid JSON.

```
[
  {% for spot in spots %}
  {
    "id": "{{ spot.key.id }}",
    "name": "{{ spot.name }}",
    "lat": {{ spot.point.lat }},
    "lng": {{ spot.point.lon }}
  }
  {% if not forloop.last%},{% endif %}
  {% endfor %}
]
```

Listing 4.5: *Django template*

4.3.3 Controller

Controllers mediate between views and models. A controller consists in a dispatcher and a request handler. After the dispatcher has passed a request to the relevant handler, the handler

1. gets relevant data from the models;
2. parses the data to the view; and
3. returns the rendered view in a response.

GAE and Django handles much of the controller logic, GAE maps URIs to an application, and Django maps URIs to request handlers by regular expressions. The request handlers, together with the regular expression mappings from URI to request handler, is the part of the controller that the developer is responsible for. An example of a request handler is given in Listing 4.6. The request handler gets all the spots from the datastore, by using the query object interface of the Spot model class. The data is parsed on as a dictionary to the template which is rendered, and the result returned to the client in a response object.

```

"""Returns spots in JSON format"""
def index(request):
    data = {'spots': models.Spot.all()}
    response = shortcuts.render_to_response('spots/spots.js',
        data)
    return response

```

Listing 4.6: *Spot service handler*

4.4 Putting it together

Given the model snippets, the controller snippets and the view snippets in the last sections, we almost have a complete GAE Web application⁴. What remain are to

1. register the application with an URI; and
2. register the event handler with an URI.

In Listing 4.7 the URI to request handler file is shown.

```

urlpatterns = patterns('',
    (r'^spots/$', index),
)

```

Listing 4.7: *URI to request handler*

Each application in the GAE has an application descriptor named `app.yaml`. An example of such a file is given in Listing 4.8. The file registers the main module with every URI request to the application. As shown in Figure 4.1, several applications can exist; `app.yaml` is the place where we route to the appropriate, specified under `handlers`. In addition, the application name and version are stated here, which are used when uploading the application to the server.

```

application: welovewind
version: 1
runtime: python
api_version: 1

handlers:
- url: /*
  script: common/appenginepatch/main.py

```

Listing 4.8: *app.yaml configuration file*

After registering the application at `http://appengine.google.com`, the application is put into production by running a single command; type

```
manage.py update
```

in the directory of the application, and the application is online, see Figure 4.3.

4.5 Further GAE prerequisites

This section presents GAE prerequisites for later sections. The reader can skip the section for now and revisit it when the prerequisites are used in later sections.

⁴We assume that `app-engine-patch` is downloaded and setup

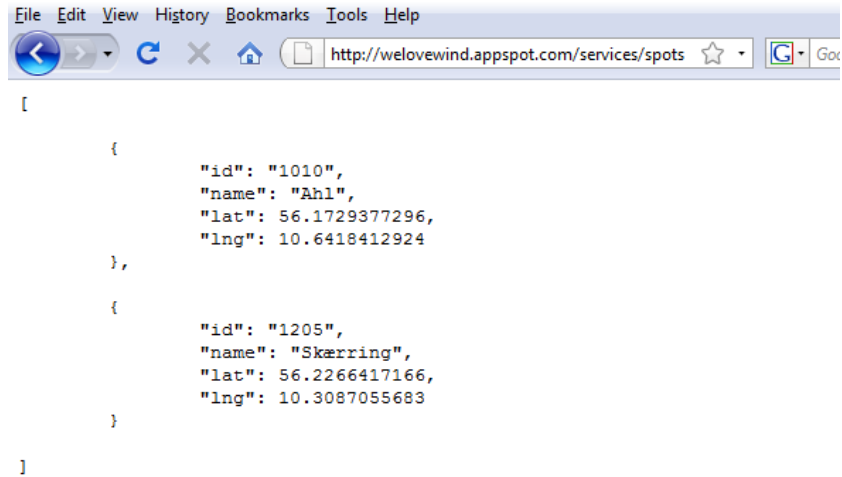


Figure 4.3: *Spots Service*

4.5.1 Unique Properties

In the datastore there is no built-in option that enforces unique properties on the models. Per design there is only one unique property namely the key of the model. This leaves two possibilities for implementing unique constraints:

1. simulate unique properties in Python⁵; or
2. incorporate the unique property in the key name.

Simulating unique properties results in a fairly clean solution, however, it has the huge flaw that there is no possible way of ensuring that the properties are indeed unique since they are simulated. This is a problem if the solution is applied in a concurrent environment and uniqueness guarantees are essential. The second solution is obviously a solution to a recurring problem and we describe this in the following.

As mentioned, the datastore ensures exactly one unique property per instance: the key name of the instance. The Datastore API provides a transactional method on models which either gets or inserts an instance based on the key name. We use the method to ensure uniqueness by incorporating the unique constraints in the keyname.

Listing 4.9 shows the implementation of unique properties for the **Forecast** weather data class. It is based on the concept that a forecast point only should have one forecast for each time delta (the difference between calculation time and forecast time). The delta is incorporated into the key name to ensure the uniqueness.

The class method `key_name` returns the unique key name to use for the entity; the key name is the URI of the forecast resource. The method is used in `update_or_insert`, which as the first step calls it to get the key name from the given arguments. The key name is used as input to `get_or_insert` that gets or creates a forecast with the specified time delta and forecast point, now incorporated in the key name.

Instead of only getting an entity when the entity exist the authors implementation updates the entity both when entities are updated and created. To find out

⁵An example of a Python programmatic implementation of unique properties is available at: <http://appengine-cookbook.appspot.com/recipe/get-or-insert-entity-by-unique-properties/>

whether the entity exists a `is_new` key word is passed to the entity model. Since the key word is used only when the entity is unsaved the property is only set to `True` in this case.

```
class Forecast(AbstractWeatherData):

    def __init__(self, is_new = False, **kwds):
        self.is_new = is_new
        super(Model, self).__init__(**kwds)

    (...)

    @classmethod
    def key_name(cls, calculation_time, forecast_time,
                forecast_point):
        time_delta = wlttime.timeDifferenceInHours(calculation_time
            , forecast_time)
        owner = forecast_point.key().name()
        return '%s/forecasts/time_delta/%s' % (owner, time_delta)

    @classmethod
    def update_or_insert(cls, calculation_time, **kwds):
        '''Update or insert forecast weather data.

        Args:
            calculation_time: forecast calc. time (datetime).
            **kwds: Key word arguments to pass to the instance and
                update / init the value of.

        Returns:
            Updated / new instance.
        '''
        key_name = Forecast.key_name(calculation_time,
            kwds['time'], kwds['forecast_point'])
        entity = super(Forecast, cls).get_or_insert(key_name,
            is_new=True, **kwds)
        if not entity.is_new:
            for prop in kwds.keys():
                setattr(entity, prop, kwds[prop])
            entity.put()
        return entity
```

Listing 4.9: *Forecast with unique properties*

4.5.2 Pagination

Pagination is the process of splitting up content into several resources. The GAE limits the number of returned results of a query and restricts the duration of a request. Because of the restrictions pagination is a must have.

Django comes with a built-in paginator class that handles pagination. The paginator takes a complete list of objects and provides methods to get specific pages. However, since the class takes a list of all the objects that should be paginated it causes a slowdown as more and more objects are in the list; in addition, since the datastore is limited to a result set of 1000, the paginator is not able to access list objects beyond the limit.

The author's initial (un-scalable) solution was a slightly changed version of the paginator that only retrieved objects that belonged to the current page plus one

more object. The extra object is used to evaluate whether the current page has a next page. The new paginator instead of taking a list of objects, takes a query as argument. The paginator then puts the relevant filters on the query object to retrieve only the objects needed for the page requested plus the additional entity to evaluate whether the page has an additional page. The modified paginator fetches objects with the following query:

```
objects = query.fetch(limit=per_page + 1,
                      offset=(self.page_number-1) * per_page)
```

This query is un-scalable. Specifying an offset merely chops off entities in the result set, however, the entities are still fetched from the datastore. That the offset argument is provided is just confusing.

The only way to do scalable pagination is paging based on an indexed property and limiting the number of entities. Following pages continues from the last value (in the previous page) in terms of the ordering using an in-equality filter on that property. We put it to use in Section 6.1.

4.5.3 Timezones

All times in the datastore are stored in Coordinated Universal Time (UTC). When displaying times to the user the application should convert the UTC times to the local times. The easiest way to convert between UTC and the timezone of the user is by using JavaScript. However, JavaScript is not always available and then a conversion on the server is needed.

Converting between timezones on the server is complicated since Python does not come with logic to handle the different timezones. It merely provides an abstract interface. What is needed is a complete database of the politically decided timezone offsets and daylight saving times. The tz database, or Olson Timezone database, is a continuously maintained database of timezone info. In Python the `pytz`⁶ module wraps the tz database and provides access to the Olson database. The `pytz` module is not included in the Python Standard Library.

`pytz` includes over 500 hundred files. The GAE has a restriction of 1000 uploaded files; the limit is quickly crossed when including `pytz`. All files except four are binary timezone definitions belonging to the tz database. A blog entry⁷ describes how to circumvent the problem by zipping all the timezone files, and include logic to unzip them during runtime.

The GAE provides a distributed memory cache – memcache – which is a distributed hash table. Unzipped timezones are obvious subjects for caching in the distributed memory. Saving timezones in the memcache avoids the overhead in unzipping data from the tz database repeatedly.

Listing 4.10 shows the single method in `pytz` we overwrite to incorporate the two changes: the changed method reads timezone definitions from the zipped archive of timezones `zoneinfo.zip`, and caches the result in the memcache. The code is based on an example from the Google App Engine Cookbook⁸.

```
def open_resource(name):
    """Open a resource from the zoneinfo subdir for reading.

    """
    import zipfile
    from cStringIO import StringIO
```

⁶<http://pypi.python.org/pypi/pytz/>

⁷<http://takashi-matsuo.blogspot.com/2008/07/using-newest-zipped-pytz-on-gae.html>

⁸<http://appengine-cookbook.appspot.com/recipe/caching-pytz-helper/>

```

name_parts = name.lstrip('/').split('/')
for part in name_parts:
    if part == os.path.pardir or os.path.sep in part:
        raise ValueError('Bad path segment: %r' % part)
cache_key = "tzinfo/%s/%s" % (OLSON_VERSION, name)
zoneinfo_contents = memcache.get(cache_key)
if zoneinfo_contents is None:
    zoneinfo = zipfile.ZipFile(os.path.join(os.path.dirname(
        __file__),
        'zoneinfo.zip'))
    zoneinfo_contents = zoneinfo.read(os.path.join(
        'zoneinfo', *name_parts).replace('\\', '/'))
    memcache.add(cache_key, zoneinfo_contents)
return StringIO(zoneinfo_contents)
(...)

# We don't need this, since we are including all tzinfo files
#all_timezones = [
#    tz for tz in all_timezones if resource_exists(tz)]

```

Listing 4.10: *Altered pytz*

`pytz` iterates through all timezone files when it is initialized, checking whether every timezone file exist, we remove this as we have included the whole tz database; avoiding the latency of unzipping over 500 files everytime the server is started. The removed code is also shown in Listing 4.10.

With `pytz` in place we are armed to convert between timezones. Listing 4.11 shows how to convert from the times stored in the datastore, UTC, to a local timezone. Datetimes stored in the datastore does not have timezone information in them; before converting to local time it must be added. After the datetime is enriched with timezone information it is converted to local time with `astimezone(tz)`: the function adjusts a datetime according to the given timezone.

```

>>> from pytz import timezone
>>> from datetime import datetime
>>> def local_time(utc_dt, tz_string):
...     utc = timezone('UTC')
...     utc_dt = utc_dt.replace(tzinfo=utc)
...     tz = timezone(tz_string)
...     return utc_dt.astimezone(tz)
...
>>> utc = datetime.utcnow()
>>> utc
datetime.datetime(2009, 5, 5, 6, 30, 53, 793000)
>>> lt = local_time(utc, 'Europe/Copenhagen')
>>> lt
datetime.datetime(2009, 5, 5, 8, 30, 53, 793000, tzinfo=<
    DstTzInfo 'Europe/Copenhagen' CEST+2:00:00 DST>)

```

Listing 4.11: *Converting datastore time to localtime*

`pytz` is put into action in Section 6.1, where we also describe how to get the relevant timezone string, e.g., `'Europe/Copenhagen'`.

4.6 Summary

In this chapter we have presented the GAE and the Web application framework, Django. GAE and Django is used in the rest of the dissertation as the basis upon which we shall build several elements that will support the aspiration of assisting wind- and kitesurfers. The chapter also presented GAE solutions to the common problems of pagination, unique properties of data models, and conversion of times to another timezone.

Part II

The Web Services

5

Resources and their Representations

Based on the scenarios in Section 1.1 we derive the resources for our Web Service which is the corner stone of the application. The nouns of the scenarios define the resources in the Web Service: spots, weather stations, and forecasts. All representations of the resources are in the JSON format.

In the sections about resources we summarize the resources in what we call a resource view. Together the views describe the resource view of the software architecture from the resource viewpoint. The concept of views and viewpoints is presented in [?].

Definition 5.1 *The resource viewpoint on the software architecture is concerned with what resources the system contains and what parts of the uniform interface they expose.*

5.1 Spots

A spot is a geographical point where people wind- and kitesurf. A point is naturally encoded by latitude and longitude. All spots have a name and information about its wind directions suitable for surfing; we denote the property of a spot being apt for surfing (suitable wind in a suitable direction) as surfability. The directions of surfability are described in a wind diagram that captures in which directions – North, North-East, East, etc. – the spot is surfable.

Listing 5.1 shows an example of the representation of the spots resource. The resource is a JSON object that among others contains the following fields:

- **next** is an indicator for *pagination*.
- **items** a list containing individual spots.
- **uri** the URI of the concrete spot.
- **forecast_point** the nearest point where forecasts for the spot are calculated for.

The spots resource is the first of three list resources. The list resources have the same structure: a **next** field which is the URI of the next page or the empty string

if the page is the last and an items field containing the concrete items. The uniform structure will come in handy when retrieving resources in the clients.

Spots are connected to the nearest forecast point (soon presented) and the individual spots, hereby satisfying connectedness.

```
{
  "next": "",
  "items": [
    {
      "name": "Ahl",
      "country_code": "dk",
      "lat": 56.1729377296,
      "lng": 10.6418412924,
      "wind_diagram": {
        "S": "YES",
        "SW": "YES",
        "W": "YES",
        "NW": "YES"
      },
      "uri": "/spots/ahl/",
      "forecast_point": "/forecasts/56.0,10.5/"
    },
    (...)
  ]
}
```

Listing 5.1: *Spots representation*

In addition to paging, for performance reasons the list of spots must be subject to different filters reducing the total number of spots returned; there must exist filters to reduce the list to

1. only spots in the proximity of a geographical point; and
2. only spots in a certain country.

The preceding representation is served to the client. We now turn to the representations sent from the client.

Spots are created by any user in the system. The application receives concrete spot resource representations to either create or update spot resources. The representation of spots accepted from the client at the server is different than those sent to the client. The representations accepted from the client are in the format that browsers submit data in: the official name is **application/x-www-form-urlencoded** [?, sec.17], we denote it uri-encoding from now on. The difference in format is because we rely on Django forms on the server side. Django forms gives much for free, including validation of fields, error messages to return to the user, and acceptance of uri-encoded data directly. A uri-encoded spot representation looks like the following:

name=Ahl&lat=56.172&lon=10.641&country_code=dk&S=YES

An overview from the resource viewpoint of the resources discussed in this section is given in Table 5.1. The URIs support the naming conventions stated in [?, p.117-118]; this means,

- / indicates a **parent/child** resource hierarchy;

Resource URI	Action
/spots/	GET all spots at page 1
/spots/{country_code}/	GET all spots in a certain country
/spots/	POST a new spot
/spots/{spot name}/	PUT a new representation of an existing spot
/spots/{spot name}/	DELETE an existing spot

Table 5.1: *Resource view of spot resources*

- , or ; indicate no hierarchy. The former indicates that the order is significant, e.g., the latitude and longitude pair 56.0,10.5, whereas the latter indicate that the order is insignificant. Inputs that conceptually are arguments to an algorithm are indicated with query parameters.

5.2 Weather Stations

A weather station is a place where observations are created. Observations consist of a time, a wind velocity measurement, and a temperature measurement. The application currently has three resources for live weather data DCA, DMI, and WB.

DCA and DMI

Live weather data – wind speed and gust, temperature, and direction measurements – are available at DCA (of the west coast in Jutland) and DMI (most regions in Denmark).

The latest observation is interesting and is exposed as a resource. In addition, it is relevant to know whether the wind picks up or slacks down based on the last observations. Therefore a list of weather observations belonging to a specific station is a resource. The latest observations is incorporated into the weather stations resource; the reason is that the last observation is needed in most cases, and this is not the case for the rest of the observations.

Listing 5.2 shows an example of the representation of weather stations. The resource is a JSON object that among others contains the following fields:

- In the JSON list, `items`, the `uri_extern` field specifies the URI of the weather station at the external location: the external page to link to. The page to scrape is identified by `uri_scrape`.
- The last two fields in the list are links to the weather station and its observations respectively, thereby satisfying the constraint of connectedness.
- The `timezone` field is used to convert the UTC observation time stored in the database to the time where the weather station is located when JavaScript is not available.

```
{
  "next": "/api/weather_stations/?geohash_offset=swgv3sxf003p"}
  "items": [
    {
      "name": "Skagen",
```

```

    "lat": 57.723952722299998,
    "lon": 10.5997467041016,
    "timezone": "Europe/Copenhagen",
    "type": "DCAWeatherStation",
    "latest": {
      "direction": 242.90000000000001,
      "icon": "/images/wind_arrows/SW_4to6.png",
      "speed": 5.9000000000000004,
      "temp": 13.0,
      "time": "2009-06-19T12:00:00"}
  },
  "uri": "/api/weather_stations/dk/skagen/",
  "uri_extern": "http://www.kyst.dk/sw3029.asp",
  "uri_observations": "/api/weather_stations/dk/skagen/
    observations/",
  "uri_scrape": "http://www.kyst.dk/custom_asp/defaultjs.
    asp?id=3003&targetStation=1100"
}, (...)
]
}

```

Listing 5.2: *Weather stations representation*

The representation for observations for a specific resource is a list of objects containing weather data; an example is shown in Listing 5.3. The `time` field indicates the time of the observation in UTC time, the `speed` and `gust` field indicates the speed in m/s, the `direction` field is the direction in degrees where the wind is blowing from, and the `icon` field is a link to a representative direction and speed icon.

```

{
  "owner": "/weather_stations/dk/skagen/",
  "page": 1,
  "has_next": true,
  "items": [
    {
      "direction": 242.9,
      "icon": "/images/wind_arrows/SW_4to6.png",
      "speed": 5.9,
      "temp": 13.0,
      "time": "2009-06-19T12:00:00"
    },
    (...)
  ]
}

```

Listing 5.3: *Observations representation*

WeatherBug

DCA and DMI produce weather observations only for Denmark; additional resources are needed.

WeatherBug provides an API to access weather information worldwide. We use the following functions from the API:

- retrieve an XML list of stations in the proximity of a latitude and longitude; and,

- retrieve an RSS feed of live weather data for a specific weather station.

Representations of the weather resources are obviously already created. The existing representations, at first sight, open the option to completely bypass the application server and let JavaScript retrieve the relevant weather data directly from WeatherBug. The representations, however, are not in JavaScript format which makes it impossible to load data directly into browser clients because of the same origin policy described in Section 3.2.

Another option is storing the data on the GAE server and update the content as it becomes stale. This approach is the only option that serves data fast enough and is therefore chosen. Following the approach weather stations are subject to frequent updates of observations. Updates are uploaded from a Web Scraper located at an external machine. We expose the POST method of the observations resource to post observations to since the concrete URI of the posted resource is unknown. The observation representation accepted from the client is in JSON format, and is just a single item of the observations list shown in Listing 5.3.

We use the same representation for WeatherBug stations as for DCA and DMI weather stations. An overview of the resources is given in Table 5.2.

Resource URI	Action
/weather_stations/	GET all weather stations
/weather_stations/{country_code}/	GET all weather stations in a certain country
/weather_stations/{id}/observations/	GET all weather observations
/weather_stations/{id}/observations/	POST a new weather observation

Table 5.2: *Resource view of weather station resources*

5.3 Forecast Points

A forecast point is a specific point on the Earth where forecasts are calculated for. A forecast is a prediction of the weather at some specific time for a specific forecast point.

A natural source for forecasts is DMI; however, their forecast data is unavailable and they have no interest in joining in on a collaborative project. The National Oceanic Atmospheric Administration¹ (NOAA) in the US, however, by law puts weather data into the public domain². This means use of their data is allowed in almost any context, and that we can be certain that the service continues.

NOAA produces many highly detailed forecast for the US, and in addition global forecasts using the Global Forecast System³ (GFS). The forecasts are made four times daily at 00, 06, 12 and at 18 o'clock UTC. At these four times detailed forecasts are calculated for the next 48 hours with a three hour interval between them; actually there are forecast with a longer prediction in the future but these are of decreasing detail. The most detailed GFS forecasts produced at NOAA are produced at a degree of 0.5°.

A note on accuracy: the earth has latitudes in the interval $[-90^\circ; 90^\circ]$. This gives a total of 360 points along *meridians* that are subject for forecasts. Longitudes are

¹<http://www.noaa.org/>

²<http://www.weather.gov/disclaimer.php>

³Info about the GFS model <http://www.emc.ncep.noaa.gov/modelinfo/index.html>

values in the interval $[-180^\circ; 180^\circ]$. This yields 720 points. Looking at a map we can see that in this resolution, e.g., Aarhus and Skanderborg probably are associated with the same grid location.

Detailed NOAA GFS forecasts are located at <http://nomad3.ncep.noaa.gov/pub/gfs/rotating-0.5/>. In the directory there are three important types of files:

1. Forecasts in GRIB2 format named `gfs.t{TT}z.pgrb2f{XX}`;
2. inventories for forecasts named `gfs.t{TT}z.pgrb2f{XX}.inv`; and
3. descriptor files for a specific run named `gfs_t{TT}z.ctl`;

In the above `TT` is the time of the calculation and `XX` is a delta in hours to add to the time of calculation to get the time of the prediction. In the descriptor file all the fields in the forecasts are defined. For now we are interested in the wind forecast which is given by a vector split up into a `u` and a `v` component of the wind 10 meters above the ground. In addition, we are interested in the temperature forecast near the surface given by a single field where the value is reported using the Kelvin scale. The value of the fields is retrieved by, e.g., `wgrib2` introduced in Chapter 8.

```
{
  "next": "",
  "items": [{
    "lat": 36.5,
    "calculation_time": "2009-07-01T06:00:00",
    "lon": 29.0,
    "uri": "/api/forecast_points/36.5,29.0/",
    "forecasts": [{
      "direction": 252.53999999999999,
      "icon": "/images/wind_arrows/W_6to8.png",
      "speed": 6.5300000000000002,
      "temp": 27.050000000000001,
      "time": "2009-07-01T12:00:00"},
      (...)
    ]
  },
  (...)
]
```

Listing 5.4: *Forecast points representation*

Listing 5.4 shows the representation of forecasts for a specific point. It contains the time of forecast calculation, and the forecasts. An overview of the resources is given in Table 5.3. Updates of a point's forecasts are supported by exposing the

Resource URI	Action
<code>/forecast_points/</code>	GET all points
<code>/forecast_points/{lat},{lon}/</code>	GET a forecast point representation
<code>/forecast_points/{lat},{lon}/</code>	PUT a new forecast point representation

Table 5.3: *Resource view of the forecast points resources*

PUT method of the uniform interface. The representation accepted is in the same JSON format.

In comparison to weather station resources, forecast point resources have no weather data resources belonging to it. Forecast point resources consist of both

forecast point information and all the forecasts. The reason for deviating is that all weather data of the forecast point belong to a specific calculation time of the forecast point making it more natural to think of all the content as a single resource.

5.4 Sorting out the resource requirements

In the last section we pointed out the resources in the Web Service, not how to process the content in them. Filtering points on a proximity basis is non-trivial. This section reviews our solution.

The application will eventually contain a huge data set: spots, weather stations, forecast points, etc. Returning the whole data set is infeasible in terms of latency at the client. In addition, we have a *location-based service* requirement. To sum up, filtering by location reduces the data set, and increases the relevancy of the data set. Therefore, the application needs a filter on points which returns a subset of points in the proximity of a certain place.

Filtering on proximity sounds like a rather easy problem but it is not. To limit the result set, in a usual database several inequality filters are applied to the latitude and longitude values creating a bounding box. As mentioned in Section 4.3.1, the datastore can only apply inequality filters on a single property; a bounding box demands two: an index on latitude and an index on longitude. Listing 5.5 shows a dysfunctional GAE query that one could think would solve the problem.

```
>>> from wlv.models import Spot
>>> from google.appengine.ext import db
>>>
>>> point_1 = db.GeoPt(lat=50, lon=-10)
>>> point_2 = db.GeoPt(lat=60, lon=0)
>>> q = db.Query(Spot) \
...     .filter('point >', point_1) \
...     .filter('point <', point_2)
>>> for s in q:
...     print 'spot: %s \npoint: (%s)' % (s.key().name(), s.
...         point)
...
spot: ahl
point: (56.1729377296,10.6418412924)
```

Listing 5.5: *Dysfunctional GAE bounding box query*

The filters in the query apparently set up a bounding box, but, because of the inequality filter restriction the query ignores the longitude boundary given by the `GeoPt` property, and only filters on the latitude. The problem results in GAE returning *all* spots that lie within the given latitudes, in this case only one: Ahl.

We present two approaches to the problem, where the first is an obvious solution, but it does not scale as the data set gets large. The other, is a scalable solution based on geohash.

5.4.1 Calculating proximity points with haversine

Finding proximity points is possible by first calculating the distance between a reference point and all other points; afterwards, filtering out points over a certain threshold distance to the reference point.

In the following we assume the Earth is a sphere ignoring the fact that best approximation is an oblate spheroid⁴.

⁴<http://en.wikipedia.org/wiki/Earth>

A great circle is a circle on a sphere which cuts the sphere into two equally sized halves, i.e, the circle has radius equal to the radius of the sphere and the same center as the sphere. An area on the sphere bounded by three arcs of great circles that do not all have a common point is called a spherical triangle. The (shortest) distance between two points on a great circle is the shortest possible distance between the two points, since the arc on the great circle is the arc with the smallest deviation from a straight line. Looking at Figure 5.1, which shows a spherical triangle, the arc c along the great circle is the shortest distance between A and B . In a spherical




Figure 5.1: *Illustration of a spherical triangle*

triangle in a unit sphere the law of cosines relates the sides and angles of the triangle, in the following way:

Theorem 5.1 *Spherical law of cosines* [?, p.130]

$$\cos(c) = \cos(a)\cos(b) + \sin(a)\sin(b)\cos(C) \quad (5.1)$$

Isolating c in the formula gives the distance between the points A and B . However, when c (the distance) is small rounding errors are stated to be significant [???], therefore, instead the haversine formula is used [?]. The formula is based on the haversin function, defined as:

$$\text{haversin}(\theta) = \sin^2\left(\frac{\theta}{2}\right) \quad (5.2)$$

The following identity between \cos and haversin exists:

$$\begin{aligned} \cos(2\theta) &= 1 - 2\sin^2(\theta) \\ &\quad \text{(double-angle formula [?, p.A23])} \\ \cos(\theta) &= 1 - 2\sin^2\left(\frac{\theta}{2}\right) = 1 - 2\text{haversin}(\theta) \end{aligned} \quad (5.3)$$

We now derive the law of haversines from the spherical law of cosines. Substituting

\cos with haversin using (5.3) in the law of cosines, we get:

$$\begin{aligned}
1 - 2\text{haversin}(c) &= \cos(a)\cos(b) + \sin(a)\sin(b)(1 - 2\text{haversin}(C)) \\
1 - 2\text{haversin}(c) &= \cos(a - b) - 2\sin(a)\sin(b)\text{haversin}(C) \\
&\quad \text{(Using the subtraction formulas [?, p.A23])} \\
1 - 2\text{haversin}(c) &= 1 - 2\text{haversin}(a - b) - 2\sin(a)\sin(b)\text{haversin}(C) \\
&\quad \text{(Substituting } \cos \text{ with } \text{haversin}) \\
2\text{haversin}(c) &= 2\text{haversin}(a - b) + 2\sin(a)\sin(b)\text{haversin}(C)
\end{aligned} \tag{5.4}$$

Theorem 5.2 *The law of haversines*

$$\text{haversin}(c) = \text{haversin}(a - b) + \sin(a)\sin(b)\text{haversin}(C)$$

We note that a , b , and c given in radians are equal to the length of the arcs, since we are working on a unit sphere. Now let C be placed on the North Pole, then by definition its latitude is 90° or in radians $\frac{\pi}{2}$. Then, the length of the arcs a and b are:

$$\begin{aligned}
a &= \frac{\pi}{2} - B_{lat} \\
b &= \frac{\pi}{2} - A_{lat}
\end{aligned}$$

Let Δlat and Δlon be the latitude and longitude differences between A and B respectively, then since,

$$\sin(\phi) = \cos\left(\frac{\pi}{2} - \phi\right) \tag{5.5}$$

we have

$$\text{haversin}(c) = \text{haversin}(\Delta lat) + \cos(a_{lat}) * \cos(b_{lat}) * \text{haversin}(\Delta lon) \tag{5.6}$$

In the general case for a sphere with arbitrary radius the unit circle distance c is related to the real distance d on the sphere by: $c = \frac{d}{\text{radius}}$

Theorem 5.3 *The haversine formula*

$$\text{haversin}\left(\frac{d}{\text{radius}}\right) = \text{haversin}(a_{lat} - b_{lat}) + \cos(a_{lat}) * \cos(b_{lat}) * \text{haversin}(b_{lon} - a_{lon})$$

By isolating d the distance is found:

$$\sin\left(\frac{d}{2R}\right) = \pm \sqrt{\text{haversin}(\Delta lat) + \cos(a_{lat}) * \cos(b_{lat}) * \text{haversin}(\Delta lon)} \tag{5.7}$$

$$\sin\left(\frac{d}{2R}\right) = \sqrt{\text{haversin}(\Delta lat) + \cos(a_{lat}) * \cos(b_{lat}) * \text{haversin}(\Delta lon)} \tag{5.8}$$

$$\frac{d}{2R} = \arcsin(\sqrt{\text{haversin}(\Delta lat) + \cos(a_{lat}) * \cos(b_{lat}) * \text{haversin}(\Delta lon)}) \tag{5.9}$$

$$d = 2R * \arcsin(\sqrt{\text{haversin}(\Delta lat) + \cos(a_{lat}) * \cos(b_{lat}) * \text{haversin}(\Delta lon)})$$

(the distance formula)

In (5.9) \arcsin returns the value of $\frac{d}{2R}$ in the interval $[0, \frac{\pi}{2}]$. There is also a solution to (5.8) in $[\frac{\pi}{2}, \pi]$ but this solution is not relevant since we are interested in the value of $\frac{d}{R}$ in $[0, \pi]$.

Distance formula in Python

The distance formula is easily implemented in Python, shown in Listing 5.6. An example of calculating the distance between two points is shown in Listing 5.7.

```
from math import sin, atan2, cos, sqrt, pi

earth_radius = 6371
radian = pi / 180

def distance(point1, point2):
    '''Calculates the distance between two points, using the
       haversine formula.
    '''
    lat_a, lon_a = to_radians(point1)
    lat_b, lon_b = to_radians(point2)
    dlon = lon_b - lon_a
    dlat = lat_b - lat_a

    haversin_c = (sin(dlat/2))**2 + cos(lat_a) * cos(lat_b) * (
        sin(dlon/2))**2

    # unit sphere distance
    c = 2 * asin(sqrt(haversin_c))

    distance = earth_radius * c
    return distance

def to_radians(point):
    '''Convert to radians.
    '''
    lat = point.lat * radian
    lon = point.lon * radian
    return lat, lon
```

Listing 5.6: *Python implementation of distance calculation using the haversine formula*

```
>>> import geo, math
>>> class Point:
...     def __init__(self, lat, lon):
...         self.lat = lat
...         self.lon = lon
...
>>> aarhus = Point(56, 10)
>>> aalborg = Point(57, 10)
>>> geo.distance(aarhus, aalborg)
111.12511347447912
```

Listing 5.7: *Calculating the distance between Aarhus and Aalborg*

Why haversine? A small experiment

We derived haversine based on the statements that using the spherical cosines relation to calculate distance would cause significant rounding errors with small distances. What is small in a modern context? In order to reason about this we have implemented the direct approach to the distance calculation as well and compared it with distance calculation based on the haversine formula.

Point	Point	Haversine Distance (m)	Cosines Distance (m)
(56.0,10.0)	(57.0,10.0)	111125.113474	111125.113474
(56.0,10.0)	(56.0000009537,10.0)	0.105977166514	0.0948756933212
(56.0,10.0)	(56.0000004768,10.0)	0.0529885829037	0.0

Table 5.4: Comparing distances calculated with haversine and cosines

Table 5.4 summarizes the results using Python on a 32-bit machine. There is a maximum difference of 5 cm between the two approaches. We conclude that the reasoning for the haversine formula is irrelevant in a modern context for most distance calculations on the Earth. Since we have derived haversine we continue using it. The reader can refer to Appendix C for the implementations and the whole transcript of the output of the experiment.

Sorting by distance

Sorting points by distance is easy exploiting the haversine function, as shown in Listing 5.8. Python’s included `sort` function can sort using any comparator. In this case the comparator calculates the distance, using the distance formula, between the current point and the reference point.

```
def sort_by_distance(ref_point, entities):
    '''Sort entities with a point property, by distance wrt. a
        reference point.

    Args:
        entities: list of entities with a point property.
        ref_point: reference point to calculate distance from.
    Returns:
        list of entities sorted by distance. The distance is
            added as a member of the entity.
    '''
    def _distance_to(entity):
        d = distance(ref_point, entity.point)
        entity.distance = d
        return d

    entities.sort(key=_distance_to)
    return entities
```

Listing 5.8: Sorting points with haversine

There is a problem with the approach using the haversine formula; it is impossible to calculate an index in the Google datastore to support the distance formula. Therefore finding spots within a certain distance involves a full table scan, iterating over all spots in the database, and for each spot calculating the distance to the reference point.

Haversine is useful for calculating the distance between few points, but increasing the number of points demands another approach that reduces the data set and exploits indexing and preserves scalability. In the next section, we present an alternative approach.

Calculating proximity points with Geohash

Geohash, invented by Gustavo Niemeyer, is an algorithm to convert between latitude and longitude coordinates and a hash value [?]. In essence, geohash is just a z-order space-filling curve. Therefore a database can benefit from a usual index on geohashed points when performing searches for points in the proximity of each other.

Geohash is based on a custom *Base32* alphabet for encoding and decoding binary data. The custom alphabet is given by the string "0123456789bcdefghjkmnpqrstuvxyz". The index of a character in the string is equal to the value of the character, e.g., 'z' has the associated value 31 which is 11111₂.

Geohash supports arbitrary precision in the encoding of geographical points by interleaving the latitude and the longitude in the bit string. A bit string such as $b = 11111$ is decoded to a latitude longitude point as follows:

- Even positioned bits pertain to longitude: $b_{lon} = b_0 b_2 b_4 = 111$
- Odd positioned bits pertain to latitude: $b_{lat} = b_1 b_3 = 11$

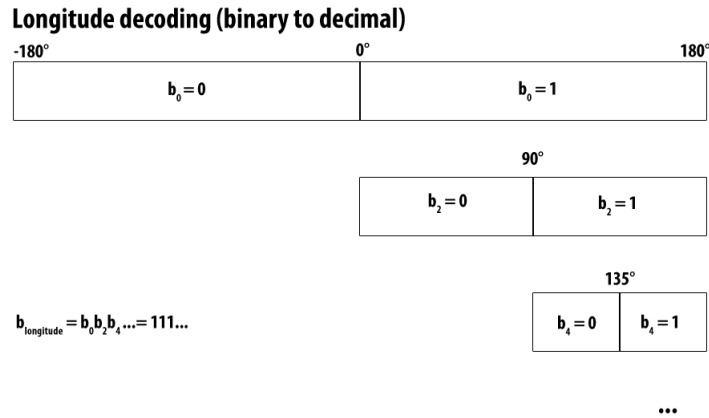


Figure 5.2: *Geohash longitude calculation*

Every bit in b_{lon} and b_{lat} halves the point's location space. When the bit is set the upper half is chosen. The longitude of geohash 'z' is therefore between $[135^\circ, 180^\circ]$, as shown in Figure 5.2, and the latitude is between $[45^\circ, 90^\circ]$. All points with prefix 'z' are in this area. A specific point is calculated by taking the mean value of the range; thus, the geohash 'z' equals the point $(67.5^\circ, 157.5^\circ)$.

There are two Python geohash implementations available: [?] which is in the public domain and [?] which is licensed under GNU Affero General Public License (AGPL)⁵. The former is a bit more advanced; it supports addition of geohashes, and has two formats for the hash: binary and character based. The result of an addition is the minimal bounding box of the two geohashes. An example of the usage of the first implementation is shown in Listing 5.9. Note that the usual convention of the sequence latitude and longitude is reversed to longitude and latitude. The last line in the example shows that the implementation is buggy when decoding hashes. We have just shown above that the point represented by 'z' was $(67.5^\circ, 157.5^\circ)$ and this

⁵AGPL is a somewhat more restrictive license than GPL. A Web application based on GPL software must only open source the software if the software running the service is distributed to others. AGPL, in essence, demands that the software running the Web application is open sourced also if the software is not distributed.

is not the case here, the longitude is wrong. On the server in the application we encode geographical points using [?].

```
>>> import geohash
>>> london_west = geohash.Geohash((-0.12, 51.50))
>>> str(london_west)
'gcpuvr295zcd2'
>>> london_east = geohash.Geohash((0.12, 51.50))
>>> str(london_east)
'u10hfyr3hpy6q'
>>> world = london_east + london_west
>>> str(world)
'',
>>> geohash.Geohash('z').point()
(180.0, 67.5)
```

Listing 5.9: *Geohash example usage*

Listing 5.9 also shows a general problem with geohashing: the property of similar prefixes of points in the proximity of each other is not always true. The problem lies within proximity points with different prefix. Take for instance the area just West of London, Greenwich; the hash of the western side (bit 0 unset), has a complete different prefix than the hash of the eastern side of Greenwich (bit 0 set) because boxes have a borderline here: namely the prime meridian. Restricting proximity points to a certain prefix results in lost proximity points.

Lost proximity points: a solution

The Base32 alphabet and the algorithm of geohash result in a conceptual grid that is overlayed on the Earth, shown in Figure 5.3. This grid is essential in solving the problem of lost proximity points.

A solution to the proximity problem is augmenting the returned data set with points from all the grid boxes surrounding the requested one: the neighbors directly to the north, northeast, east, etc.

The geohash grid has a recursive structure: every char in a geohash specifies a grid box container, and the next char in the geohash specifies a grid box within the first grid box container, and so on. The position of the char, even or odd, decides which of the two containers in the figure is representative for the char, since even positioned chars begin with an even bit, and odd positioned chars begin with an odd bit.

We have implemented a solution where the two representative grids are represented as matrices: allowing for easy navigation to the neighbors of a grid box. Navigating to neighbors falls into three cases:

1. the neighbor is in the same container;
2. the neighbor is in a neighbor container; or,
3. the neighbor is on the other side of the cylindrical projection, i.e., crossing the ends of the square cylindrical projection.

Locating a neighbor in the first case is easy: for all neighbors return their geohash, which is the current geohash with the last char replaced by the char representing the new location in the container.

Locating the neighbors in the second case is more tricky. The case is identified by going out of bounds of the matrix. By wrapping around the indexes we identify

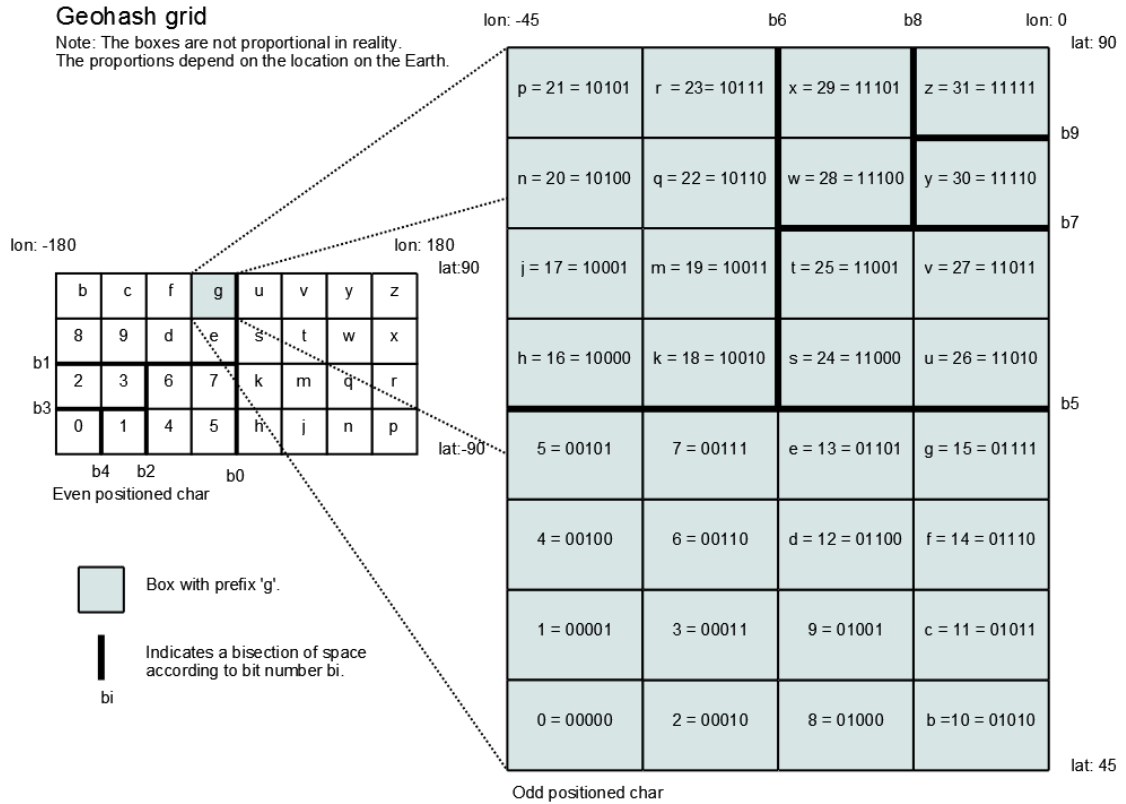


Figure 5.3: *Cylindrical projection of Geohash grid*

the new suffix of the grid box we are after, i.e., if the request is for the Eastern neighbor of geohash 'z' the index is wrapped around to the other side which is 'b' (in the even case). However, as we are out of bounds, we need to zoom out and still find the eastern neighbor. Zooming out and tracking the neighbor is equal to doing the same query on the geohash with the last char stripped.

If the container of 'z' has an eastern neighbor the tracking succeeds and we can return the char of the eastern container concatenated with the suffix identifying the grid box within it, found when wrapping around.

The third case happens when we are fully zoomed out: in this case there exist no neighbors in the grid; however, the logical neighbor is still found by wrapping around (this only makes sense for East-West directions). We can think of this as connecting the ends of a square map of the world assuming it is created by a cylindrical projection.

An excerpt of the implementation of geohash neighbors is shown in Listing 5.10. The main method is `step_in_direction` which returns the grid box, at the same zoom level, which is the neighbor to the box of the given geohash in the requested direction. The function calls itself recursively when it must zoom out to find the neighbor. Under the recursive call it keeps the found suffix of the neighbor in the `new_geohash_suffix` argument, and when a neighbor is found the char of its grid box is concatenated with the `new_geohash_suffix` resulting in another geohash suffix. At last the suffix is prefixed with the geohash of the current zoom level. We note that the time used to find neighbors is in the worst case proportional to the length of geohash (the number of times to zoom out).

```
def neighbors(geohash):
```



```

'''
Calculate surrounding geohashes.
Precondition:
    len(geohash) > 0
Returns:
    A list of neighbors as a list of geohashes, in the
    order N, NE, E, SE, ...
'''
geohash_north = step_in_direction(north, geohash)
geohash_east = step_in_direction(east, geohash)
geohash_south = step_in_direction(south, geohash)
geohash_west = step_in_direction(west, geohash)

geohash_north_east = step_in_direction(east, geohash_north)
geohash_north_west = step_in_direction(west, geohash_north)

geohash_south_east = step_in_direction(east, geohash_south)
geohash_south_west = step_in_direction(west, geohash_south)

return [
    geohash_north,
    geohash_north_east,
    geohash_east,
    geohash_south_east,
    geohash_south,
    geohash_south_west,
    geohash_west,
    geohash_north_west
]

def step_in_direction(fct, geohash, new_geohash_suffix=''):
'''
Step to the neighbor grid box of the grid box given by the
geohash.

Args:
    geohash: geohash string to find neighbor of.
    fct: the direction to step in as a function (north,east
        ,west,south)
    new_geohash_suffix: DON'T set this, it is used to bring
        the values around
        in the recursive calls.

Returns:
    the geohash of the neighbor grid box.
'''
# geoh = c0c1c2
# len(geoh) % 2 = 1 # which is True
even_index = len(geohash) % 2
prefix = geohash[:-1] # cut off last
suffix = geohash[-1] # last elem

# convert to matrix indexes
i,j = reverse_geohash(suffix[0], even_index)
new_i, new_j = fct(i,j)

# check for out of bounds

```

```

zoom_out = False
if new_i == -1:
    # west out of bounds
    new_i = (_MAX_ROW_EVEN if even_index else _MAX_ROW_ODD)
    zoom_out = True
elif new_i == (_ROWS_EVEN if even_index else _ROWS_ODD):
    # east out of bounds
    zoom_out = True
    new_i = 0

if new_j == -1:
    # north out of bounds
    zoom_out = True
    new_j = (_MAX_COL_EVEN if even_index else _MAX_COL_ODD)
elif new_j == (_COLS_EVEN if even_index else _COLS_ODD):
    # south out of bounds
    zoom_out = True
    new_j = 0

if zoom_out:
    if len(prefix) == 0:
        # we are falling off the flat earth!
        # east / west fall off: wrap around earth
        if fct == east or fct == west:
            return geohash_matrix(new_i, new_j, even_index) +
                new_geohash_suffix
        # NOTE: It does not make sense to wrap around from
        # What we are doing instead doesn't make much sense
        # either
        # Think off the pole as a pie with each piece
        # identified by a geohash char
        # we return the piece on the opposite side
        if fct == north or fct == south:
            # pie has 8 pieces j-4 = the one on the other
            # side
            return geohash_matrix_even[i][j-4] +
                new_geohash_suffix
        return step_in_direction(fct, prefix, geohash_matrix(
            new_i, new_j, even_index) + new_geohash_suffix)

    return prefix + geohash_matrix(new_i, new_j, even_index) +
        new_geohash_suffix

```

Listing 5.10: *Geohash neighbors implementation*

With our geohash neighbors implementation it is possible to retrieve all the neighbors in the grid. Listing 5.11 shows an example usage which gets the geohash of London and retrieves the grid boxes around.

```

>>> import geohashneighbors as ghn
>>> import geohash
>>>
>>> london = geohash.Geohash((0, 51))
>>> str(london)
'u1040h2081040'
>>> ghn.neighbors('u10')
['u12', 'u13', 'u11', 'u0c', 'u0b', 'gbz', 'gcp', 'gcr']

```

Listing 5.11: *Retrieving geohash grid box neighbors*

Earth Surface Area	Geohash Length	Approximate Covered Area
510,072,000 km^2	1	15939750 km^2
	2	498117 km^2
	3	15566 km^2
	4	486 km^2

Table 5.5: *Approximate surface area covered by geohashes of different lengths*

What is proximity in this context? Proximity is specified by the length of the geohash; every bit in the geohash chars halves the proximity area, thus, every char reduce the current area by a factor of 2^5 .

The surface area of the Earth is 510,072,000 km^2 the area that a prefix of length i covers is approximated by the formula:

$$area(i) = \frac{510,072,000km^2}{(2^5)^i}$$

The formula is a rough estimate since the distance between longitudes is inconsistent (starting at 111km at the equator and going towards 0 at the poles), but good enough to give an idea of the covered surface area. Table 5.5 shows the covered surface area for geohash lengths from 1 to 4 based on the formula.

Resource URI	Action
/forecast_points/?geohash={geohash}	GET forecast points in the given geohash grid box
/forecast_points/?geohash_offset={geohash}	GET all forecast points starting from the given geohash offset
/spots/?geohash={geohash}	GET
/spots/?geohash_offset={geohash}	GET
/weather_stations/?geohash={geohash}	GET
/weather_stations/?geohash_offset={geohash}	GET

Table 5.6: *Resource view of proximity points calculating resources*

Three resources are relevant for proximity searches: forecast points, spots, and weather stations. The proximity searches are limiting the result data set, and the geohash prefix is clearly input into an algorithm; therefore, the geohash argument is given in the URIs as a query parameter. Eventually, the resources will need a split up because of the 1000 entity limit on the GAE. The `geohash_offset` is a field for paging. It specifies the geohash value to start returning entities from. The resource view in Table 5.6 shows the additional resources of the application.

5.5 Summary

In this chapter we presented all the resources of the application. Resources were summarized in what we call the resource view of the software architecture. In the last section we presented theory and implementations for distance calculations, and proximity queries. In the application we use a combination of haversine and geohashing: first reducing the returned points to a specific area with geohash and

afterward using haversine to sort and reduce the number of points to the ones closest to the user.

*"If you have built castles in the air,
your work need not be lost; that is
where they should be. Now put
the foundations under them."*

Henry David Thoreau
(1817 – 1862)

6

The We Love Wind Web Service

In this chapter we describe the implementation of the corner stone in the We Love Wind application: the We Love Wind Web Service. We implement the service using Google App Engine/Django, described in Chapter 4.

The last chapter covered all the resources of the Web Service; we now turn to the logic backing the resources. Since the architecture is a typical MVC architecture, the implementation is presented in terms of this style.

The interested reader can access the Web Service by pointing a browser to the following URIs:

- <http://www.welovewind.com/api/spots/>
- http://www.welovewind.com/api/weather_stations/
- http://www.welovewind.com/api/forecast_points/

6.1 Implementation of Models

The models are not just defining the structure of data; the models contain all the business logic of the application separating and encapsulating the data of the application from the presentation of it.

Figure 6.1 shows the data models of the application in a UML static structure diagram. Arguments to methods is left out for brevity. Methods underlined in the diagram are Python class methods. Since all data models inherit from the `db.Model` class it is omitted. The figure is a good mental model to have in mind during the descriptions of the different model classes.

6.1.1 AbstractGeoHash model

Listing 6.1 shows the AbstractGeoHash model. The purpose of the model is to provide geohash structure and logic for any model that inherits from it; in that way all properties and methods relevant for geohashing is factored out of the other models, i.e., spot, weather station, and forecast point.

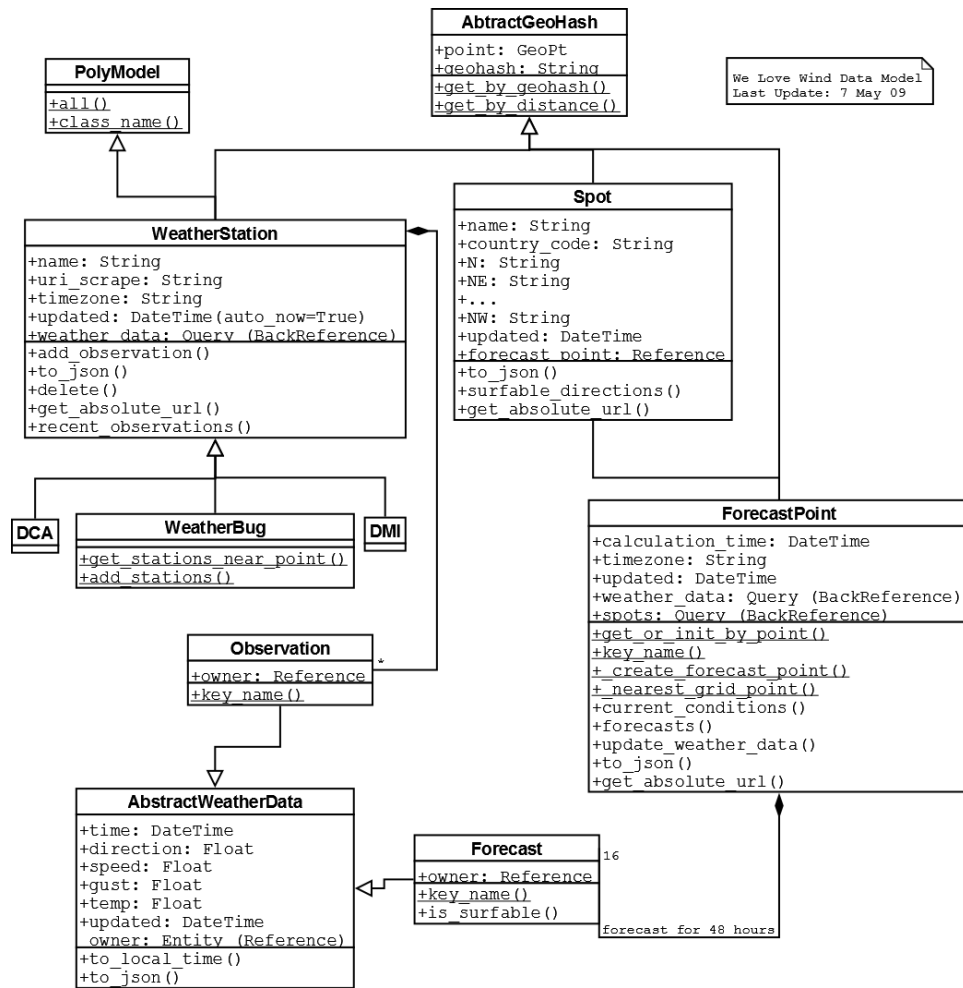


Figure 6.1: Data model (UML notation)

Structure

The model contains two properties: **point** and **geohash**. Point is the latitude and longitude for a place, and geohash is the geohash encoding of the latitude and longitude.

Logic

The model contains logic that gets all points of the current class that have the same geohash prefix as the given: manifested in **get_near_geohash**. GQL lacks the option to search for patterns in strings (LIKE in SQL), instead two inequality filter are applied to match the geohash prefix pattern; the last filter sets the upper boundary of the range of results to return since the prefix is concatenated with the largest possible unicode character: the replacement character which is encoded as `u"\ufffd"`.¹

In addition, the model includes logic which retrieves all points within a threshold distance of a given point: manifested in **get_near_point**. It applies the geohash neighbors technique to get neighbors. The method uses **get_near_geohash** to get

¹http://en.wikipedia.org/wiki/Replacement_character

all points in the geohash areas; afterward pruning the list to points within the threshold. Currently, the method assumes that the radius (distance) from the center point is within the geohash neighbors.

Instances of the `GeoHash` model automatically get a geohash value, since the `put` method is overwritten, and before saving the instance to the datastore the geohash value is calculated and set.

```
class AbstractGeoHash(db.Model):

    point = db.GeoPtProperty(required=True)
    geohash = db.StringProperty()

    @classmethod
    def get_near_geohash(cls, ghash_prefix, geohash_offset):
        ghash_prefix = ghash_prefix.lower()
        query = db.Query(cls)
        query.filter('geohash >=', geohash_offset)
        query.filter('geohash <', ghash_prefix + u"\ufffd")
        return query

    @classmethod
    def get_near_point(cls, point, distance,
                      geohash_prefix_length = 3):
        entities = []
        center = geohash.encode(point.lat, point.lon)[:
            geohash_prefix_length]
        ghashes = geohashneighbors.neighbors(center)
        ghashes.append(center)

        for h in ghashes:
            query = cls.get_near_geohash(h,h)
            for entity in query:
                entities.append(entity)

        def prune(list):
            '''Prune a sorted list by distance'''
            for i, point in enumerate(list):
                if point.distance > distance:
                    return list[:i]
            return list

        return prune(geo.sort_by_distance(point, entities))

    def put(self):
        self.geohash = geohash.encode(self.point.lat, self.
            point.lon)
        return super(AbstractGeoHash, self).put()
```

Listing 6.1: *GeoHash model*

6.1.2 Spot model

Listing 6.2 shows the Spot model.

Structure

The spot model inherits from the `AbstractGeoHash` model: getting both a `point` and a `geohash` property. Every spot has a `name`, and, as discussed in the resource section, every spot belongs to the nearest `forecast_point`. The many-to-one relationship between spots and forecast points is implemented with the reference property. The wind diagram is defined as string attributes that must have one of four values: NO, MAYBE, YES, or UNDEFINED.

The model has an `updated` property which is a pervasive property in the data models. The property is automatically updated to the current datetime each time the model is saved. This is useful for tracking the times of last modification of entities, e.g., essential for implementing HTTP caching.

Logic

There are two pervasive methods located on concrete models: `to_json` and `get_absolute_url`. The `to_json` method returns a simple Python object representation of the instance ready for JSON serialization. The method will prove its worth when we construct resource representations in the controllers in Section 6.2.

The `get_absolute_url` method return the URI for the concrete model; the advantage is that changing the URI can now be done by just changing that method, instead of refactoring all views.

As in the case with the `AbstractGeoHash` model, the `put` method is overwritten to wire in some data. After saving a spot it is always associated with the nearest forecast point and, in addition, the `geohash` value is wired into the spot.

```
class Spot(AbstractGeoHash):
    WIND_DIAGRAM_CHOICES = ('NO', 'MAYBE', 'YES', 'UNDEFINED')

    name = db.StringProperty(required=True)
    updated = db.DateTimeProperty(auto_now=True)
    country_code = db.StringProperty(required=True)
    forecast_point = db.ReferenceProperty(ForecastPoint,
        collection_name='spots')

    # wind diagram
    N = db.StringProperty(default='UNDEFINED', choices=
        WIND_DIAGRAM_CHOICES)
    (...)
    NW = db.StringProperty(default='UNDEFINED', choices=
        WIND_DIAGRAM_CHOICES)

    def put(self):
        self.forecast_point = ForecastPoint.
            get_or_init_by_point(self.point)
        return super(Spot, self).put()

    @classmethod
    def key_name(self, country_code, name):
        return '/spots/%s/%s/' % (country_code, name)

    def to_json(self):
        return {
            'name': self.name,
            'country_code': self.country_code,
            'lat': self.point.lat,
            'lon': self.point.lon,
```



```

        'wind_diagram': self.surfable_directions(),
        'uri': self.get_absolute_url(),
        'forecast_point': self.forecast_point.get_absolute_url()
    }

    def get_absolute_url(self):
        return self.key().name()
(...)

```

Listing 6.2: *Spot model*

6.1.3 AbstractWeatherData Model

The `AbstractWeatherData` model, shown in Listing 6.3 is a common supermodel for forecasts and observations. Forecasts and observations have similar structure except for the fact that forecasts belong to a forecast point and observations belong to a weather station.

Structure

The `WeatherData` model contains all the relevant weather data types mentioned before: `speed`, `gust`, `temp`, `direction`, and a `time` of when the weather data applies. The model is abstract delegating the references to its subclasses: `Observation` and `Forecast`.

```

class AbstractWeatherData(db.Model):
    '''Abstract weather data model.

    Attributes:
        is_new: indicates if this model is just created. Used
            by update_or_insert()

    Properties:
        updated: time of update on this server
        time: the observation / forecast time
    '''
    updated = db.DateTimeProperty(auto_now=True)
    direction = db.FloatProperty()
    speed = db.FloatProperty()
    temp = db.FloatProperty()
    gust = db.FloatProperty()
    time = db.DateTimeProperty(required=True)
    #abstract owner = db.ReferenceProperty()

    def __init__(self, is_new = False, **kwargs):
        self.is_new = is_new
        super(AbstractWeatherData, self).__init__(**kwargs)

    @classmethod
    def update_or_insert(cls, **kwargs):
        key_name = cls.key_name(**kwargs)
        entity = super(AbstractWeatherData, cls).get_or_insert(
            key_name, is_new=True, **kwargs)
        if not entity.is_new:
            for prop in kwargs.keys():
                setattr(entity, prop, kwargs[prop])

```

```

        entity.put()
    return entity
(...)
```

Listing 6.3: *AbstractWeatherData model*

Logic

In an environment where weather data is continuously added eventually the datastore will fill up. The datastore limit demands that some house keeping is done. GAE cron jobs are limited in their number and duration time [?] and are not a scalable solution. An alternative to cleaning up the datastore with cron jobs is needed.

An idea is continuously maintaining the datastore as new data is fed into it by throwing out the oldest data. If we let, e.g., the last 24 hours of observations and only the forecasts from the last calculation be in the datastore the amount of weather data is pruned.

We let the weather data have a unique property, pertaining to the owner (weather station or forecast point) and the time, e.g., a time delta such as total time of day in minutes. When these properties are unique, inserting new data with the same time delta and owner updates the existing entity. Implementation of unique properties in the context of forecasts was presented in Section 4.5.1 on page 32.

If the incoming data is consistently produced at certain times of the day the technique is *First In, First Out* with a buffer size proportional to the number of atomic units in the maximum time delta. In the case of observations the buffer size is one day and the unit size is in minutes; that gives a maximum of 60 minutes * 24 hours = 1440 entities in the datastore per weather station. In the case of forecasts the buffer size is 48 hours, however, since the time of forecasts are always in hours the unit size is in hours; that gives a maximum of 48 entites per forecast point.

6.1.4 WeatherStation Model

Listing 6.4 shows the WeatherStation model.

Structure

The model inherits from the `PolyModel` and the `AbstractGeohash`. The `PolyModel` class provides a programmatic implementation of object inheritance in the datastore similar to using NULL values to combine hierarchy relations in relational databases [?, sec.4.6]. The advantage is that a query on the superclass includes all the entities from the subclasses, and a query on the subclass only includes entities from the subclass. The model has a list of observations. This is not visible in the model since it is the reference property in the observation model that defines the one-to-many relationship between weather stations and observations.

Logic

The model is the owner of weather observations. As mentioned we overwrite those observations on a circular basis. The model includes a method, `update_or_insert_observation`, that takes an observation resource in deserialized JSON format and either inserts the observation or overwrites an existing based on the unique properties of the observation model.

```

class WeatherStation(polymodel.PolyModel, AbstractGeoHash):
    """Super model class for weather stations.
    Properties:
        name: name of weather station.
        updated: time of last update.
        country_code: country code of country where weather
            station is.
        weather_data: associated weather data
        uri_extern: URI to weather station
        uri_scrape: URI to scrape
    """
    name = db.StringProperty(required=True)
    uri_scrape = db.StringProperty()
    country_code = db.StringProperty(required=True)
    timezone = db.StringProperty(required=True)
    updated = db.DateTimeProperty(auto_now=True)
    # weather_data ref. from observation

    @classmethod
    def key_name(cls, country_code, slug):
        return '/weather_stations/%s/%s/' % (country_code, slug
        )

    def update_or_insert_observation(self, json):
        """Update or insert observation to this station's
        weather data.

        Args:
            json: deserialized json weather data
                time (required): isoformat time
                speed (not required):
                gust (not required):
                direction (not required):
        Returns:
            the updated/new observation.
        """
        dt = wlttime.ISO2dt(json['time'])

        entity = Observation.update_or_insert(
            owner=self,
            time=dt,
            speed = json.get('speed', None),
            gust = json.get('gust', None),
            temp = json.get('temp', None),
            direction = json.get('direction', None)
        )
        return entity

    def to_json(self):
        return {
            'name': escape(self.name),
            'type': self.class_name(),
            'timezone': self.timezone,
            'lat': self.point.lat,
            'lng': self.point.lon,
            'uri_scrape': self.uri_scrape,
            'uri': self.get_absolute_url(),
            'uri_observations': '%sobservations/' % self.

```

```

        get_absolute_url()
    }
    (...)

```

Listing 6.4: *WeatherStation model*

6.1.5 ForecastPoint model

Listing 6.5 shows the `ForecastPoint` model.

Structure

The forecast point resource, presented in Section 5.3, contains a number of forecasts for a specific forecast point. This is a many-to-one relationship between the forecast point and the forecasts of the point. The relationship is manifested by the reference property from the `Forecast` Model to the `ForecastPoint` model.

A forecast point is uniquely identified by its latitude and longitude; we incorporate the point information into the key name of every instance hereby ensuring that no two instances can have the same latitude and longitude.

Logic

It is useful if the application does not load the datastore with all forecast points in the world, since the forecast points are continuously updated. Surf spots will for sure not use the majority of the forecast points. Therefore the forecast points are added just-in-time when a spot is created. The `get_or_init_by_point` is the method that contains the logic to add forecast points just-in-time. The method takes the following sequence of steps:

1. Since forecast points are located in a granularity of 0.5° degrees, it calculates the nearest grid point in the gfs grid.
2. Checks for an existing point and returns the point if so; else, if the point does not exist, the point is created. Creating the point includes retrieving the timezone of the point from geonames.org and retrieving forecasts from the author's weather service at DAIMI, presented in Chapter 7.

The process includes calls to two external resources which obviously includes a significant latency. However, it is only the first time the application access a particular point that the latency is present.

The `update_weather_data` method takes a weather forecast representation, shown in Listing 5.4 on page 44, and updates the forecast point with all the forecast in the representation. It iterates through all the forecasts using the `update_or_insert` method everytime putting the updated forecast to the datastore.

```

class ForecastPoint(AbstractGeoHash):
    timezone = db.StringProperty(required=True)
    calculation_time = db.DateTimeProperty()
    # weather_data ref.

    @classmethod
    def get_or_init_by_point(cls, point):
        '''Get the forecast point nearest the given lat, lon.

        Returns:
            the (new saved) forecast point.

```

```

'''
nearest = ForecastPoint._nearest_grid_point(point)

# Get the forecast point for the given point.
forecast_point = ForecastPoint.get_by_key_name(
    ForecastPoint.key_name(nearest.lat, nearest.lon))

if forecast_point:
    return forecast_point
else:
    return ForecastPoint._create_forecast_point(nearest
)

@classmethod
def _create_forecast_point(cls, nearest):
    '''Create new forecast point.

    Args:
        nearest: nearest grid point (GeoPt)
    Returns:
        forecast point with forecasts added.
    '''
    try:
        forecasts = daimiweather.get_forecasts(nearest)
    except Exception:
        forecasts = []
        calc_time = datetime.datetime.utcnow()
        tzinfo = {'timezoneId': 'UTC'}

    if forecasts:
        # timezone from geonames.org
        try:
            response = geoinfo.geonames_timezone(nearest)
            tzinfo = simplejson.loads(response.content)
            calc_time = wlttime.ISO2dt(forecasts['
                calculation_time'])
        except WSEException:
            tzinfo = {'timezoneId': 'UTC'}
            calc_time = datetime.datetime.utcnow()

    forecast_point = ForecastPoint(
        key_name=ForecastPoint.key_name(nearest.lat,
            nearest.lon),
        calculation_time=calc_time,
        point=nearest,
        timezone=tzinfo['timezoneId'])

    forecast_point.put()

# add all forecasts
if forecasts:
    for f in forecasts['weather_data']:
        Forecast.update_or_insert(
            calculation_time=calc_time,
            owner = forecast_point,
            speed = f.get('speed', None),
            temp = f.get('temp', None),
            direction = f.get('direction', None),

```

```

        time = wlwtime.ISO2dt(f.get('time'))

    return forecast_point

def forecasts(self, limit=19):
    """
    Returns:
        the weather data valid now and in the future.
    """
    now = datetime.datetime.utcnow()
    now -= datetime.timedelta(hours=3)
    return self.weather_data.filter('time >=', now).order('
        time').fetch(limit)

def update_weather_data(self, post_data):
    """Update this point's weather data.

    Args:
        post_data: dictionary of key / value pairs from
            POST.
        calculation_time: iso time string of
            calculation time.
        weather_data: the weather data to update with.
    """
    self.calculation_time = wlwtime.ISO2dt(post_data['
        calculation_time'])

    for data in post_data['weather_data']:
        Forecast.update_or_insert(
            calculation_time=self.calculation_time,
            owner=self,
            speed = data.get('speed', None),
            temp = data.get('temp', None),
            direction = data.get('direction', None),
            time = wlwtime.ISO2dt(data.get('time')))
    self.put()
(...)

```

Listing 6.5: *Forecast model*

6.2 Implementation of Controllers

The resources of the Web Service are composed of the data models in our GAE/Django application. The resources are not equal to the data models, however, the resources are equal to the URIs exposed, which are all backed by a controller. We summarized the resources in the resource views in Chapter 5; we now realize them in Django.

URIs are modeled in Django without any restrictions by the framework. URIs are mapped to controllers by regular expressions that matches on the request's URI. The regular expressions together with the controllers are registered in the Python module `urls.py` shown in Listing 6.6. The code defines patterns that map from an URI to a controller.

The regular expression functionality used is somewhat esoteric; therefore, we give a short description of it. `\w` is a predefined set of characters, namely the character class consisting of all alphanumeric characters. A character class is defined by the

content inside the metacharacters [and]. In the case of [\w-] the character class is the set of alphanumeric characters and a dash. (?P<any_name>...) defines a named group; the matched content of these can later be referred to by name.

```
urlpatterns = patterns('',
    (r'^api/weather_stations/(?P<country_code>\w+)/(?P<id>[\w-]+)/$', WeatherStationsController()),
    (r'^api/weather_stations/(?P<country_code>\w+)/$', WeatherStationsController()),
    (r'^api/weather_stations/$', WeatherStationsController()),

    # weather observations
    (r'^api/weather_stations/(?P<country_code>\w+)/(?P<ws_id>[\w-]+)/observations/$', weather_observations.
        ObservationsController()),
)

# forecasts
urlpatterns += patterns('',
    (r'^api/forecast_points/info/$', forecasts.info),
    (r'^api/forecast_points/(?P<id>.+)/$', forecasts.
        ForecastPointsController()),
    (r'^api/forecast_points/$', forecasts.
        ForecastPointsController()),
)

# spots
urlpatterns += patterns('',
    (r'^api/spots/(?P<country_code>\w+)/(?P<id>[\w-]+)/$', spots.
        SpotsController()),
    (r'^api/spots/(?P<country_code>\w+)/$', spots.SpotsController
        ()),
    (r'^api/spots/$', spots.SpotsController()),
)
```

Listing 6.6: *urls.py*

The patterns are quite powerful: they capture all requests to individual and list resources of spots, forecast points, weather stations, and weather observations, and route them to a corresponding controller. The controllers are usually Python functions, but we take advantage of the fact that the controller can be any callable object. The reason is that we want separation of concerns with regard to the HTTP method of the requests. Thus, instead of handling the different HTTP method types in each controller function we define a common REST controller class that route to a Python method corresponding to the HTTP method.

6.2.1 REST controller

Table 6.1 shows the methods of the REST controller class and how the methods correspond to URIs and HTTP method. As mentioned in Section 2.4.3 the POST method can be used to support PUT and DELETE from (X)HTML by overloading it. The approach is supported by the REST controller. The URI pattern in the table is common in Web development; having an items resource, and a resource for the specific items.

Listing 6.7 shows the REST controller. Each of the Python methods `get`, `index`, `post`, `put`, and `delete` correspond to a resource action that must be overridden

Resource action	URI	HTTP Method
<code>delete</code> , subordinate resource	<code>/ {res.} / {res. id} /</code>	DELETE
<code>index</code> , of all subordinate resources	<code>/ {res.} /</code>	GET
<code>post</code> , new subordinate resource	<code>/ {res.} /</code>	POST
<code>put</code> , existing subordinate resource	<code>/ {res.} / {res. id} /</code>	PUT
<code>show</code> , subordinate resource	<code>/ {res.} / {res. id} /</code>	GET

Table 6.1: *REST controller-action to URI and HTTP method mapping*

in the subclass inheriting from the REST controller in order to support the action. Invoking a method that is not overridden in the subclass causes an HTTP 405 `Method Not Allowed`. A 405 must include the allowed methods [?, sec.10.4], therefore the allowed methods are extracted and returned.

The controller must distinguish between resources and subordinate resources. The subordinate case is distinguished from the other by having the key name in its URI. This is used in the method `_get_permitted_methods` and `_handle_get` and `_handle_post` to distinguish the two cases.

The Django patterns take any callable object as input, e.g., a function or a method. An instance is callable if it contains the `__call__` method. This is used in the REST controller: upon start up of the Django application an instance of the `RESTController` is created. A pattern match triggers a call to the corresponding instance, hereby executing the `__call__` method.

The first argument given by Django to the callable object is always an `HttpRequest`, additional values captured by the regular expression are parsed as subsequent arguments. Therefore the `__call__` method takes a request and an arbitrary number of keyword arguments: the values from the named groups in the regular expressions. The method saves relevant values in the instance and does a reflective look-up based on the request's HTTP method. The method look-up resolves to one of four methods: `_handle_delete`, `_handle_get`, `_handle_post`, or `_handle_put`. `_handle_post` is interesting since it handles overloaded POST. If POST is overloaded yet another reflective look-up is made according to the value of the `_method` query parameter.

The controller supports two representation types from the client: the usual uri-encoding type, and the `application/json` type. All browsers upon submitting forms with POST convert the form data to uri-encoding, and put the data in the message body of a POST request.

In Django, on the first access to the POST or GET data of a `HttpRequest`, Django parses the request data into a query dictionary that expects uri-encoded data. This is useful if data is indeed in that format. This is not the case when JSON is used as representation format. Parsing of the data is bypassed by accessing the `raw_post_data` member of the request, used in the `_load_post_data` that parses data according to the type of the content: JSON or uri-encoded.

```
import os
import cgi

from webob.multidict import MultiDict

from google.appengine.ext import webapp
from google.appengine.api import users

from django.utils import simplejson
```



```

from helpers.decorators import require_login

class RestController(webapp.RequestHandler):
    """Route HTTP requests to the corresponding method wrt. the
        HTTP method and argument.

    In addition to supporting usual PUT and DELETE, PUT and
    DELETE are also handled by use of
    overloaded POST:
        resource/{id}/?_method=PUT
        resource/{id}/?_method=DELETE

    Args:
        **kwargs
            id: if specefied taken to be the id of a concrete
                resource
    returns:
        response - HTTP response
    """
    def doIndex(self, *args):
        self.error(405)

    def doShow(self, *args):
        self.error(405)

    def doPost(self, *args):
        self.error(405)

    def doGet(self, *args):
        self.error(405)

    def doDelete(self, *args):
        self.error(405)

    # overwritten get, put, post
    def get(self, key=None):
        if key:
            # GET the concrete item
            return self.doShow(key)
        else:
            # GET the items
            return self.doIndex()

    @require_login
    def put(self, *args):
        self._load_post_data()
        return self.doPut(*args)

    @require_login
    def post(self, *args):
        self._load_post_data()
        method = self.request.get('_method')
        if method:
            try:
                callback = getattr(self, 'do%s' % method.
                    capitalize())
                return callback(*args)

```

```

        except ValueError:
            self.error(405)
    else:
        return self.doPost()

    def _load_post_data(self):
        content_type = self.request.content_type
        if content_type == 'application/json':
            self.data = simplejson.loads(self.request.body)
        elif content_type == 'application/x-www-form-urlencoded':
            # LOOK INTO THIS
            fs = cgi.FieldStorage(fp=self.request.body_file,
                                environ=self.request.environ.
                                    copy(),
                                keep_blank_values=True)
            self.data = MultiDict.from_fieldstorage(fs)
        else:
            raise "Unsupported content type: %s" % content_type

```

Listing 6.7: *REST controller*

With the base controller in place we now turn to implementing the concrete controllers of the resources; all inheriting from the REST controller.

6.2.2 Spots controller

The spots controller implements the spot resources from Table 5.1 on page 41. Clients create, read, and update spots reflected in the controller that implements `index`, `put`, and `post`.

`index` serves all spots potentially filtered by country in a JSON representation. Several query parameter filters can be applied: `geohash`, and `geohash_offset`.

The JSON resource representation is created by first assembling a simple Python object structure and then using `simplejson`² to convert the structure to a JSON encoded string. By simple we mean that the structure uses only simple Python constructs like dictionaries, lists, unicode, etc. that have a corresponding construct in JSON; thus, allowing serialization.³

The representation of spots includes a maximum 100 spots. Paging of spots is handled with the `geohash_offset` query parameter. The application always returns spots in ascending order sorted by geohash value. When more than 100 spots are present the representations point to the next geohash value in the datastore; an example:

```

(...)
"next": "/api/spots/?geohash_offset=swgv3sxf003p"
(...)

```

A request for that `next` resource includes the 100 next spots with geohash greater than or equal to `geohash_offset`. This somewhat odd approach is the only way to page geohash models. In-equality filters in queries are restricted to a single property; in this case used on geohash; therefore, no other property can act as the property for paging.

New spots are created with `post`. The incoming representation is validated with Django forms and saved if it is valid. A JSON representation is returned both on

²JSON encoder/decoder for Python

³The documentation for `simplejson` includes all the allowed constructs <http://simplejson.googlecode.com/svn/tags/simplejson-2.0.9/docs/index.html>

success and failure caused by validation errors. The representation fits the format of the jQuery forms plugin, which we use in Ajax frontend.

put is similar to post, except that it demands existence of the resource.

```
PAGE_SIZE = 100

class SpotsController(RestController):

    def index(self):
        geohash = self.request.GET.get('geohash')
        geohash_offset = self.request.GET.get('geohash_offset',
            '')

        if geohash:
            spots = Spot.get_near_geohash(geohash,
                geohash_offset)
        else:
            spots = db.Query(Spot).order('geohash').filter('
                geohash >=', geohash_offset)

        country_code = self.kwargs.get('country_code')
        if country_code:
            spots.filter('country_code =', country_code)

        # fetching one more to check for has_next
        entities = spots.fetch(PAGE_SIZE + 1)

        qs = None
        has_next_page = (len(entities) > PAGE_SIZE)

        if has_next_page:
            qs = self.request.GET.copy()
            qs['geohash_offset'] = entities[-1].geohash

        repr = {
            'next': '%s?%s' % (self.request.META['PATH_INFO'],
                qs.urlencode()) if has_next_page else '',
            'items': [s.to_json() for s in entities]
        }

        return http.HttpResponse(simplejson.dumps(repr),
            mimetype='application/json')

    def put(self, id):
        (...)

    def post(self):
        form = SpotForm(self.request.POST)

        if not form.errors:
            spot = form.save()

            # add weather stations
            stations = WeatherBugWeatherStation.
                get_stations_near_point(spot.point)
            keys = WeatherBugWeatherStation.add_stations(
                stations)
```

```

        result = {'type': 'success', 'message': 'Spot %s
                created' % spot.name, 'spot': spot.to_json()}
    if len(keys) != 0:
        result['message'] += '. Added Weather stations:
                %s' % keys

    json = simplejson.dumps(result)
    response = http.HttpResponse(json, mimetype='
        application/json')
    response.status_code = 201
    response['Location'] = spot.get_absolute_url()
    return response
else:
    result = {'type': 'error', 'message': form.errors.
        as_ul()}
    return http.HttpResponse(simplejson.dumps(result),
        mimetype='application/json')

```

Listing 6.8: *Spots controller*

An interesting thing with `post` is that it adds weather stations from WB in the area of the new spot to the application. Section 8.2.2 on page 85 presents how.

6.2.3 Weather stations controller

The weather station controller implements the resources in Table 5.2 on page 43. It merely serves JSON representations of the weather station list resource, corresponding to the `index` method.

The `index` method differs from the corresponding method of the spots controller only in adding filters on the type of weather station, and on the type of the time format used in the representation; therefore, the code is omitted.

6.2.4 Observations controller

We still need to realize the last resource from Table 5.2 on page 43: the observations resource. The key to realize the resources is the regular expression in Listing 6.6 on page 67: the regular expression attach the observations controller under weather stations URIs. This means that the `country_code` and the weather stations id (`ws_id`) is available for the observations controller. The two values are used to fetch the weather station from which observations are fetched.

The resource expose `GET` and `POST` from the uniform interface manifested in `index` and `post`, shown in Listing 6.9. `index` is similar to the previous `index` methods, however, it fetches the weather station as the first thing in order to fetch the right observations. `post` creates or updates observations with the resource in the request. It uses the `update_or_insert_observation` of the weather station model which realizes First In, First Out for observations.

```

class ObservationsController(RestController):
    def index(self):
        '''Get list of observations.

        Query Args:
            limit: the number of observations to fetch.
            time_format: (ietf|iso) format of times. Default is
                ISO format.
        Returns:

```

```

        latest observations.
    '''
    ws_key_name = WeatherStation.key_name(self.kwargs['country_code'],
        self.kwargs['ws_id'])
    ws = dbutils.get_object_or_404(WeatherStation, key_name=ws_key_name)

    observations = ws.weather_data.order('-time')
    (...)

def post(self):
    '''POST create or update 1 weather observation to this
        weather station.

    Args POST:
        json weather data:
            time (required): isoformat time
            speed (not required)
            gust (not required)
            direction (not required)
    '''
    ws_key_name = WeatherStation.key_name(self.kwargs['country_code'],
        self.kwargs['ws_id'])
    ws = dbutils.get_object_or_404(WeatherStation, key_name=ws_key_name)
    ws.update_or_insert_observation(self.post_data)
    return HttpResponse("Weather observation created")

```

Listing 6.9: *Observations controller*

6.2.5 Forecast points controller

The forecast point controller, shown in Listing 6.10, realizes the resources in Table 5.3 on page 44. The only conceptually new method is `show`, which returns a JSON representation of concrete forecast points including its forecasts. `show` takes the following sequence of steps:

1. extracts the latitude and longitude from the URI;
2. finds the forecast point (or returns 404 if un-existing); and
3. creates a JSON representation and returns it.

```

class ForecastPointsController(RestController):

    def index(self):
        (...)
    def put(self, id):
        (...)
    def show(self, id):
        '''Show concrete forecast point.

    Args:
        id: id of forecast point
    Args GET:
        time_format: (ietf|iso) iso is standard

```

```

    '''
    key_name = _extract_key_name(id)
    forecast_point = dbutils.get_object_or_404(
        ForecastPoint, key_name=key_name)
    time_converter = wlttime.dt2IETF if self.request.GET.
        get('time_format') == 'ietf' else wlttime.dt2ISO
    repr = simplejson.dumps(forecast_point.to_json(
        include_forecasts=True, time_converter=
        time_converter))
    return http.HttpResponse(repr, content_type='
        application/json')

def _extract_key_name(id):
    '''Extract the lat and lon given as 10.0,10.0
    and return the key_name of the forecast point id'ed by the
    point.

    Returns:
        lat,lon: the latitude and longitude.
    '''
    latlng = id.split(',')
    lat = latlng[0]
    lon = latlng[1]
    return ForecastPoint.key_name(lat,lon)

```

Listing 6.10: *Forecast points controller*

6.3 Implementation of Views

The functionality of the Web Service is separated into the three layers of MVC. Creating the views of the application is, however, so simple that they are incorporated into the controllers. The controllers assembled a simple Python object, converted it to serialized JSON using simplejson, and returned that in the HTTP response.

6.4 Enhancing Performance

The performance of the Web Service is essential; applying explicit HTTP caching can significantly increase the performance. The We Love Wind Web Service take advantage of HTTP caching in most resources. In the following, we focus on how the `index` method in the Spot controller is altered to support caching.

Updates of the spots resources are infrequent. The spots resources are valid in caches until a user adds a new spot or modifies a spot, however, an inconsistent cached copy is not a significant problem, it should just be brought up to date relatively fast. The `index` method of the spots controller explicitly set the duration of caches to one day.

The application can avoid transferring data over the network even if the cache is expired by using **ETags**. A natural **ETag** for spots resources are the last time of modification of spot entities for the current page. The representation of the spots resource sent back to clients is always attached with an **ETag** of the last modification time. Additional requests includes the **ETag** which is validated by the controller and only upon **ETag** mismatch the representation is sent back. Upon match the representation is unmodified and the controller sends back an HTTP 304 to the client which then uses its cached representation. The altered `index` method

of the spots controller is shown in Listing 6.11. In Django HTTP headers are set directly in the response object.

```
class SpotsController(RestController):

    def index(self):
        geohash = self.request.GET.get('geohash')
        geohash_offset = self.request.GET.get('geohash_offset',
        geohash)

        if geohash:
            spots = Spot.get_near_geohash(geohash,
            geohash_offset)
        else:
            spots = db.Query(Spot).order('geohash').filter('
            geohash >=', geohash_offset)

        country_code = self.kwargs.get('country_code')
        if country_code:
            spots.filter('country_code =', country_code)

        entities = spots.fetch(PAGE_SIZE + 1)

        # very long ago, day 1
        last_modified = datetime.datetime.fromordinal(1)
        items = []
        for s in entities[:PAGE_SIZE]:
            items.append(s.to_json())
            if s.updated > last_modified:
                last_modified = s.updated

        if str(last_modified) == self.request.META.get("
        HTTP_IF_NONE_MATCH"):
            return http.HttpResponseNotModified()

        qs = None
        has_next_page = (len(entities) > PAGE_SIZE)

        if has_next_page:
            qs = self.request.GET.copy()
            qs['geohash_offset'] = entities[-1].geohash
        repr = {
            'next': '%s?%s' % (self.request.META['PATH_INFO'],
            qs.urlencode()) if has_next_page else '',
            'items': items
        }

        response = http.HttpResponse(simplejson.dumps(repr),
        mimetype='application/json')

        response['Cache-Control'] = 'public, max-age=86400'
        response['Etag'] = last_modified
        return response
```

Listing 6.11: *Caching spots controller*

The complexity is increased but the speedup is significant. Spots are instantly loaded in browsers on additional requests within 24 hours. ETag caching, however,

does not offer a significant speed up: the following is a small test with cURL and time.⁴

```
$ time curl -i "http://www.welovewind.com/api/spots/?1"
HTTP/1.1 200 OK
Etag: 2009-06-15 08:14:58.829396
Content-Type: application/json
Cache-Control: public, max-age=3600
Date: Wed, 24 Jun 2009 11:25:10 GMT
Server: Google Frontend
Transfer-Encoding: chunked
(...)

real    0m0.622s
user    0m0.003s
sys     0m0.001s
```

ETag caching is tested, by setting the `If-None-Match` header field to the value of the ETag.

```
$ time curl -i -H "If-None-Match:2009-06-15 08:14:58.829396" \
> "http://www.welovewind.com/api/spots/?2"
HTTP/1.1 304 Not Modified
Content-Type: text/html; charset=utf-8
Date: Wed, 24 Jun 2009 11:25:19 GMT
Server: Google Frontend

real    0m0.555s
user    0m0.001s
sys     0m0.004s
```

In the example, the server returns a HTTP 304 which means that the response has an empty body. The duration of the requests, however, are almost equal. To investigate the case we issued a series of tests, however, with the same result. The average of both are within 620 – 660 ms. We note that the test was issued with a different URI each time, bypassing intermediate caches and forcing both request to hit the server.

We conclude that ETag caching is not worth the trouble in terms of latency if the representations are small and if it is not essential to reduce the used bandwidth.

6.5 Summary

In this chapter, we described the design and implementation of a RESTful Web Service with Google App Engine/Django. The implementation was described in terms of the MVC style. Finally, we presented an example of how HTTP caching is applied in the Web Service to reduce latency.

⁴cURL is a tool to, e.g., issue GET requests to a specified URI. Headers are specified with `-H` and included in the output with `-i`

7

The DAIMI Forecast Web Service

In the last chapter we described the We Love Wind Web Service. That service included external calls to another Web Service: The DAIMI Forecast Web Service, found at

- <http://daimi.welovewind.com/>

There are important requirements in the application that demands an additional forecast Web Service:

- it should be possible to get forecasts for arbitrary locations; and,
- after creating new spots, the forecasts of the nearest forecast point should be in place immediately.

If the logic to get forecasts were placed on GAE, the requirements are trivial. The application is based on binary forecasts from NOAA. The decoding software is incompatible with GAE; the decoding of the binary forecast files demands that the decoding logic is placed elsewhere. It is infeasible to continuously update all forecast points – ca. 260000 points – therefore, the forecasts of points are only updated when relevant for surf spots, and that demands an extra service for forecasts. In essence, we are wrapping a modern Web Service around the binary forecasts files from NOAA.

In this chapter we describe the implementation of a RESTful weather service with Python/CGI.¹ The focus of this chapter is on the implementation of the Web Service; we postpone the description of retrieving and converting weather data to the next part.

The resource of the service is already defined: it is the forecast points resources described in Section 5.3 on page 43. Therefore, we jump directly to the implementation; again described in terms of the MVC style.

¹The Web Service is located at the department of computer science, Aarhus University; that left the author with two options for generating the dynamic output: PHP or CGI. To keep things consistent we implement the forecast service in a Python CGI script.

7.1 Model

The forecast weather service contains a single model: the forecast point model. The purpose of the model is to encapsulate the access to the weather forecasts located in files.

Structure

The model contains three properties: the latitude and longitude of the point, and the forecasts of the point.

Logic

The forecast point model is able to return a representation of itself in JSON format. The main method in the model is a class method that retrieves the current and all future forecast from the gfs files; since decoding forecasts from the forecast files takes about 0.5 seconds the method takes a maximum argument that reduces the number of returned forecasts.

7.2 Controller

The controller is implemented as a Python/CGI script. CGI is a standard that connects a server with any programming language. Responses are crafted by printing the headers of the HTTP response which must include a **content-type** header, the header is ended with an empty line and after that the program prints the body of the response.

Listing 7.1 shows the controller. **handle_request** is the entry method; it starts by extracting the parameters of the request. The values of those parameters are later given to an external shell program as arguments; therefore, the parameters are sanitized and parsed as floats avoiding security issues. If any error occur during this process the response signals the error with a HTTP 400 Bad Request. Secondly, the extracted parameters are given to the model which populates itself with forecast data according to the arguments. Finally, if caching does not apply the forecast point is printed in JSON format.

```
def handle_request():
    try:
        params = extract_params()
        if not params:
            print_intro()
            return
        forecasts = ForecastPoints(**params)
        if etag_match(forecasts):
            not_modified_304()
            return
        print_forecasts(forecasts)
    except (ValueError, IndexError):
        print_intro(400)

def extract_params():
    input = os.environ.get('REDIRECT_QUERY_STRING')
    lat, lon, maximum = None, None, 3 # maximum 3 if not specified
    if input:
        input = input.strip('/').split(',')
```

```

        lat = float(input[0])
        lon = float(input[1])
        if lat < -90.0 or lat > 90.0 or lon < -180.0 or lon >
            180.0:
            raise ValueError
        match = re.search(r'max=(\d+)', os.environ.get('
            REQUEST_URI'))
        if match:
            maximum = int(match.group(1))
        return {'lat':lat, 'lon':lon, 'maximum':maximum}
    else:
        return None

def print_intro(status=None):
    print 'Content-type: text/plain'
    if status:
        print 'Status: %s' % status
    print
    print 'Welcome to the DAIMI Weather Forecasts Service'
    print 'forecasts are accessed by ../forecasts/{lat},{lon}'
    print 'the number of returned forecast is specified by
        setting the max query parameter: ../?max=10'

def print_forecasts(forecasts):
    print 'Content-type: text/plain'
    print 'ETag: %s' % forecasts.last_modified()
    print
    print forecasts.to_json()

def not_modified_304():
    print 'Status: 304 Not Modified'
    print

def etag_match(forecasts):
    etag = os.environ.get('HTTP_IF_NONE_MATCH')
    if etag:
        if forecasts.last_modified() == etag:
            return True
    return False

handle_request()

```

Listing 7.1: *Forecast point CGI controller*

7.3 Summary

In this chapter, we described the implementation of a forecast Web Service. We designed two Web Services since NOAA forecasts was incompatible with GAE. At the time of implementation the GAE lacked support for cron jobs, offline processing, and Java. All the necessary components are now in place² making a future solution that is more clean and effective possible. Java is needed since there are limited options to decode GRIB2 files: the unidata Java decoder tools³, `wgrib2`,

²Offline processing of tasks is still experimental on GAE.

³<http://www.unidata.ucar.edu/software/decoders/>

and `pygrib2`⁴. The latter two are incompatible with GAE since they are based on C modules.

This ends the description of the Web Services (or the server-side) of the We Love Wind mashup. In the following chapters we turn to the clients of those Web Services.

⁴<http://pygrib2.googlecode.com/>

Part III

The Clients

*"It is a capital mistake to theorize
before one has data."*

Sherlock Holmes, Sir Arthur
Conan Doyle (1859 – 1930)

8

Weather data

8.1 Getting weather data

The resources incorporate weather data – forecasts and observations – from four external resources: DCA, DMI, WB, and NOAA. The weather data must be up to data when the user requests the information. None of the sources provide JSON(P) Web Services therefore we must either use the GAE as a proxy or save the data locally and periodically update the content. 7 April 2009 Google announced support for scheduling cron jobs¹, giving three options for keeping the data up to date: either update the data

1. by an internal cron job at the GAE;
2. on-demand upon user HTTP requests; or,
3. by an external cron job.

Internal cron jobs: The natural choice would be an internal cron job at the GAE, making the GAE responsible for consuming weather data from the external resources at regular intervals. The cron jobs at Google are limited to the same constraints as usual HTTP requests, e.g., the request duration limit. In addition, the number of jobs are limited to 20, and at minimum there must be one minute between individual cron jobs [?]. With hundreds of weather stations a naive cron job times out, and the changes done in the transaction are rolled back. Since it is possible to execute cron jobs every minute, a workaround is to limit the number of processed weather stations. Every data retrieval lasts about one second which means that about 20 resources can be processed per minute under the request duration limit. Weather stations are updated ca. every 20 minutes, which yields only 400 stations that can be updated on a continuously basis per cron jobs, and which makes this approach infeasible.

On-demand: By on-demand we mean weather data is updated just-in-time when the user needs the weather data. Because of the request duration limit the approach would include requests from JavaScript at the client side for every resource of interest, and the controller would update the GAE version of the weather data if it is

¹<http://code.google.com/p/googleappengine/issues/detail?id=6>

stale. The approach, however, includes a significant latency. Getting weather data must be a low latency operation making the approach infeasible.²

External cron jobs: The last option is letting external cron jobs pull weather data from all the weather resource and push the data to the GAE. The external updaters are tightly coupled to the GAE in a point-to-point solution with much redundant data transfer, however, it is the best of the worst at the moment.

8.2 Weather Observations

Unfortunately we have no Swiss Army Knife for the problem of pulling information from the different sources. The resources are diverse and in different formats; therefore, each case is handled on an individual basis.

8.2.1 Scraping DMI and DCA

The only option to get weather data from DCA and DMI is by scraping their Web pages for content. Web scraping is the process of programmatically retrieving content from a Web page meant for human consumption. The data at DCA is inaccessible as a Web Service, but DCA has agreed to let their pages be accessed by our Web Scraper in order to retrieve the data.

DMI prohibits incorporation of their data in other sites. The author has several times asked DMI for a collaboration, however, un-successful. We might be crossing ethical boundaries, but, we have decided to Web scrape their site, just like Google does.

Content in the case of DCA is located at sub-pages on a station and data type basis (speed, direction, and temperature are located at different pages). Content at DMI is located in sub-pages on a weather station basis.

All resources are invalid HTML and will cause an XML parser to fail. This is a case where BeautifulSoup³, a Python module to Web scrape Web pages, has a great force; to quote their Web page:

You didn't write that awful page. You're just trying to get some data out of it. Right now, you don't really care what HTML is supposed to look like.

Neither does this parser.

Scraping the latest speed, direction, temp, and time from DMI Aarhus, the code in Listing 8.1 is all that is needed (after downloading the BeautifulSoup module).

```
months = ['', 'januar', 'februar', 'marts', 'april', 'maj', 'juni', 'juli', 'august', 'september', 'oktober', 'november', 'december']

def _scrape(ws):
    page = urlopen(ws['uri_scrape']).read()
    soup = BeautifulSoup(page)

    # all text from page
```

²GAE is an emerging technology. A promising GAE offering (announced 18 June 2009) is the Task Queue. The Task Queue process tasks independently of the originating request. The only problem is getting the updated information back to the user, and not just to the next user accessing the page. Such a requirement is typically realized with Comet (or Ajax Push), however, this is not supported on GAE. If (or when) it is supported the Task Queue together with Comet is a solution for the update problem.

³<http://www.crummy.com/software/BeautifulSoup/>


```

text = soup.findAll(text=True)

# time
contents = text[6].split()
day = int(contents[2].strip('.'))
month = int(months.index(contents[3])) # januar -> 1
year = int(contents[4])
hour = int(contents[5].split(':')[0])
minute = int(contents[5].split(':')[1])
dt = datetime(year, month, day, hour, minute)
iso_time_utc = convert_local_time_to_utc(dt.timetuple())

# wind
contents = text[7].split()
direction = float(contents[1])
speed = float(contents[3])

# temp
contents = text[8].split()
temp = float(contents[1])

return {'time':iso_time_utc, 'direction':direction, 'temp':
        temp, 'speed':speed}

def convert_local_time_to_utc(time_tuple):
    secs = time.mktime(time_tuple)
    utc = time.gmtime(secs)
    iso = time.strftime("%Y-%m-%dT%H:%M:%S", utc)
    return iso

if __name__ == '__main__':
    print _scrape({'uri_scrape':'http://servlet.dmi.dk/bv/
                  servlet/bv?stat=6074'})

```

Listing 8.1: *DMI Web scraper*

In the example we use the Python library `urllib2` to fetch the page. This is given as an argument to `BeautifulSoup` when it is instantiated. The `BeautifulSoup` class has many methods that assist in extracting relevant data from the page; we use the `findAll` method to find all elements, and extract content from that list.

Scraping DCA's site is similar except that it needs tailoring for the specific layout of that site.

The Web scraper is fragile to modifications in the user interface, e.g., modification to the placement of the different data, but we have no other option to retrieve the data.

8.2.2 Scraping WB

In contrast to the ill-formed data from DCA the data from WB is in XML formats, and therefore more apt for computer consumption. Listing 8.2 shows an example of a WB XML document. The documents are subject to a single operation: extraction of node content and node attribute values.

From an abstract viewpoint, there are two types of XML parsers: tree-based parsers and event-based parsers. Tree-based parsers read in the whole XML document at once and create an internal data structure, while event-based parsers consider XML documents as a stream of events serially generating events while

parsing the document. In the previous section, we used a tree-based parser, BeautifulSoup. The advantages of event-based frameworks over tree-based are that they are fast and efficient since they avoid building a tree data structure; the advantage, however, is at the expense of simplicity in the programming model.

XML parsing of WB data merely has the purpose of extracting content of specific nodes: the application does not call for any transformations, and not for processing dependent on former seen nodes, and not for fixing invalid XML; therefore, we choose the event-based parser, SAX.⁴

```
(...)
<aws:station requestedID="EKAH" id="EKAH"
  name="Tirstrup Airport"
  city="Tirstrup" state="Denmark"
  citycode="60090" country="Denmark"
  latitude="56.2999992370605"
  longitude="10.6166667938232" />
<aws:current-condition
  icon="http://deskwx.weatherbug.com/images/Forecast/icons/
    cond024.gif">
  Mostly Cloudy
</aws:current-condition>
<aws:temp units="&deg;C">9.0</aws:temp>
<aws:rain-today units="mm">0.00</aws:rain-today>
<aws:wind-speed units="km/h">22</aws:wind-speed>
<aws:wind-direction>ESE</aws:wind-direction>
<aws:gust-speed units="km/h">21</aws:gust-speed>
<aws:gust-direction>ESE</aws:gust-direction>
</aws:weather>
(...)
<pubDate>Wed, 15 Apr 2009 13:50:00 GMT</pubDate>
(...)
```

Listing 8.2: *WeatherBug Weather RSS feed for weather station Tirstrup Airport*

Parsing with Python SAX is done by inheriting from the class `ContentHandler` and overwriting relevant methods. The methods correspond to events generated when SAX parses the document.

A generic content handler that extract content must overwrite three methods in the `ContentHandler`: `startElement`, `endElement`, and `characters`. During parsing the `startElement` method is called for every XML start tag, the `endElement` method is called for every XML end tag, and the `characters` method is called for each chunk of character data.⁷

Listing 8.3 shows the code for the XML parser. The parser is instantiated with a list of the type of XML elements from which to retrieve the content. The parser keeps track of when it is inside relevant nodes with the `in_element` variable. The relevant node content is put in a dictionary of XML-element keys and node content values.

```
from xml.sax.handler import ContentHandler

class XMLParser(ContentHandler):
    '''Simple XML Parser.
```

⁴The Python standard library contains both types of XML parsers: tree-based, manifested in the Document Object Model⁵ (DOM), and event-based, manifested in the Simple API for XML⁶ (SAX).

⁷For more info see <http://docs.python.org/library/xml.sax.handler.html>

```

        Extracts node content from the elements in the
        xmlelements list. These
        elements MUST NOT be nested in the XML document.
        Note: Duplicate named elements will result in the
        content of the
        last processed element.
        '''
def __init__(self, values={}, xmlelements=[]):
    '''
    Initializes instance of XMLParser.

    Args:
        values: dictionary where result is put.
        elements: list of XML element names to extract
        values from.
    '''
    ContentHandler.__init__(self)
    self.values = values
    self.xmlelements = xmlelements
    self.node_content = []
    self.in_element = False

def startElement(self, name, attrs):
    if name in self.xmlelements:
        self.in_element = True

def endElement(self, name):
    if self.in_element:
        self.values[name] = ''.join(self.node_content)
        self.node_content = []
        self.in_element = False

def characters(self, string):
    if self.in_element:
        self.node_content.append(string)

```

Listing 8.3: *SAX XML parser*

Extracting the content of the RSS feed is now just a matter of picking the elements to retrieve the node content from, and parsing the XML string with the content handler, the code is shown in Listing 8.4.

```

XML_ELEMENTS = ['aws:temp', 'aws:wind-speed', 'aws:wind-
    direction', 'aws:gust-speed', 'aws:gust-direction', 'pubDate'
]

def _scrape(ws):
    response = urllib2.urlopen(ws['uri_scrape']).read()
    values = {}
    parseString(response, XMLParser(values, XML_ELEMENTS))

    time = convert_local_time_to_utc(parse_wb_time(values['
        pubDate']))
    temp = parse_float(values.get('aws:temp'))
    gust = parse_float(values.get('aws:gust-speed'))
    direction = values.get('aws:wind-direction')
    speed = parse_float(values.get('aws:wind-speed'))

```

```

    if speed:
        speed = speed / 3.6 # from km/h to m/s

    if gust:
        gust = gust / 3.6

    if direction:
        direction = convert_direction(direction)

    return {'time':time, 'direction':direction, 'temp':temp, '
           speed':speed, 'gust':gust}

(...)

if __name__ == '__main__':
    print _scrape({'uri_scrape':'http://api.wxbug.net/
                  getLiveCompactWeatherRSS.aspx?ACode=A5580954882&
                  stationid=ETGG&unitttype=1'})

```

Listing 8.4: *Extracting XML content from WB station RSS feed*

8.2.3 Integrating with the Web Service

Web scraping the three types of weather stations is a similar process:

1. fetch all weather stations of a certain type from the Web Service; and for all weather stations,
2. scrape it and post the new data to the Web Service.

The only difference between the weather station types are in the concrete scraping code. Therefore, we apply the template method pattern [?] and create a general scraper skeleton, deferring the concrete scraping code to its subclasses.

`update_weather_data` is the template method, shown in Listing 8.5, that implements the process described above. `_scrape` is the deferred method, it uses the `scrape_uri` from the fetched representation and returns the scraped data as a simple Python object compatible with the observation representation format (see 5.3 on page 42). `_update_server` converts the object to JSON, sets the content type (used by the controller in the Web Service), and post the representation to the server. `_scrape` continues recursively if the paginator contains more pages. The individual `_scrape` methods were presented in the last section and fit directly into the skeleton when they are put in a class that inherits from `Scraper`.

```

URI_BASE = 'http://www.welovewind.com'
URI_WEATHER_STATIONS = 'http://www.welovewind.com/api/
                        weather_stations/'
class Scraper:

    def __init__(self, type):
        '''Craete instance of scraper
        Args:
            type: weather station type ('dmi'|'wb'|'dca')
        '''
        self.type = type

    def update(self):

```

```

        self.update_weather_data(URI_WEATHER_STATIONS + '?type
                                =%s' % self.type)

    def update_weather_data(self, uri):
        '''Template method'''
        wss = self._fetch_weather_stations(uri)
        for ws in wss['items']:
            data = self._scrape(ws)

            # update server
            uri = '%s%s' % (URI_BASE, ws['uri_observations'])
            self._update_server(uri, data)
        if wss['next']:
            self.update_weather_data('%s%s' % (URI_BASE, wss['
                next']))

    def _fetch_weather_stations(self, uri):
        response = urllib2.urlopen(uri)
        return simplejson.load(response)

    def _scrape(self, ws):
        raise Exception('Abstract Method')

    def _update_server(self, uri, data):
        try:
            json = simplejson.dumps(data)
            # putting in the values will cause the request to
            # be a POST
            request = urllib2.Request(uri, json, {'Content-Type':
                'application/json'})
            response = urllib2.urlopen(request).read()
        except urllib2.HTTPError, e:
            logging.error('trying to post data: %s' % data)

```

Listing 8.5: *Common Web scraper*

We use `cron` as the job scheduler. A table that runs the scraper each 20 minutes looks like the following:

```
0,20,40 * * * * /scraper.py
```

After installing the table with `crontab -e` in a Unix like environment weather observations are continuously updated for all the weather stations in the Web Service.

8.3 Forecasts

We stated earlier that there was no Swiss Army Knife for doing Web Scraping, this is even more true for retrieving the forecasts.

8.3.1 Retrieving forecasts

The forecasts are quite large about 45 MB for a single forecast. This means a total download of 45 MB * (48 hours / 3 hour interval) = 720 MB every 6 hours! 3 fields (temperature, u wind component, v wind component) out of 250 are interesting for the application. It is possible to reduce the download to interesting data only. This is done by first downloading a small – around 26KB – inventory of the forecast; it

specifies the content of the forecast and at which byte ranges the different content is located. This makes it possible to only fetch the parts we are interested in. A perl script that does this is available.⁸

A small shell script using perl scripts that downloads a subset of fields, now with a size under 1 Mb, looks as follows:

```
#!/bin/sh

url=http://nomad3.ncep.noaa.gov/pub/gfs/rotating-0.5/gfs.t00z.pgrb2f00
get_inv.pl "${url}.inv" | \
egrep "(:UGRD:10 m above|VGRD:10 m above|:TMP:surface)" | \
get_grib.pl "${url}" pgb.grb
```

The NOAA server contains forecasts from four calculation runs which are circularly updated as the forecasts calculations are finished. The application keeps a local version of all forecasts files, and also updates them circularly. Local forecasts are updated when they are stale by comparing the last modified time – found with a HTTP HEAD request – of the remote files with the local files.⁹ In the application we check the last file in each run (since it is produced by NOAA after the other files in the run) and update all forecasts in the run if the last is stale. The updating logic checks last modification times twice every hour by a cronjob.

8.3.2 Extracting data from GRIB2 files

`wgrib2`¹⁰ is a program that retrieves data from GRIB2 files. `wgrib2` supports extracting records out of a GRIB2 file based on latitude and longitude. We can type in the following command to get all data fields about Skagen from a GRIB2 file.

```
$ wgrib2 forecasts/gfs.t00z.pgrb2f00 -s -lon 10.59 57.7
1:115:d=2009062700:TMP:2 m above ground:anl::lon=10.500000,
   lat=57.500000,val=289.71
2.1:180652:d=2009062700:UGRD:10 m above ground:anl::lon=10.500000,
   lat=57.500000,val=-4.81
2.2:180652:d=2009062700:VGRD:10 m above ground:anl::lon=10.500000,
   lat=57.500000,val=-0.76
```

The `-lon` option gets values from a grid point closest to a specified point. The output shows the closest point, in this case it is (57.5,10.5), together with the value of the records.

The values in the GRIB2 files are in a meteorological format not applicable for the users of the Weather Web Service.

Temperature conversion

The temperature is reported in Kelvin; conversion: $[^{\circ}C] = [K] - 273.15$

Wind velocity conversion

The wind velocity forecast is a vector separated in two components `u` and `v`. `u` is the zonal – East-West – component; if positive the wind is blowing to the East. `v`

⁸http://www.cpc.noaa.gov/products/wesley/fast_downloading_grib.html

⁹We note that we have changed the perl scripts in order to preserve the external time when downloading the forecasts.

¹⁰`wgrib2` is found at: <http://www.cpc.noaa.gov/products/wesley/wgrib2/>

is the meridional – North-South – component; if positive the wind is blowing to the North.

Instead of the vectors the Web Service reports the wind speed and the meteorological wind direction; which is a value in the interval $[0;360[$ with respect to North (North=0), where the wind is coming from. The components are converted in Python, shown in Listing 8.6.

```
import math
def toSpeed(u_component, v_component):
    """Returns the norm of the vectors, i.d., the speed.
    """
    return math.hypot(u_component, v_component)

def toDirection(u_component, v_component):
    degrees = math.degrees(math.atan2(-u_component, -
        v_component))
    if(degrees < 0):
        degrees += 360
    return degrees
```

Listing 8.6: *Conversions in Python*

u and v are negated since meteorological wind direction is where the wind is coming from instead of blowing to. `atan2` returns an angle in the interval $[-\pi; \pi]$. To map this into $[0^\circ; 360^\circ]$ we must add 360° to negative angles.

8.3.3 Integrating with the Web Service

A central function for the updating process is retrieving all future forecasts from files in the latest run. The function finds the latest run, comparing the last forecast in each run. All file names in a run is available in a sorted list; the time delta from the calculation time to the current utc time is used to slice away stale forecasts from the list. `get_wind_forecast` uses `wgrib2` on forecast files and returns weather data found with regular expression on the output from `wgrib2`.

```
URI_BASE = 'http://www.welovewind.com'
URI_POINTS = 'http://www.welovewind.com/api/forecast_points/'

def update_server(uri):
    """PUT forecasts updates to server if the data at the
    server is stale.

    """
    forecast_points = _fetch_forecasts(uri)
    for p in forecast_points['items']:
        if not is_local_fresher(p):
            continue
        fp = forecasts.ForecastPoints(lat=p['lat'], lon=p['lon'],
            maximum=16)
        if p['lat'] != fp.lat or p['lon'] != fp.lon:
            raise Exception('lat/lon inconsistency')
        uri = '%s%s?_method=PUT' % (URI_BASE, p['uri'])
        representation = fp.to_json()
        request = urllib2.Request(uri, representation, {'
            Content-Type': 'application/json'})
        response = urllib2.urlopen(request).read()
    if forecast_points['next']:
```

```
        update_server('%s%s' % (URI_BASE, forecast_points['next  
        ']))  
  
def _fetch_forecasts(uri):  
    response = urllib2.urlopen('%s' % uri).read()  
    return simplejson.loads(response)
```

8.4 Summary

In this chapter, we presented and used different types of Web scraping techniques to fetch data from external resources not meant for computer consumption. In addition, we presented logic to load data into the application on a continuous basis.

*“...the goal is to transform data
into information and information
into insight.”*

Carly Fiorina (1954–)

9

Main user interface

Users are unlikely to appreciate the Web Services developed in the last part. They merely serve data; that data needs processing and a proper presentation in order to be converted to information relevant for the user. In this chapter, we describe the design and implementation subtleties regarding the main user interface of the mashup. The client is an Ajax client conveying surf weather information that assist practitioners of wind sports.

9.1 Design

An overall requirement for the user interface is conveying relevant information in a way that is easily comprehensible for the user. The key to fulfill this requirement is 1) to present the information in a graphical way, and 2) to reduce the information to the relevant only in terms of the location of the user. In essence, the main client is a geographical information system.

We embrace the requirements by a design consisting of three pages all visualizing data on a map. The map is centered around the location of its user which filters away most irrelevant information from the user interface. Surf spots are pivotal in the application and are available on all pages. Weather observations and weather forecasts are separated on two different pages due to their conceptual difference. Otherwise, forecasts and observations are displayed in the same way: an arrow in the wind direction and with a color indicating the speed. Spots are shown on all maps with a green circle with a cross in it. Finally, we have a page where spots are input or edited.

Figure 9.1 shows an early prototype of the user interface. The current application still has a design consisting of three pages. Page 1, however, shows an initial idea for visualizing: uniting spots and weather data. There exists no guarantee for a weather station near spots, and forecasts are conceptually inappropriate for flag illustration; therefore, spots and weather data is kept separate in the current design. Page 2 shows the observations on a map. Page 3 shows the page where users enter new spots in the application.




Figure 9.1: Paper prototype of main user interface

9.2 Finding the location of the client

A pivotal requirement for the weather service is delivering location-aware geographic information. This demands that the location of the client is known. Ideally, the user client tracks the user's location by any available means, e.g., IP-address, GPS, or GSM localization, sends the geographic information to the server that responds with a resource representation relevant for the location. W3C is working on the Geolocation API Specification that embrace the requirement [?]. The specification is novel but browsers vendors have begun supporting it in their beta versions, e.g., Opera¹ and Mozilla².

¹<http://my.opera.com/core/blog/geolocation-enabled-build>

²https://developer.mozilla.org/En/Using_geolocation

The API is obviously the future, however, the API is not yet supported to an extent where we would consider using it. At the moment a more viable approach is using the Google Gears Geolocation API³, which among others the W3C Geolocation API specification builds upon. Google Gears also support mobile devices, however, only a handful.

JavaScript that use Google Gears triggers a popup in browsers. The popup ask users to grant permission to the Web application in order to store information on their computer. The popup is intimidating for users and a sole reason for another solution.

`ipinfodb.com` is a free geolocation service. The service converts from IP address – uri-encoded in a GET request – to a geographical position returned in either XML or JSON(P). An example: a request such as

`http://ipinfodb.com/ip_query.php?output=json&callback=foo`

returns a JSON object wrapped in the `foo` function indicating the approximate location of the callers computer. Since the returned format is valid JavaScript requests are directly loaded from the source into client, it just have to be dynamically inserted into the `<script>` tag. Luckily, jQuery abstracts that away in the `getJSON` function.

Listing 9.1 shows how the We Love Wind client uses `ipinfodb`. The client avoids calling `ipinfodb` repeatedly by storing the result of the first request as a cookie that is valid for one day.⁴ `getLocation` is the entry function; when a location cookie is not present the client fetches its own location from `ipinfodb` using `geoIP` and JSON(P).

```
// get the lat/lon of this user ip address
function geoIP(callback){
    jQuery.getJSON("http://ipinfodb.com/ip_query.php?output=
    json&callback=?", function(data){
        var lat = data.Latitude;
        var lon = data.Longitude;
        callback({
            'lat': lat,
            'lon': lon
        });
    });
}

// get the lat/lng of the client
// if location cookie is set get this
function getLocation(callback){
    var location = readCookie('location');
    if (location) {
        var latlng = location.split(',');
        callback({
            'lat': latlng[0],
            'lon': latlng[1]
        });
    }
    else {
        geoIP(function(obj){
            createCookie('location', obj.lat + ',' + obj.lon,
                1); // 1 day
        });
    }
}
```

³<http://code.google.com/apis/gears/mobile.html>

⁴A thorough introduction to cookies in JavaScript is available at <http://www.quirksmode.org/js/cookies.html>

```

        callback(obj);
    });
}
}
// set location cookie
function setLocation(point) {
    createCookie('location', point.lat + ',' + point.lon,
        1);
}

```

Listing 9.1: Tracking client location with *ipinfodb* and *JSONP*

9.3 Reducing the data set

The size of the application's data set that the client downloads will eventually reach a significant size. A size that will hurt the client in terms of increased latency and poor performance. The geohash grid is a natural means to reduce the data set. We decrease the amount of data to download by reducing data to geographical points that are in the neighbors geohash area. Figure 9.2 shows the concept; only points within the boxes are downloaded.

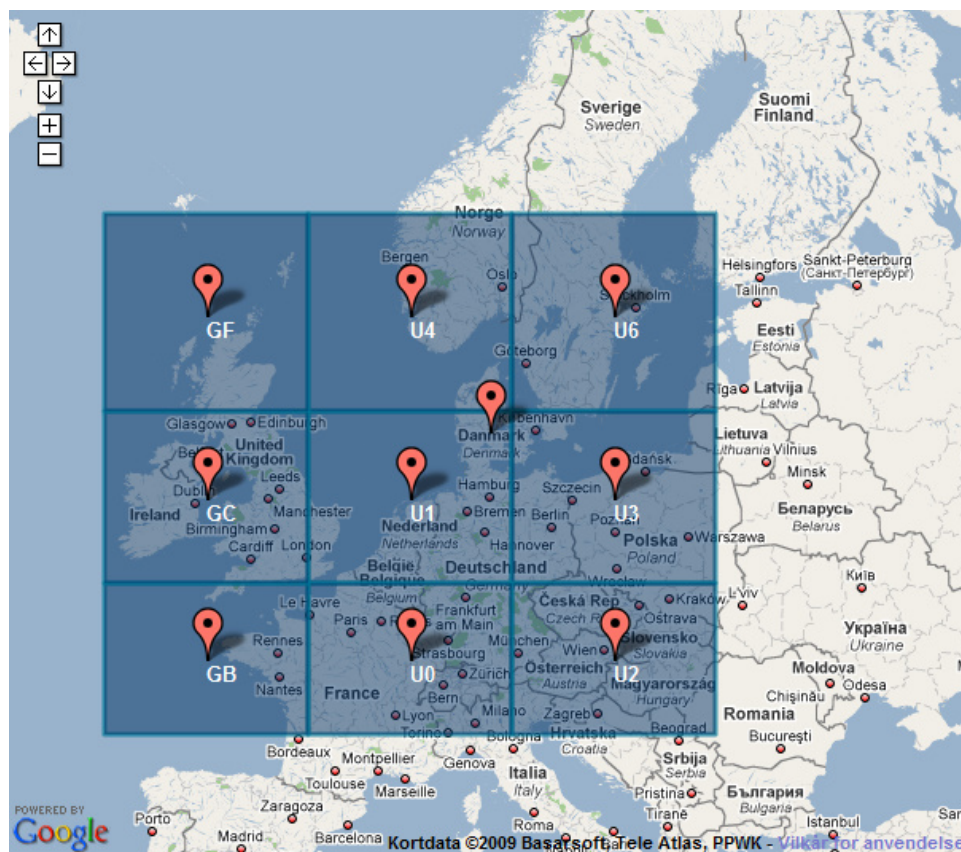


Figure 9.2: Geohash grid

The Web Service already supports geohash queries. However, the approach demands that geohash encoding and grid neighbor finding is also in place at the client. [?] is an open source geohash and grid neighbor finding implementation. We have

made the implementation compatible with most browsers (removing non-standard ECMAScript) and added logic to handle case three of navigating to neighbors (see Section 5.4.1) and included it in the client.

With geohash in place the client needs a transparent way of handling 1) downloading points of different types: spots, weather stations, and forecasts points, and 2) panning of the map which should trigger download of data from untouched grid boxes.

The two requirements is accomplished with the publish-subscribe pattern [?]. A geohash publisher is connected to the map subscribing to map movement finished events. Upon such events the geohash publisher encodes the geohash value of the new center, calculates all the neighbors, and publishes all new geohashes to the subscribers.

We encourage the reader to access <http://www.welovewind.com/examples/geohash/index.html> for a demonstration of the concept.

9.4 Seamless points loading

With the geohash publisher in place we turn to the problem of seamlessly downloading points in the published geohash grid cells and displaying the downloaded points on the map.

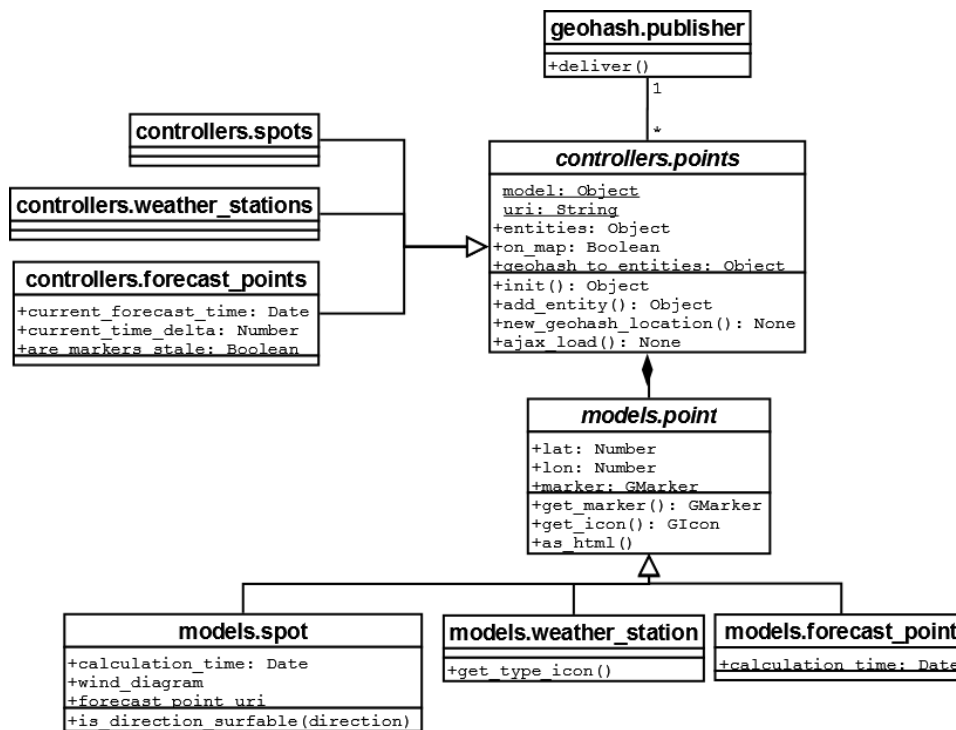


Figure 9.3: Client architecture view (UML static structure diagram)

All active geographical points are handled by a points controller, shown in Figure 9.3. The points controller contains logic to retrieve points from the server given a geohash. In addition, other objects (not included in the figure) can subscribe to the points controller, in which case the controller delivers downloaded points to them. The application currently has three different types of points: spots, weather stations, and forecast points that all inherit from the points controller. The “subclass” defines the URI of the concrete Web Service resource, and defines the model

for the concrete objects in the JSON items list (present in all JSON list resources). The model is used to augment the loaded JSON objects with logic, e.g, to display themselves on a map, and determine the surfable direction of a spot.

An example of a subscriber of the points controllers is the map controller. The map controller's responsibility is showing / removing categories of markers on the map. Figure 9.4 shows the runtime architecture of points loading, and their display on the map. First, all the publish-subscribe relationships are setup. In addition, a map movement finished subscriber is setup in the geohash publisher; the subscriber is the source that starts the following process:

1. Upon a map movement finished event the geohash for the center of the map is calculated.
2. The neighbors are found and subscribers (the points controllers) are delivered all new touched geohash grid cells.
3. The points controllers then contact the server for points in the geohash cells.
4. Upon return the points are de-serialized and augmented with logic. The augmented points are delivered to the subscribers, in this case, the map controller.
5. The map controller adds the points to the map.

The architecture decouples objects and separates concerns. An advantage is that any other controller interested in points can just add itself as a subscriber to the different type of points it is interested in.

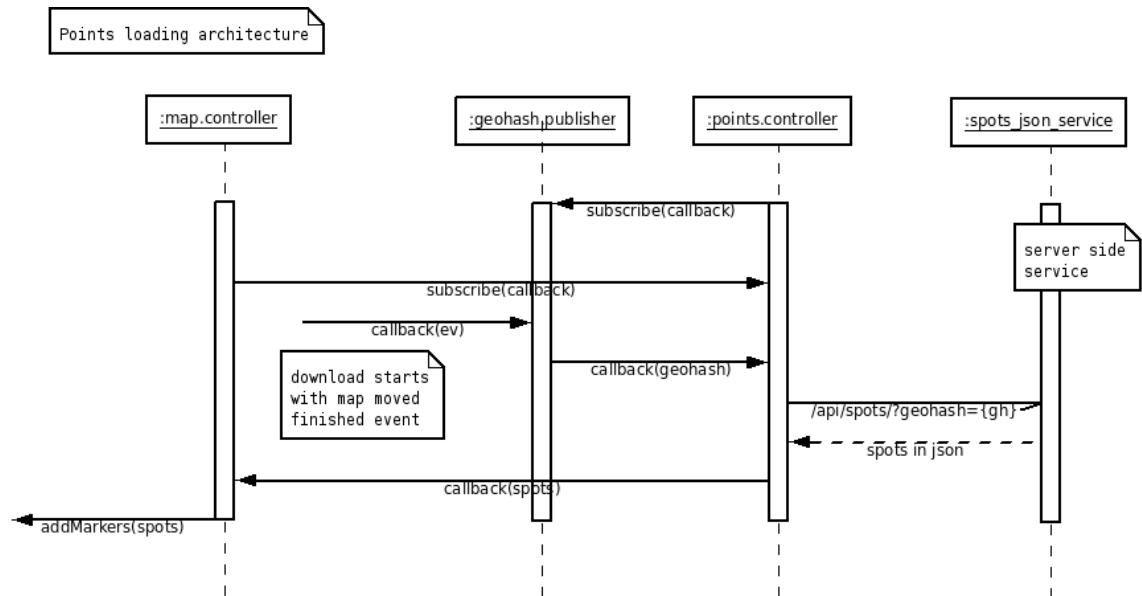


Figure 9.4: Points loading architecture view (UML sequence diagram)

9.5 Coordination between Ajax calls

A requirement for the application is visualizing relevant forecasts coupled to the surfability of spots. We have come up with a graph to visualize the current and

future surfability of spots, shown in Figure 9.5.⁵ An upward bar in the figure indicates that the wind is in the correct direction (input by the users), and downward direction means the wind is in the wrong direction. The colors correspond to the color map in Figure 1.2 on page 3, except that when the wind direction is unsuitable the color is black.




Figure 9.5: *Surfability graph*

The graph demands the presence of both the spot and its corresponding forecast point. With Ajax the ordering of returned calls is arbitrary. Thus, the client should somehow coordinate the dependent requests and first when both points are loaded create the graph.

Every points controller maintain a map from geohash to concrete point entities. Whenever points from a new geohash are loaded they are inserted into the map. When controllers subscribe to a points controller they specify a list of dependencies in the form of references to other points controllers. When points controllers publish updates to its subscribers the dependencies are first checked and only if all have downloaded points from the geohash the callback is triggered. Listing 9.2 shows the `publish` function in the common points controller that handles dependencies.

```
// publish new points loaded
// Args:
//   geohash_prefix: geohash prefix of loaded models
points.publish = function(geohash, the_points){
  // publish points to subscribers
  subscribers_loop: for (var i = 0; i < points.subscribers.
    length; i++) {
    // check if dependencies are loaded
    for (var j = 0; j < points.subscribers[i].dependencies.
      length; j++) {
      var controller = points.subscribers[i].dependencies[j];

      // if other not loaded continue with next subscriber
      if (!controller.geohash_to_entities[geohash]) {
        continue subscribers_loop;
      }
    }
    // callback(points)
    points.subscribers[i].callback({
      'geohash': geohash,
      'points': the_points
    });
  }
}
```

Listing 9.2: *Dependency aware publish function*

⁵The graph is created with the open source jQuery-based graph tool flot: <http://code.google.com/p/flot/>

9.6 Drawing a circle with Google Maps

The user interface creates surfability graphs for spots within a certain radius of the user's location. The user interface needs some component that indicates this radius to the user and let them adjust it. We apply a circle as the means to indicate to users the radius of graph creation.

The Google Maps API provides no built-in functionality to draw circles on a sphere. The API provides a `GPolygon` that can be used to create arbitrary geometry. There are examples of how to create a circle out of polygons⁶, however, none derive nor explain the theory of the circle creation. In [?] many geographical information systems formulas for JavaScript are presented, however, it contains no explanation. In the following, we derive the formula behind circle creation.




Figure 9.6: *Spherical triangle on sphere*

The first requirement is given a point, that will be the center of the circle, find another point in a certain distance, and bearing from the center point. The law of cosines, Theorem 5.1, can be used to find points in a certain distance and bearing from another point. First we find the latitude of that point.

Latitude

Given the spherical triangle in Figure 9.6, then

$$\cos(c) = \cos(a)\cos(d) + \sin(a)\sin(d)\cos(C) \quad (\text{law of cosines})$$

Let N be the North Pole, then the distances a and c is given by the latitudes P_0 and P_1 respectively. Therefore,

$$\cos\left(\frac{\pi}{2} - P_{1_{lat}}\right) = \cos\left(\frac{\pi}{2} - P_{0_{lat}}\right)\cos(d) + \sin\left(\frac{\pi}{2} - P_{0_{lat}}\right)\sin(d)\cos(C)$$

Since $\cos(\frac{\pi}{2} - \phi) = \sin(\phi)$,

$$\sin(P_{1_{lat}}) = \cos(P_{0_{lat}})\cos(d) + \sin(P_{0_{lat}})\sin(d)\cos(C)$$

⁶Google Maps circle example: http://maps.forum.nu/gm_sensitive_circle2.html

That is, given a point P_0 , another point P_1 in bearing C and distance d from P_0 , the latitude of the point P_1 is given by

$$P_{1_{lat}} = \text{asin}(\cos(P_{0_{lat}})\cos(d) + \cos(P_{0_{lon}})\sin(d)\cos(C))$$

Longitude

From the last subsection the latitude of the new point is now known. Given that the angle B on the figure is known, the longitude of the new point, $P_{1_{lon}}$ is given by

$$P_{1_{lon}} = P_{0_{lon}} + B \quad (9.1)$$

We now derive a formula for the angle B . Given the spherical triangle in Figure 9.6, then

$$\cos(d) = \cos(a)\cos(c) + \sin(a)\sin(c)\cos(B)$$

Let N be the North Pole, then the distances a and c is given by the latitudes of P_1 and P_0 respectively. Since $\cos(\frac{\pi}{2} - \phi) = \sin(\phi)$,

$$\cos(d) = \sin(P_{0_{lat}})\sin(P_{1_{lat}}) + \cos(P_{0_{lon}})\cos(P_{1_{lon}})\cos(B)$$

Isolating B ,

$$\begin{aligned} \cos(B) &= \frac{\cos(d) - \sin(P_{0_{lat}})\sin(P_{1_{lat}})}{\cos(P_{0_{lon}})\cos(P_{1_{lon}})} \\ B &= \text{acos}\left(\frac{\cos(d) - \sin(P_{0_{lat}})\sin(P_{1_{lat}})}{\cos(P_{0_{lon}})\cos(P_{1_{lon}})}\right) \end{aligned}$$

The value for B can now be inserted in (9.1). Since acos returns the value in $[0; \pi]$ we have to handle the case when the bearing is above π . In that case,

$$P_{1_{lon}} = P_{0_{lon}} - B \quad (9.2)$$

The very observant reader might have noticed something peculiar: the formula derived does not match with the formulas used in the referred examples and [?]. Deriving the formulas used in those are much more complicated and since they lack explanation it remains uncertain why they did it that way.

From a mathematical perspective an interesting point is that the whole theory, both for the distance formulas and the circle formulas breaks down when one of the points is the North Pole. In that case the “triangle” is a line. In practice, however, it is not a problem since on Google Maps it is not possible to get a point at 90° latitude. It returns ca. 89.6°.

JavaScript implementation

With the formulas in place it is easy to implement the function that finds a new point in a certain bearing and distance from an originating point, Listing 9.3 shows how. The function directly reflects the formulas; because of small rounding errors $\cos(B)$ might be larger than 1, in that case we set it to one to avoid invalid input to acos .

```
// Calculate a new point in the given bearing and distance
// from the given point
// Args:
//     point: GLatLng to calculate from
//     bearing: bearing in radians
//     distance: distance to point
function destination(point, bearing, distance){
```

```

var R = 6367; // earth's mean radius in km
var d = distance / R; // unit sphere distance
var P0_lon = point.lngRadians();
var P0_lat = point.latRadians();

var P1_lat = Math.asin(
    Math.sin(P0_lat) * Math.cos(d) +
    Math.cos(P0_lat) * Math.sin(d) * Math.cos(bearing)
);

var cos_B = (Math.cos(d) - Math.sin(P0_lat) * Math.sin(
    P1_lat)) /
    (Math.cos(P0_lat) * Math.cos(P1_lat));

// secure against rounding errors.
if(cos_B >= 1){
    cos_B = 1;
}

if(bearing > Math.PI) {
    var P1_lon = P0_lon - Math.acos(cos_B);
}else{
    var P1_lon = P0_lon + Math.acos(cos_B);
}
// convert to degrees
P1_lat = P1_lat * (180 / Math.PI);
P1_lon = P1_lon * (180 / Math.PI);
return new GLatLng(P1_lat, P1_lon);
};

```

Listing 9.3: *Finding a point in a certain bearing and distance with JavaScript*

Creating the circle on Google Maps is done by iterating from $[0; 2\pi]$ creating a list of points around the center with **destination** and connecting them in a **GPolygon**.

9.7 Summary

In this chapter we presented the design of the Ajax client and the main subtleties in the implementation of the client. The subtleties included 1) loading points on a proximity basis, 2) loading additional points as users pans a Google Map, 3) coordination of points loading, and 4) drawing a circle with Google Maps.

*"Focus on people – their lives,
their work, their dreams."*

Google

10

Interface for mobile devices

In a context where the users of the application would rather be out surfing than sitting home by their computer a mobile interface is essential. The application's interface for mobile devices lets users access the weather service while on the go. In this chapter, we describe the design and implementation of such an interface for mobile devices, available at <http://m.welovewind.com/>.

In theory the mobile application could be situated at another domain, requesting the data from the Web service over the network. However, since it is convenient and faster to directly access the data models of the application we couple the mobile interface to the Python data models, already presented.

10.1 Design

Most mobile devices have small screens, reduced computing power, reduced input capabilities, and in general fewer functionalities than a standard computer. The limitations mean that most mobile browsers are unable to render the standard pages of the application since they rely on JavaScript; in addition, if they could, the network latency and cost due to the size of the pages would be significant. Therefore, the mobile version of the weather service calls for a complete redesign and restructuring of the content, however, the overall requirement is still the same: to present weather information that assist practitioners of wind sports.

Figure 10.1 shows an early manifestation of the mobile user interface: a user interface paper prototype. The prototype handles Scenario 1 and Scenario 2. Page number 1 on the figure, is a form where the user must input his location by name. Since the user most likely need information about the same location on the next visit, the user's last entered location is automatically set in the form.

Users arrive at page 1 also when accessing the main URI with a mobile device. Users therefore are relieved of 1) guessing the URI of the mobile page, or if that is unsuccessful 2) using a search engine to find it.

The main content page of the mobile interface incorporates a concise version of interesting location-based information, in a wind sports context. The page shows spots in the proximity, weather observations from weather stations in the proximity, and the weather forecast from the nearest forecast point.

Page number 2 gives an all information about the current weather conditions.

Page 4 is a map that shows spots and weather stations in the proximity of the user. The map is important when the user needs to find the way to a new spot. Page 3 and 5 are pages that show additional information about spots and weather stations respectively.

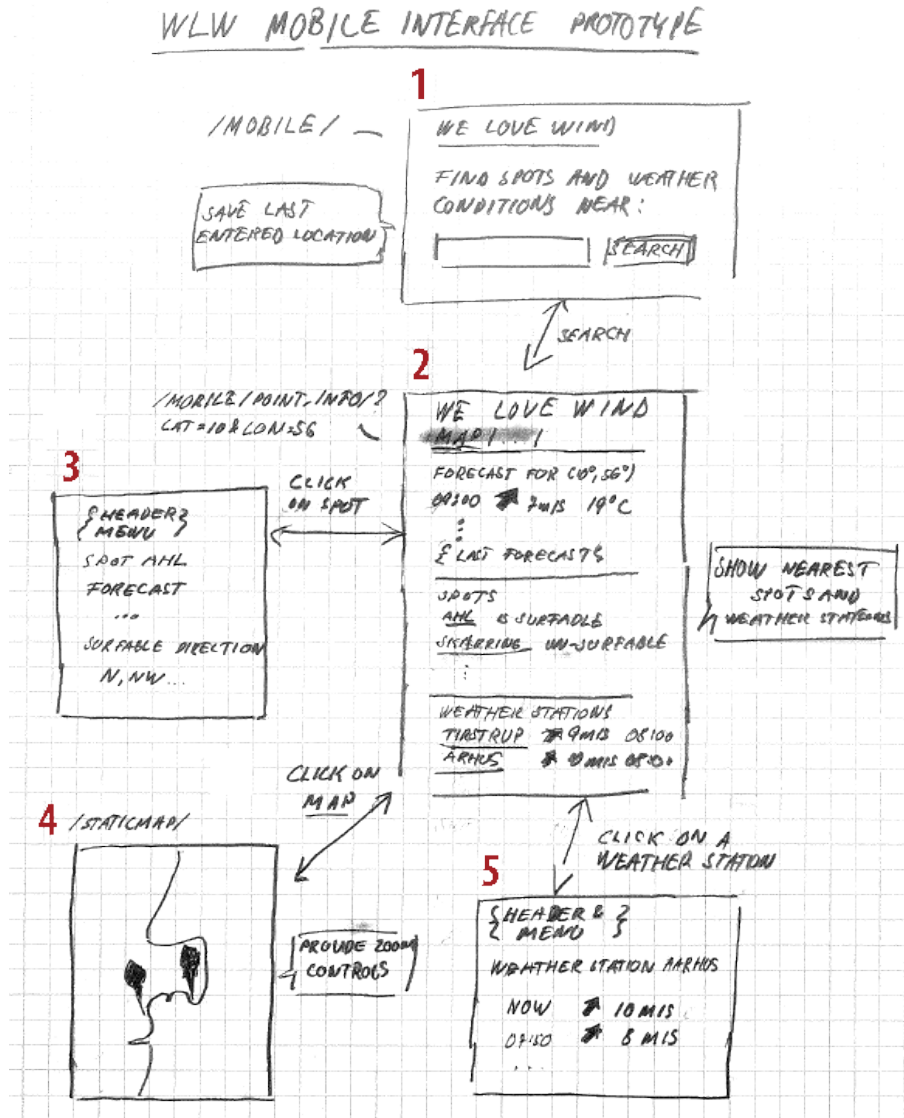


Figure 10.1: Paper prototype of mobile interface

The design discussed in the previous section adds four resources to the application; the resources are shown in Table 10.1.

10.2 Sorting out the resource requirements

In the last section we investigated the requirements and pointed out the additional resources the requirements of the mobile client added to the Web Service. Some of the resource content requires somewhat non-trivial processing, e.g., detecting a mobile client, displaying static maps apt for mobiles without JavaScript, and converting times to a user relevant timezone. This section reviews our solutions.

Resource URI	Query Parameters	Action
/mobile/		GET search page
/mobile/geo/	location name	GET geocoded location name
/mobile/info/	lat latitude, lon longitude, address address	GET all information for point
/mobile/map/	clat latitude of center, clon longitude of center, zoom zoom level of map, spots list of spot keys, wss list of weather station keys	GET map with spots and weather stations centered around the given point

Table 10.1: *Resource view of mobile resources*

10.2.1 Location-based mobile services

Making a single location-based solution accessible for all browser clients including mobile ones is impossible at the moment.

More low level solutions like [?] exists, however, they rely on technologies that are not ubiquitous either. We end up with a solution that rely on mobile browsers rendering simple HTML pages with no JavaScript. More advanced solutions could be developed, but the requirement is a widely accessible solution. In addition, some of the advanced phones, e.g., iphone, are capable of presenting the main page, see Figure 10.2. That is the mobile service is a lowest common denominator solution.

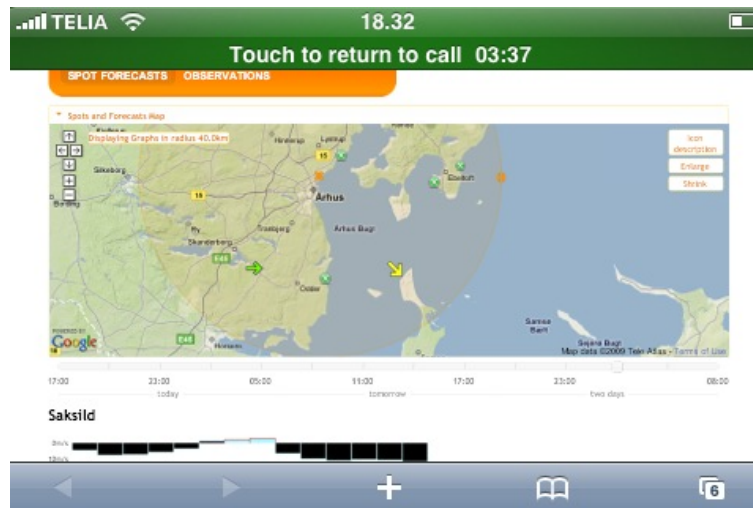


Figure 10.2: *Main page on iphone*

10.2.2 Content Negotiation

Content negotiation is a term defined in [?, sec.12] which is a mechanism to serve a representation of a resource based on the capabilities of the client. In case of the application, users accessing the site with a mobile should be served representations apt for mobile devices. All browsers populate the user agent field in HTTP headers; this can be used to decide if the user is accessing the site from a mobile device.

```

# special handling of theese
import logging
user_agents = ['ipod', 'android', 'opera mini', (...)]
# mobile user agents trimmed
more_user_agents = ['1207', '3gso', '4thp', '501i', (...)]

def mobile_device(request):
    '''Mobile device detection.

    Based on:
        http://detectmobilebrowsers.mobi/
    Args:
        request: HTTP request
    Returns:
        bool indicating whether the user agent is a mobile
        device or not.
    '''
    # e.g., User-Agent: Opera/9.60 (J2ME/MIDP; Opera Mini
    # /4.2.13918/422; U; da)
    user_agent = request.META['HTTP_USER_AGENT'].lower()
    logging.info('mobile.mobile_device: user_agent: %s' %
                 user_agent)
    if 'iphone' in user_agent:
        return False
    for ua in user_agents:
        if user_agent.find(ua) != -1:
            return True
    if user_agent[0:4] in more_user_agents:
        return True
    return False

```

Listing 10.1: *Mobile user agent detection*

Listing 10.1 shows the implementation, inspired by a PHP implementation by [?], of user agent detection. The code contains a list of mobile user agents and detects whether the user agent of the request is one of them.

When the main page of the application is accessed the content negotiation is performed. Listing 10.2 shows how the root controller take advantage of the content negotiation.

HTTP caching is used to bypass the content negotiation on additional accesses to the page from the same type of user agent.

```

from wlv.helpers import mobile

from django import http

def index(request):
    '''Render the index page'''

    if mobile.mobile_device(request):
        response = http.HttpResponseRedirect('/mobile/')
    else:
        response = http.HttpResponseRedirect('/observations/')
        response['Vary'] = 'User-Agent' # content varies by
            user-agent
        response['Cache-Control'] = 'public, max-age=604800' #
            cache for one week

```

```
return response
```

Listing 10.2: *Content negotiating controller*

10.2.3 Local time

In the Ajax client, JavaScript converted the UTC times of weather data served from the Web Service to local times using the timezone of the user's machine. The mobile interface must work without JavaScript demanding the time conversion to be done at the server. There are three ways to realize the requirement for serving time in a sensible timezone format:

- let the user input timezone and save it as a cookie;
- associate user accounts with timezone information; or
- report times in the timezone where it applies.

The first and second demands the user to input timezone; whereas, the last approach totally relieves the user for any input regarding the issue. Picking the last solution, however, puts forth additional requirements of storing timezone with all weather data owners: weather stations and forecast points.

In Chapter 6 the Forecast controller and Spot controller retrieved timezone information from an external Web Service: <http://www.geonames.org/>; it is now evident why: namely to realize option number three.

GeoNames is a free database containing extensive geographical information. GeoNames provide access to its resources through different mainly RESTful Web Services, and serves representations in XML and JSON formats. The applied service is the timezone service.

The implementation of code that retrieves a timezone from a point from GeoNames is shown in Listing 10.3. The code uri-encodes the point, fetches the resource, and handles different failures.

```
def geonames_timezone(point):
    '''Get timezone information with geonames.org.

    Returns:
        json string {'time':..., 'countryName': ..., '
                    countryCode':..., 'timezoneid': ...}
    '''
    values = {'lat':point.lat, 'lng':point.lon}
    query = http.urlencode(values)
    timezone_uri = 'http://ws.geonames.org/timezoneJSON?%s' %
        query
    try:
        response = urlfetch.fetch(url=timezone_uri)
        if response.status_code != 200:
            raise WSEException('geonames.org not available')
        if 'message' in response.content:
            json = simplejson.decode(response.content)
            raise WSEException('ws.geonames.org: %s' % json['
                status']['message'] )
        return response
    except DownloadError, e:
        raise WSEException('geonames.org not available')
```

Listing 10.3: *Using GeoNames timezone Web Service*

With the timezone in place the times are converted with the logic presented in Section 4.5.3.

10.2.4 Geocoding

Geocoding is applied in order to map from the user input location names to a latitude and longitude location. There are several geocoding services: Yahoo¹, GeoNames², and [?].

GeoNames is free, however, the author has experienced low availability of the Web service. Yahoo, restricts the number of geocoding requests to 5000 pr. day pr. IP address. Google restricts their service to uses that includes displaying markers on a Google map. The Google geocoding service is chosen, since the mobile interface presents markers on a Google map.

Listing 10.4 shows how the service is put to use from GAE. The code takes the following sequence of steps:

- Builds an URI to the Google Web service to query.
- Handles different failures, and throw a custom exception if so.
- Decodes and inserts the JSON response into a data structure.
- Uses memcache to bypass the geocoding service, and speed up additional requests for the same location.

```
def google_geocode(location):
    '''Geocode by location name.
    Args:
        location: name of location.
    Returns:
        location info as dictionary {'lat':..., 'lon':..., '
        address':...}
    '''
    data = memcache.get('/geo/location/%s' % location)
    if data is not None:
        return data
    else:
        values = {'key': settings.GOOGLE_APP_ID, 'q': location,
                  'output': 'json', 'oe': 'utf8', 'sensor': 'false'}
        query = http.urlencode(values)
        geocoding_uri = 'http://maps.google.com/maps/geo?%s' %
            query
        response = urlfetch.fetch(url=geocoding_uri)
        if response.status_code != 200:
            raise WSEException('Error finding location')
        json = simplejson.loads(response.content)
        status_code = json['Status']['code']
        if status_code >= 500:
            raise WSEException('Could not find location')
        placemark = json['Placemark'][0]
        lat = float(placemark['Point']['coordinates'][1])
        lon = float(placemark['Point']['coordinates'][0])
        address = placemark['address']

        data = {'lat': lat, 'lon': lon, 'address': address}
```

¹<http://local.yahooapis.com/MapsService/V1/geocode>

²<http://www.geonames.org/export/geonames-search.html>


```

if lat and lon and address:
    memcache.add('/geo/location/%s' % location, data)
return data

```

Listing 10.4: *Google geocoding in Python*

10.2.5 Displaying static Google maps

Google Static Maps³ is a Web Service that takes an HTTP request generates a static map images and responds with the image. The service is useful in a context where JavaScript is unavailable such as our mobile service.

It is easy to put the Static Maps Web Service to use, e.g., retrieving a map of Denmark is done by issuing a request to the URI:

```

http://maps.google.com/staticmap?center=56,10&zoom=6&size=512x512&
maptype=mobile&sensor=false

```

Since the representation format is an image format the URI can be placed in a usual `img` tag in HTML documents.

```

4</sup>

```

def info(request):
 '''Info about a point.

 Retrieve spots, weather stations in the proximity, and the
 forecasts for the point.

 Args GET:
 lat: latitude of point
 lon: longitude of point
 address: address of point
 '''
 lat = float(request.GET.get('lat'))

```

<sup>4</sup><http://www.welovewind.com/examples/geohash/index.html>

```

lon = float(request.GET.get('lon'))
if lat is None or lon is None:
 http.HttpResponseBadRequest('lat/lon must be set')
address = request.GET.get('address', '(%s,%s)' % (lat, lon))

center = db.GeoPt(lat=lat, lon=lon)

wss = WeatherStation.get_by_point(center, 50, 3)
spots = Spot.get_by_point(center, 50, 3)
forecast_point = ForecastPoint.get_or_init_by_point(center)

spot_keys = [s.key().name() for s in spots]
wss_keys = [w.key().name() for w in wss]

the 'list' is just a key string separated by commas in
the uri.
map_values = {
 'clat': lat,
 'clon': lon,
 'spots': ';'.join(spot_keys),
 'wss': ';'.join(wss_keys),
}

map_uri = '/mobile/map/?%s' % urlencode(map_values)

return shortcuts.render_to_response(
 'info_mobile.xhtml',
 {
 'map_uri': map_uri,
 'spots': spots,
 'wss': wss,
 'fp': forecast_point,
 'qs': request.GET.urlencode(),
 'address': address
 })

```

### 10.3.4 Views

All the views of the mobile interface are valid XHTML Basic 1.1: the W3C recommendation for content the format for mobile interfaces [?]. We skip presenting the views of the mobile application. The interested reader can refer to Appendix B for the views.

## 10.4 Summary

In this chapter, we extended the application with a mobile interface, to support scenarios where the user is on the go. The mobile interface had a significant impact on the application in terms of complexity. The added complexity stems from 1) a requirement for the mobile application being widely accessible, and 2) the limit of many mobile browsers. The application could not rely on JavaScript, and therefore functionality already developed in JavaScript, such as timezone conversion and display of maps, was implemented on the server side also.



## Part IV

# Conclusion & Perspectives





Through three parts in this dissertation, The Foundations, The Web Services, and The Clients, we have reached the destination, a “Web-Based Weather Service for Wind Sports”. The integration of these three parts is a mashup that assists the practitioners of wind sports.

Our dissertation gives an insight into the course we took to 1) extract data from external resources (independent of format) and, to 2) create a geographical information system mashup based on that data. The insight includes:

- The theoretical foundations for creating mashups (Chapter 2). The foundations included a presentation of the architectural style REST, the architecture ROA upon which we based our Web Services, and an infrastructure to run the mashups: the GAE.
- Practices for designing the resources in a Web Services (Chapter 5) and afterwards practices for implementation of mashups; in our application, manifested in two Web Services: the We Love Wind Web Service (Chapter 6) and the DAIMI Forecast Web Service (Chapter 7).
- At last the insight includes the theoretical foundations for creating an Ajax client (Chapter 3) and a mobile client. In addition, concrete practices for implementing such clients (Chapter 9 and Chapter 10)

Our thesis was only possible because we had access to public weather information from the US that use the ‘open access’ model [?] for public sector information. The European model, known as the ‘cost recovery’ model, restricts access to public data. During the last half year there have been European initiatives that argue in favor for open access to public data. `digitaliser.dk`, created by the Danish government, is one of the front-runners, quoting their Web site:

Digitaliser.dk aims to stimulate development and adoption of digital content and business models by utilising Web 2.0 technologies and public data and digital resources.

`digitaliser.dk` aspires to create common grounds of how to get access to public sector information. Times are thus changing in favor of mashups.



## Part V

# Appendix & Glossary



# Glossary

|                        |                                                                                                                                                                                                                                                                            |
|------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| active record          | Design pattern describing a solution to the problem of programmatic interaction with data in a database. In active record interaction is handled seamlessly by methods, such as save and delete, on objects that corresponds to entries in the database.                   |
| anti-patterns          | An ineffective practice commonly used.                                                                                                                                                                                                                                     |
| application state      | REST sees a Web application as a state machine that transitions to different states by following links and submitting forms in hypermedia. Application state is the current state of the application given by the client's sequence of followed links and submitted forms. |
| Base32                 | An alphabet based on 32 symbols.                                                                                                                                                                                                                                           |
| Big Web Services       | Term coined in [?, ch.10], for Web Services architecture based on SOAP, WSDL, and the WS-* stack.                                                                                                                                                                          |
| business logic         | The part of the application that defines the data model, the data storage, the data retrieval, and functions that retrieve certain data from the models.                                                                                                                   |
| CGI                    | A standard that connects a Web server with any programming language.                                                                                                                                                                                                       |
| cloud                  | A term for the Internet often used when describing services on the Internet where most of the infrastructure is hidden. An example is Google App Engine that hides all server details from its users.                                                                      |
| content negotiation    | HTTP mechanism that serves a certain representation of a resource based on the capabilities of the client [?, Sec. 12]                                                                                                                                                     |
| cross-site scripting   | Injection of HTML and / or JavaScript into a Web page.                                                                                                                                                                                                                     |
| de facto               | A general known but undeclared fact.                                                                                                                                                                                                                                       |
| distributed hypermedia | A non-linear media that integrates other media besides text, that can be located at different places.                                                                                                                                                                      |
| DOM                    | Tree representation format of XML documents, often used from JavaScript to access and modify a HTML document.                                                                                                                                                              |

|                          |                                                                                                                                                                                                   |
|--------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| First In, First Out      | A way of storing or processing of data where the first input goes out first.                                                                                                                      |
| geocoding                | Process of converting a location name into a latitude and longitude pair.                                                                                                                         |
| Google accounts          | Service that provides users access to several Google services.                                                                                                                                    |
| location-based service   | A service that takes the location of the user into consideration.                                                                                                                                 |
| meridian                 | A half great circle linking the North Pole and the South Pole.                                                                                                                                    |
| pagination               | Process of splitting up entities to several Web pages.                                                                                                                                            |
| quality attributes       | A measurable physical or abstract property of quality.                                                                                                                                            |
| request-response pattern | Architectural pattern for message exchange where every request is matched with a response.                                                                                                        |
| SOAP                     | A wrapper around XML documents used when transferring XML documents.                                                                                                                              |
| time to live             | The period of time some data is valid.                                                                                                                                                            |
| URI                      | Acronym for Uniform Resource Identifier: scheme to identify resource on the Web                                                                                                                   |
| WS-* stack               | A set of Big Web Service descriptions.                                                                                                                                                            |
| WSDL                     | An XML language to describe the interface of a Big Web Service.                                                                                                                                   |
| WSGI                     | A specification for the interface between Web servers and Python Web application frameworks. The rationale of the specification is to decouple the Web application framework from the Web server. |

# A

## Appendix A

### A.1 Rejseplanen.dk screenshots

When polite search engines, like Googlebot, access pages they first check the rules in the server's `robots.txt`. In Figure A.1 <http://www.rejseplanen.dk/robots.txt> is shown; all crawling is disabled, since `Disallow` is set to `/` for all agents.

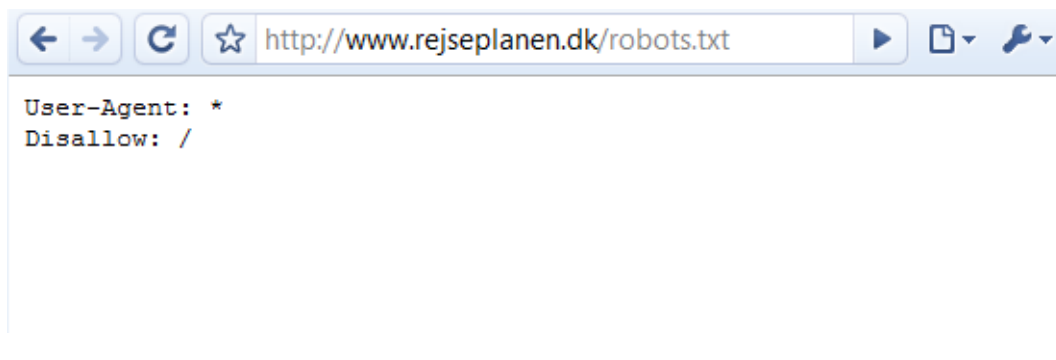


Figure A.1: `robots.txt` at *www.rejseplanen.dk*; [accessed 11-March-2009]

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Frameset//EN">
<html>
 <head>
 <meta http-equiv="Content-Type" content="text/html; charset
 =ISO-8859-1">
 <title>Rejseplanen</title>
 </head>
 <frameset rows="100%">
 <frame name="rejseplanen" src="/bin/query.exe/mn">
 <noframes>
 <body>
 Videre til
 Rejseplanen. Læs mere om
```

```

 rejseplanl&#
 230;gning med kollektiv trafik.
 </body>
 </noframes>
 </frameset>
</html>

```

Listing A.1: <http://www.rejseplanen.dk/>

## A.2 Cross-site scripting with eval

Imagine a site where user entered JSON data is de-serialized using the JavaScript `eval` function<sup>1</sup>, e.g., a site with user entered surf spots.

Listing A.2 shows how executing `eval` on the user input – hardcoded in the site in this example – executes injected JavaScript. When the data is ealed, the entered data is executed as if it were JavaScript. This is used in Listing A.2 to inject an extra field in the ealed data structure. The interesting line is this:

```

{
 "name": "Ahl", attack: alert('attack'), after:"
}

```

Because the value in the `attack` field is not surrounded by `""`, the command `alert()` is executed. We notice that this line comes from entering the following as surf spot name:

```
Ahl'', attack: alert('attack'), after:''
```

Since it is valid JavaScript syntax, but not JSON syntax the attack is realizable only if the response is not parsed as JSON. The attack can be tried out in action by going to <http://welovewind.appspot.com/examples/xss.html>.

```

<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
 "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
 <head>
 <meta http-equiv="content-type" content="text/html; charset
 =utf-8"/>
 <title>JSON eval() issues</title>

 <script type="text/javascript">
 //
 function run(id) {
 json = document.getElementById(id).
 innerHTML;
 res = eval('(' + json + ')');
 return res;
 }
 //]]>
 </script>
 </head>
 <body>
 <h1>JSON eval() issues</h1>
 <h2>Example</h2>
 </body>
</html>
</pre>
</div>
<div data-bbox="201 915 785 941" data-label="Footnote">
<p>¹https://developer.mozilla.org/en/Core_JavaScript_1.5_Reference/Global_Functions/eval</p>
</div>
<div data-bbox="481 957 514 973" data-label="Page-Footer">
<p>124</p>
</div>
```



```
 <pre id="ex1">
{
 "name": "Ahl", attack: alert('attack'), after:""
}
 </pre>
 <button onclick="run('ex1 ');">eval it!</button>
 </body>
</html>
```

**Listing A.2:** *xss.html*



# B

## Appendix B

### B.1 Implementation of Views for Mobile Devices

#### B.1.1 Base mobile template

```
<?xml version='1.0' encoding='UTF-8'?>
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML Basic 1.1//EN" "http
 ://www.w3.org/TR/xhtml-basic/xhtml-basic11.dtd">
<html xmlns='http://www.w3.org/1999/xhtml'>
 <head>
 <title>We Love Wind (Mobile)</title>
 </head>
 <body>
 <h2>We Love Wind</h2>
 {% block nav %}
 <div id="navigation">
 Search

 |

 {% block nav_map %}
 Map
 {% endblock %}

 |

 {% block nav_con %}
 Conditions
 {% endblock %}

 </div>
 {% endblock %}<hr/>
 {% block body%}{% endblock %}

 {% block nav_bottom %}
 <hr/>
 <div id="navigation_bottom">
```

```

 Search

 |

 {% block nav_map_bottom %}
 Map
 {% endblock %}

 |

 {% block nav_con_bottom %}
 Conditions
 {% endblock %}

</div>
{% endblock %}
</body>
</html>

```

**Listing B.1:** *Base mobile user interface template*

## B.1.2 Search template

```

{% extends "base_mobile.xhtml" %}
{% block nav%}{% endblock %}
{% block nav_bottom %}{% endblock %}

{% block body %}
 <form action="/mobile/geo">
 <fieldset>
 <label for="location">Get spots and
 weather info for location:</label>

 <input id="location" name="location"
 type="text" value="{% if
 latest_address %}{{latest_address
 }}{% else %}Aarhus, dk{% endif %}"
 />

 <input type="submit" value="Search" />
 </fieldset>
 </form>
{% endblock%}

```

**Listing B.2:** *Search template*

## B.1.3 Info template

```

{% extends "base_mobile.xhtml" %}
{% load datetime_filters %}
{% load inclusion_tags %}

{% block nav_con %}Conditions{% endblock %}
{% block nav_con_bottom %}Conditions{% endblock %}

```

```

{% block body %}
 <h3>Weather Stations near {{address}}</h3>
 {% if wss %}
 <table>
 {% for w in wss %}
 <tr>
 <td>
 <a href="{{ w.
 get_absolute_url }}?
 output=mobile&{{
 qs }}">{{ w.name}}
 </td>
 {% if w.current_conditions %}
 {% weather_data_as_td w.
 current_conditions %}
 <td>{{ w.current_conditions.
 time|timesince}} ago</td>
 {% endif %}
 </tr>
 {% endfor %}
 </table>
 {% else %}
 No weather stations around.
 {% endif %}
 <hr />
 <h3>Spots near {{address}}</h3>
 {% if spots %}
 <table>{% for s in spots %}
 <tr>
 <td><a href="{{s.get_absolute_url}}?
 output=mobile&{{qs}}">{{s.name
 }}</td>
 <td>{{s.forecast_point.
 current_conditions_as_object.
 to_local_time|dtformat}}</td>
 <td>{% weather_data_as_td s.
 forecast_point.
 current_conditions_as_object %}</td>
 </tr>
 {% endfor %}
 </table>
 {% else %}No spots around.{% endif %}
 <hr />
 <h3>Forecast near {{address}}</h3>
 <table>
 {% for wd in fp.get_forecasts_as_objects %}
 <tr>
 <td>{{wd.time|timeuntil}}</td>
 {% weather_data_as_td wd %}
 </tr>
 {% endfor %}
 </table>

 <p>Forecast calculated {{fp.calculation_time|timesince
 }} ago</p>
{% endblock %}

```

### B.1.4 Map template

```
{% extends "base_mobile.xhtml" %}

{% block nav_map %}Map{% endblock %}
{% block nav_map_bottom %}Map{% endblock %}

{% block body %}
<div>

 Zoom in Zoom out
</div>

{% if spots %}
<h3>Spots</h3>
{% for s in spots %}

 <p>
 {{s.name}}

 </p>

{% endfor %}

{% endif %}

{% if wss %}
<h3>Weather Station</h3>
{% for w in wss %}

 {{w.name}}

{% endfor %}

{% endif %}
{% endblock %}
```



## Appendix C

### C.1 Comparing Distances Calculated with Haversine and Cosines

```
def distance_cosines_relation(point1, point2):
 '''Calculates the distance between two points, using the
 spheric cosines formula.
 '''
 lat1, lon1 = to_radians(point1)
 lat2, lon2 = to_radians(point2)
 a = sin(lat1) * sin(lat2)
 b = cos(lat1) * cos(lat2) * cos(lon2 - lon1)
 c = acos(a + b)
 d = earth_radius * c
 return d
```

```
class Point:
 def __init__(self, lat, lon):
 self.lat = lat
 self.lon = lon

def compare():
 class Point:
 def __init__(self, lat, lon):
 self.lat = lat
 self.lon = lon
 aarhus = Point(56.0, 10.0)
 aalborg = Point(57.0, 10.0)
 for i in range(25):
 # distances in m
 cosines_distance = distance_cosines_relation(aarhus,
 aalborg)
 haversine_distance = distance(aarhus, aalborg)
 difference = cosines_distance - haversine_distance
```

```

print 'P2: (%s,%s)\nHDist: %sm\tCDist:%sm\nDiff:%sm' %
\
(aalborg.lat, aalborg.lon,
 haversine_distance * 1000, cosines_distance * 1000,
 difference * 1000)
print '-----'
half distance
aalborg.lat = aarhus.lat + (aalborg.lat - aarhus.lat) /
2

```

```

P2: (57.0,10.0)
HDist: 111125.113474m CDist:111125.113474m
Diff:-1.12549969344e-008m

P2: (56.5,10.0)
HDist: 55562.5567372m CDist:55562.5567371m
Diff:-9.63567003964e-008m

P2: (56.25,10.0)
HDist: 27781.2783686m CDist:27781.2783686m
Diff:-5.07327513333e-009m

P2: (56.125,10.0)
HDist: 13890.6391843m CDist:13890.6391843m
Diff:3.37667671602e-008m

P2: (56.0625,10.0)
HDist: 6945.31959215m CDist:6945.31959159m
Diff:-5.66108049327e-007m

P2: (56.03125,10.0)
HDist: 3472.65979608m CDist:3472.65979594m
Diff:-1.33233424293e-007m

P2: (56.015625,10.0)
HDist: 1736.32989804m CDist:1736.32989803m
Diff:-7.03592739626e-009m

P2: (56.0078125,10.0)
HDist: 868.164949019m CDist:868.164945702m
Diff:-3.31756488947e-006m

P2: (56.00390625,10.0)
HDist: 434.08247451m CDist:434.082477783m
Diff:3.27316213022e-006m

P2: (56.001953125,10.0)
HDist: 217.041237255m CDist:217.041228492m
Diff:-8.76289446561e-006m

P2: (56.0009765625,10.0)
HDist: 108.520618627m CDist:108.520583137m
Diff:-3.54903024608e-005m

P2: (56.0004882813,10.0)
HDist: 54.2603093133m CDist:54.2602708314m
Diff:-3.84819070012e-005m

```



```

P2: (56.0002441406,10.0)
HDist: 27.1301546566m CDist:27.1301354156m
Diff:-1.9241015066e-005m

P2: (56.0001220703,10.0)
HDist: 13.565077328m CDist:13.5649018138m
Diff:-0.0001755141407m

P2: (56.0000610352,10.0)
HDist: 6.78253866398m CDist:6.78211910881m
Diff:-0.00041955517375m

P2: (56.0000305176,10.0)
HDist: 3.39126933199m CDist:3.39039587702m
Diff:-0.000873454969691m

P2: (56.0000152588,10.0)
HDist: 1.69563466564m CDist:1.69453406618m
Diff:-0.00110059946489m

P2: (56.0000076294,10.0)
HDist: 0.847817332821m CDist:0.848593998715m
Diff:0.000776665893623m

P2: (56.0000038147,10.0)
HDist: 0.423908666057m CDist:0.413553559386m
Diff:-0.0103551066715m

P2: (56.0000019073,10.0)
HDist: 0.211954333029m CDist:0.212148499679m
Diff:0.000194166650125m

P2: (56.0000009537,10.0)
HDist: 0.105977166514m CDist:0.0948756933212m
Diff:-0.0111014731931m

P2: (56.0000004768,10.0)
HDist: 0.0529885829037m CDist:0.0m
Diff:-0.0529885829037m

P2: (56.0000002384,10.0)
HDist: 0.0264942914519m CDist:0.0m
Diff:-0.0264942914519m

P2: (56.0000001192,10.0)
HDist: 0.0132471457259m CDist:0.0m
Diff:-0.0132471457259m

P2: (56.0000000596,10.0)
HDist: 0.00662357286296m CDist:0.0m
Diff:-0.00662357286296m

```

